What's hidden in the tails? Revealing and reducing optimistic bias in entropic risk estimation and optimization

Utsav Sadana

Department of Computer Science and Operations Research
Université de Montréal

INFORMS Computing Society conference,

Ides of March, 2025

(joint work with Erick Delage and Angelos Georghiou)

It is not a calculated risk if you haven’t calculated it.

- Naved Abdali


What this talk is about? Tails and Bias correction

What this talk is about? Tails and Bias correction

- Uncertain loss
- Risk measure: Map loss to a real number
- Entropic risk measure:
  - mean
  - variance
  - Higher moments

- Estimation
  - True risk - Use known loss distribution
  - We have data - construct risk estimator

What this talk is about? Tails and Bias correction

- Uncertain loss
- Risk measure: Map loss to a real number
- Entropic risk measure:
  - mean
  - variance
  - Higher moments

- Estimation
  - True risk - Use known loss distribution
  - We have data - construct risk estimator

True risk

What this talk is about? Tails and Bias correction

- Uncertain loss
- Risk measure: Map loss to a real number
- Entropic risk measure:
  - mean
  - variance
  - Higher moments

- Estimation
  - True risk - Use known loss distribution
  - We have data - construct risk estimator

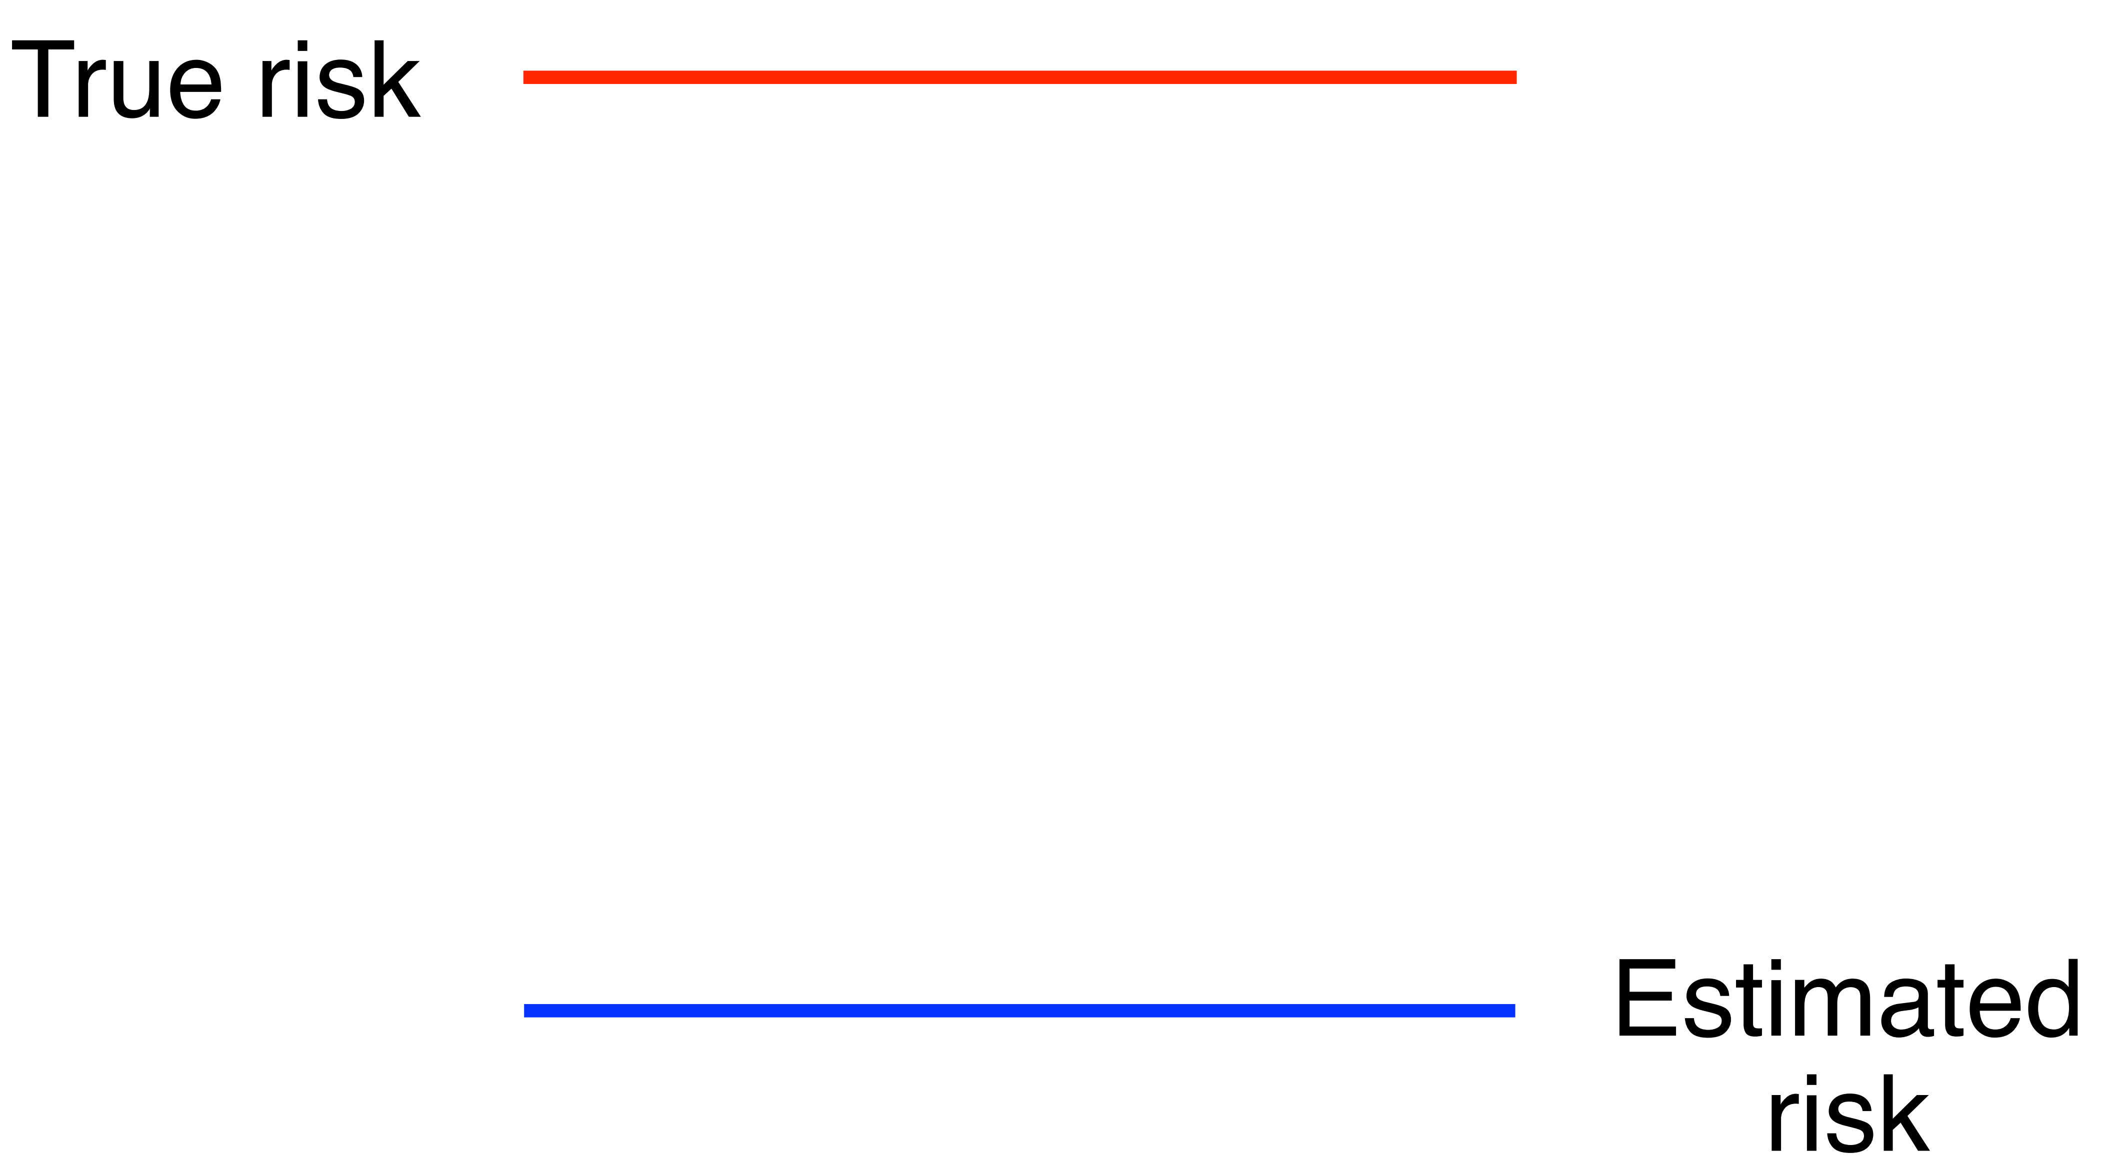

What this talk is about? Tails and Bias correction

- Uncertain loss
- Risk measure: Map loss to a real number
- Entropic risk measure:
  - mean
  - variance
  - Higher moments

- Estimation
  - True risk - Use known loss distribution
  - We have data - construct risk estimator

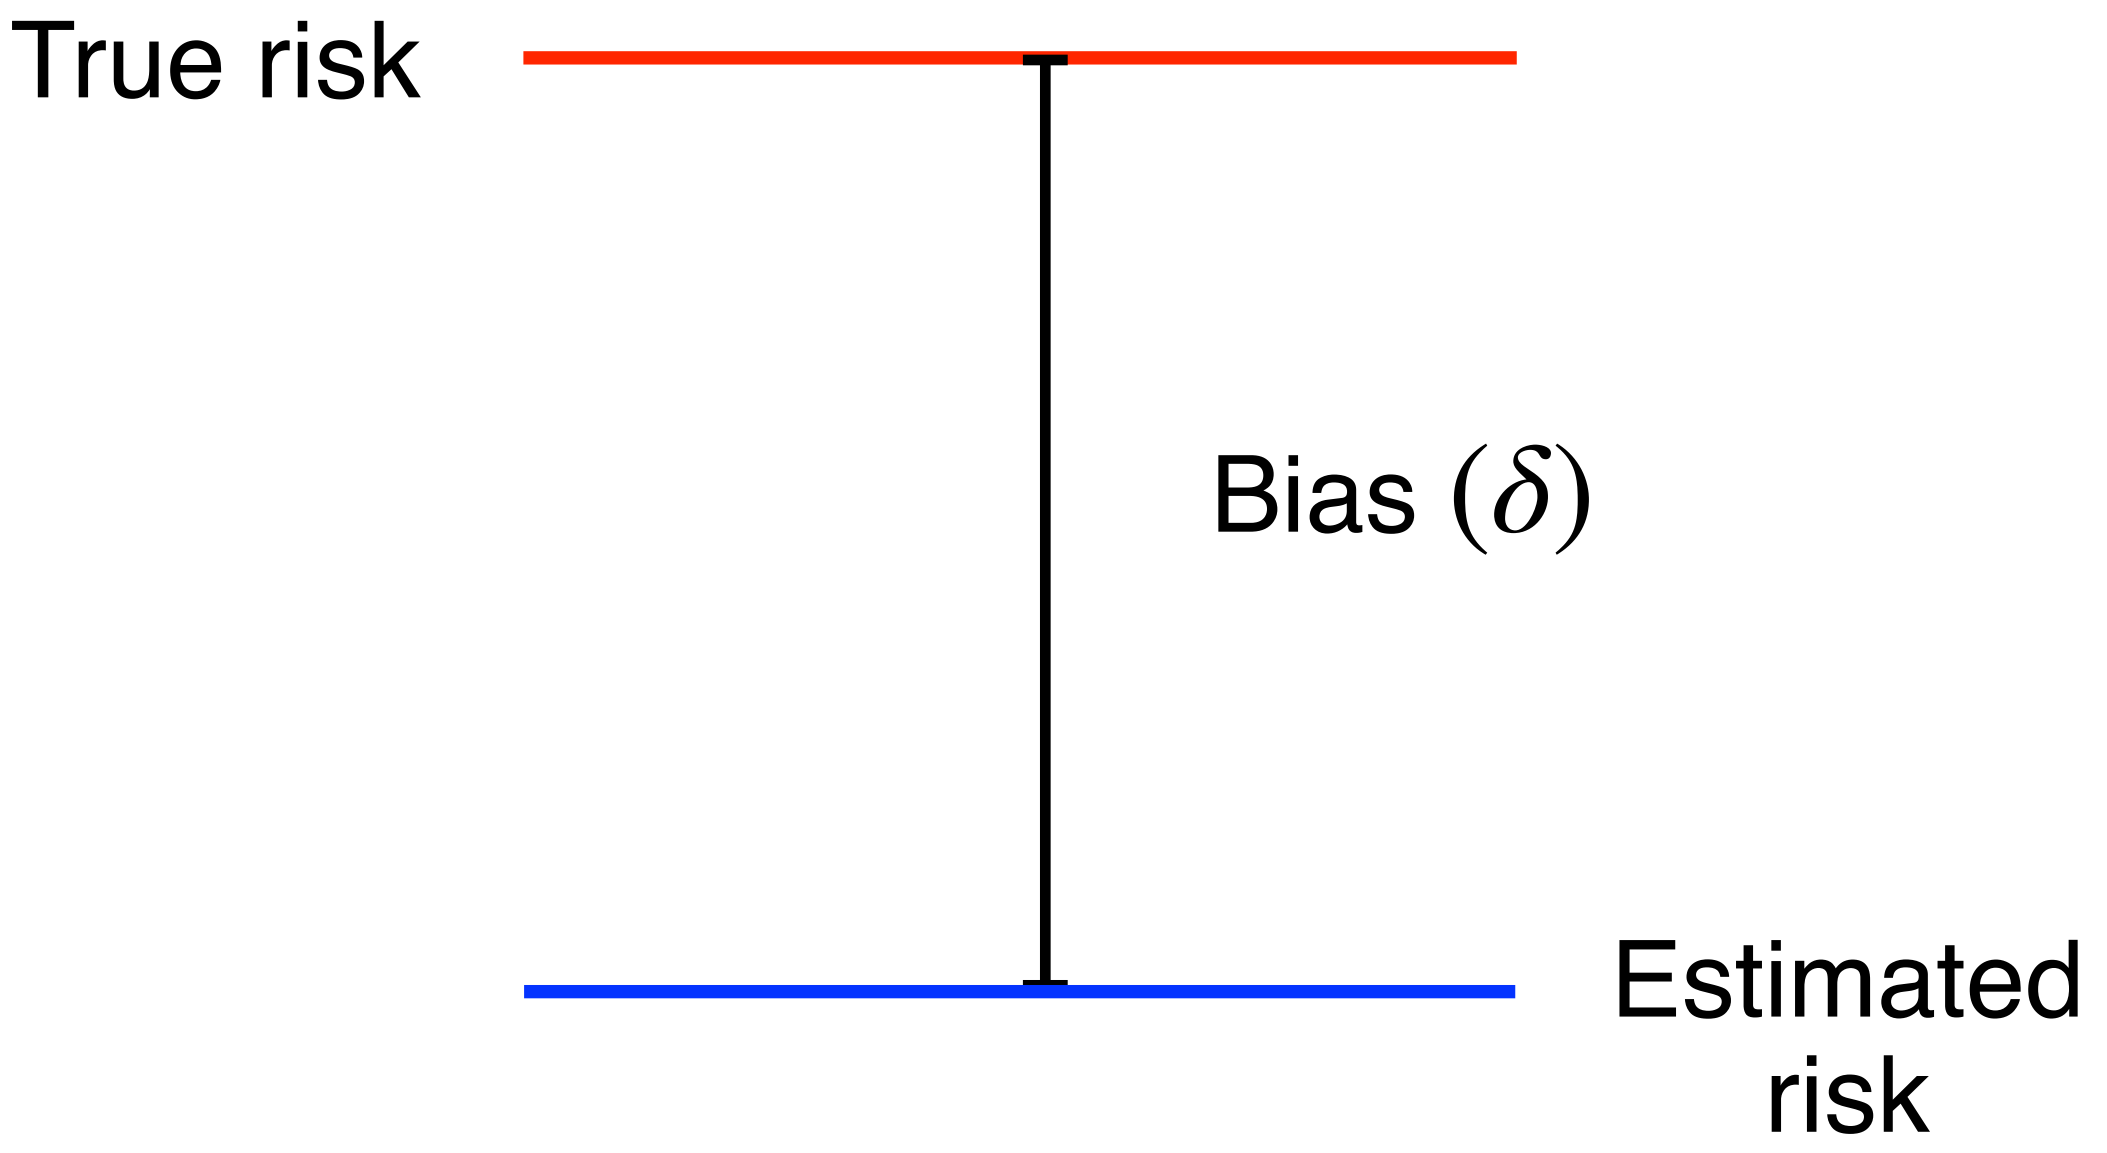

What this talk is about? Tails and Bias correction

- Uncertain loss
- Risk measure: Map loss to a real number
- Entropic risk measure:
  - mean
  - variance
  - Higher moments

- Estimation
  - True risk - Use known loss distribution
  - We have data - construct risk estimator

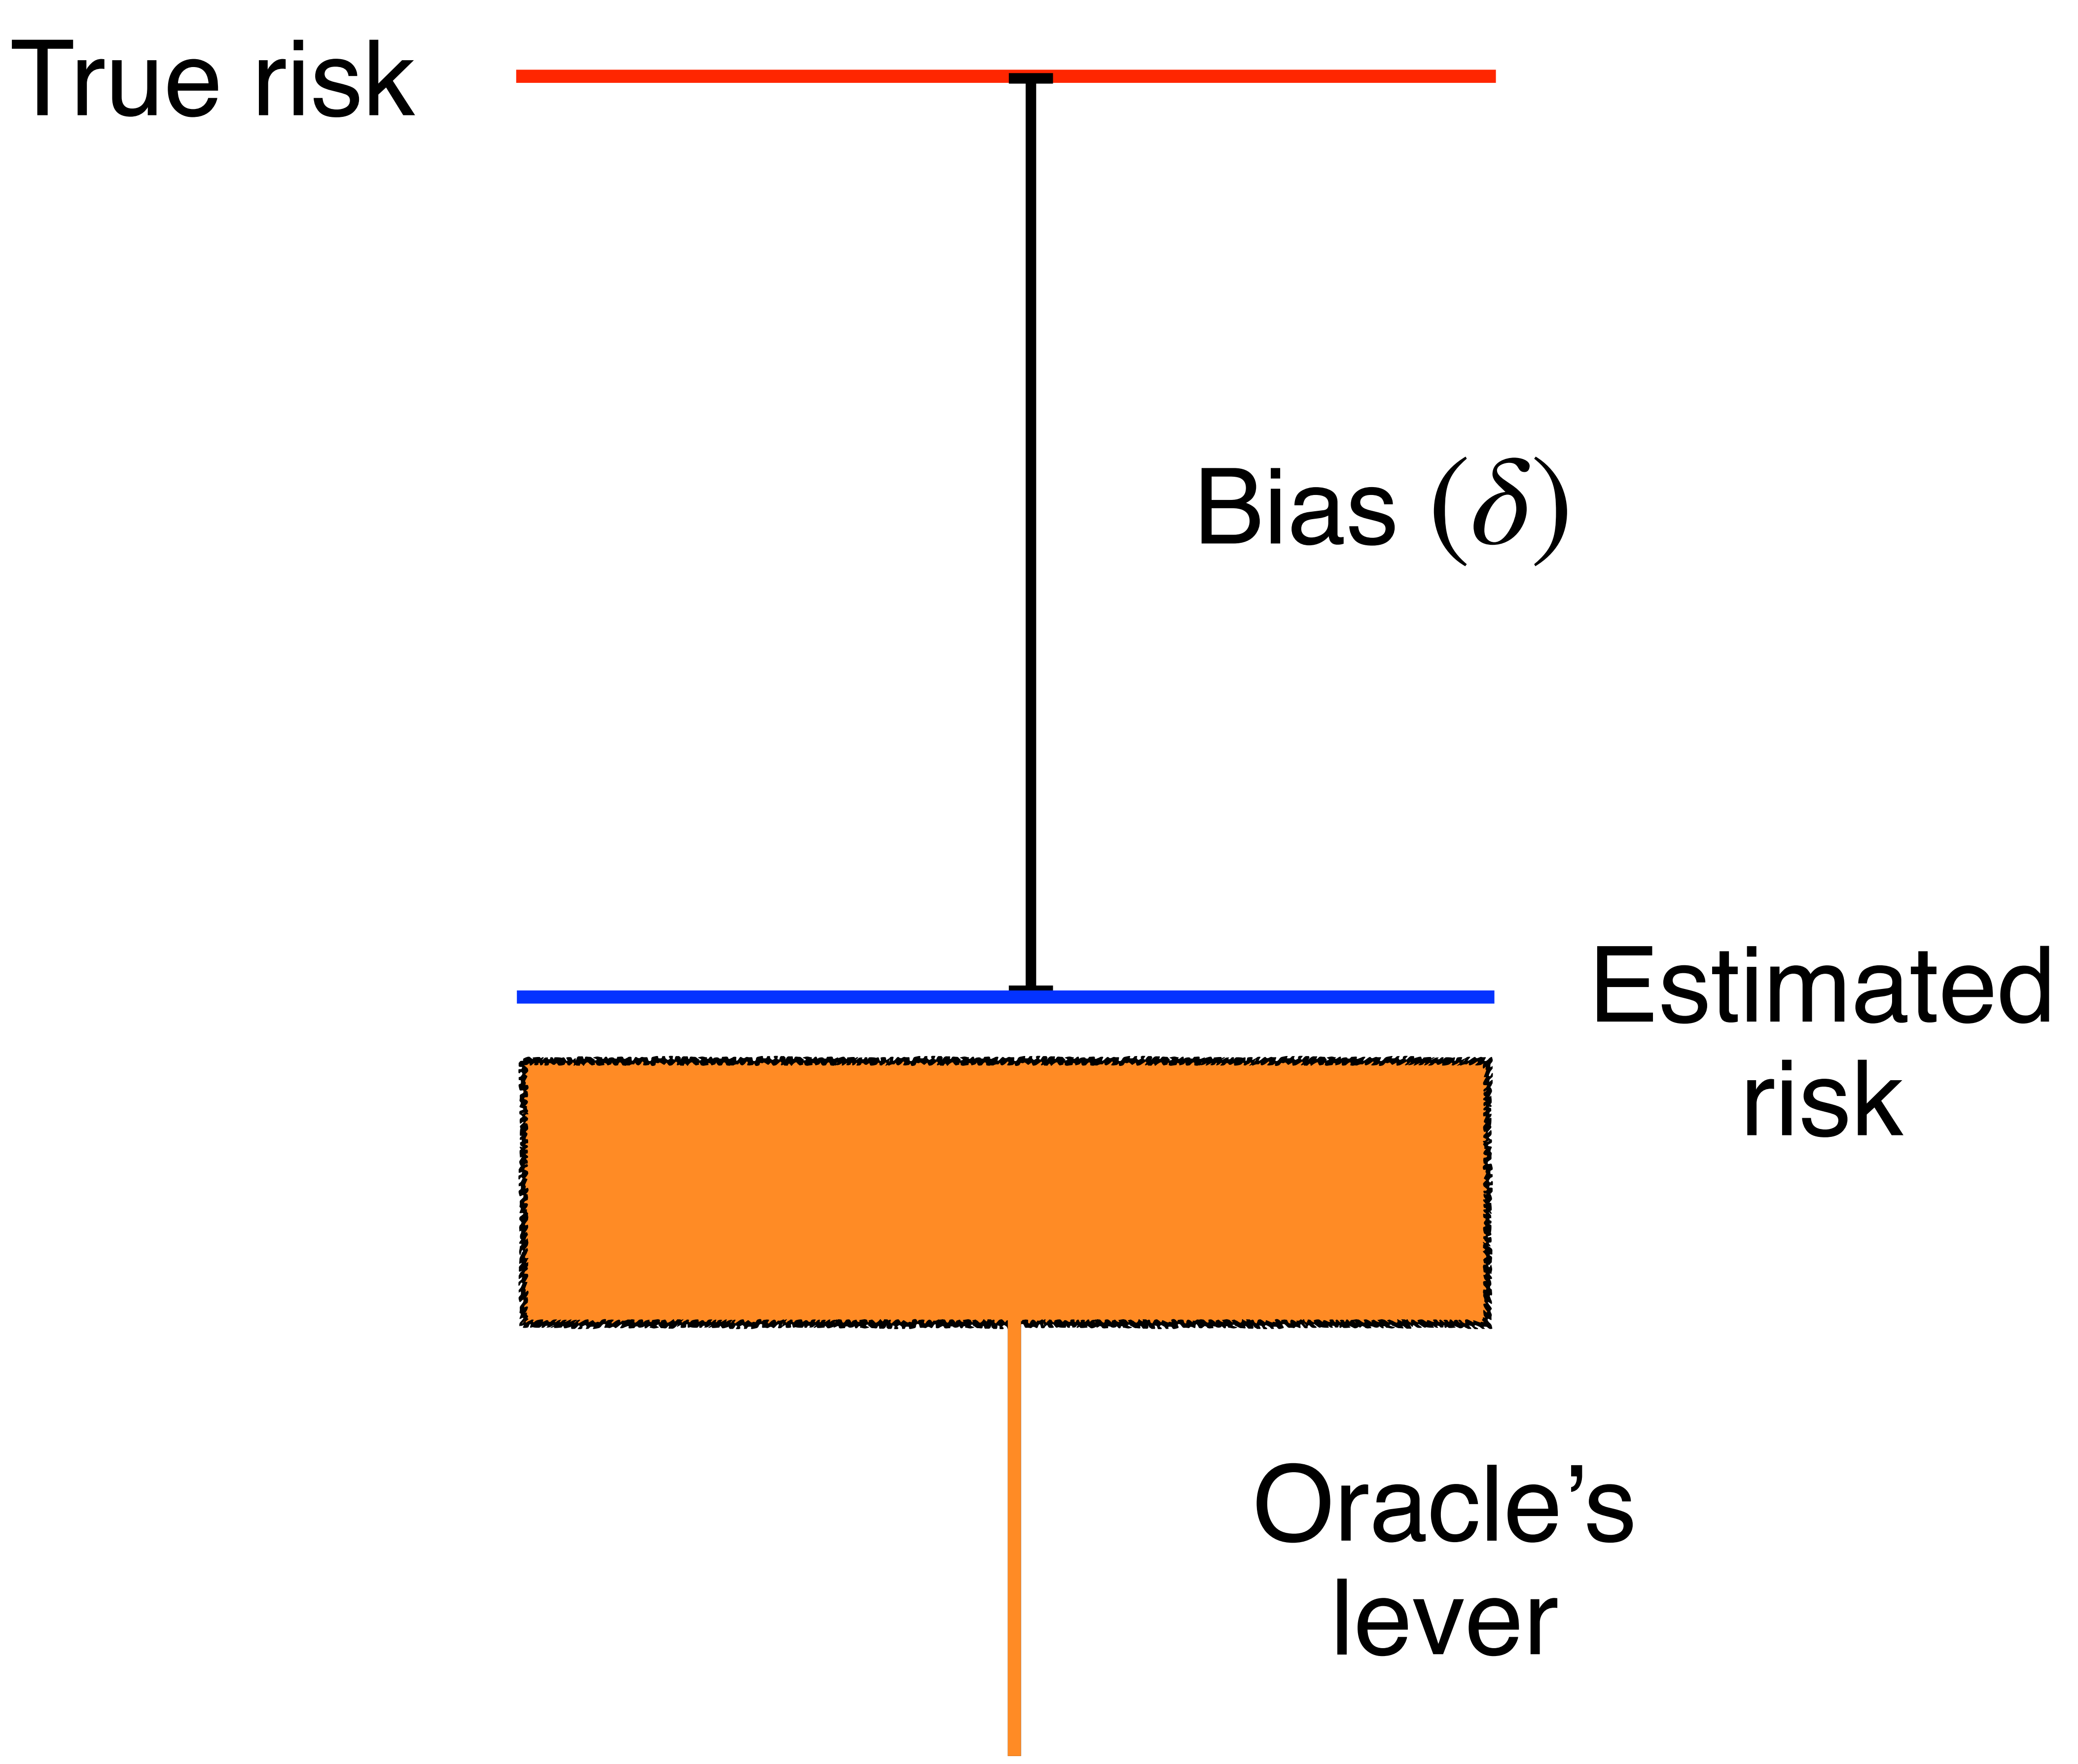

What this talk is about? Tails and Bias correction

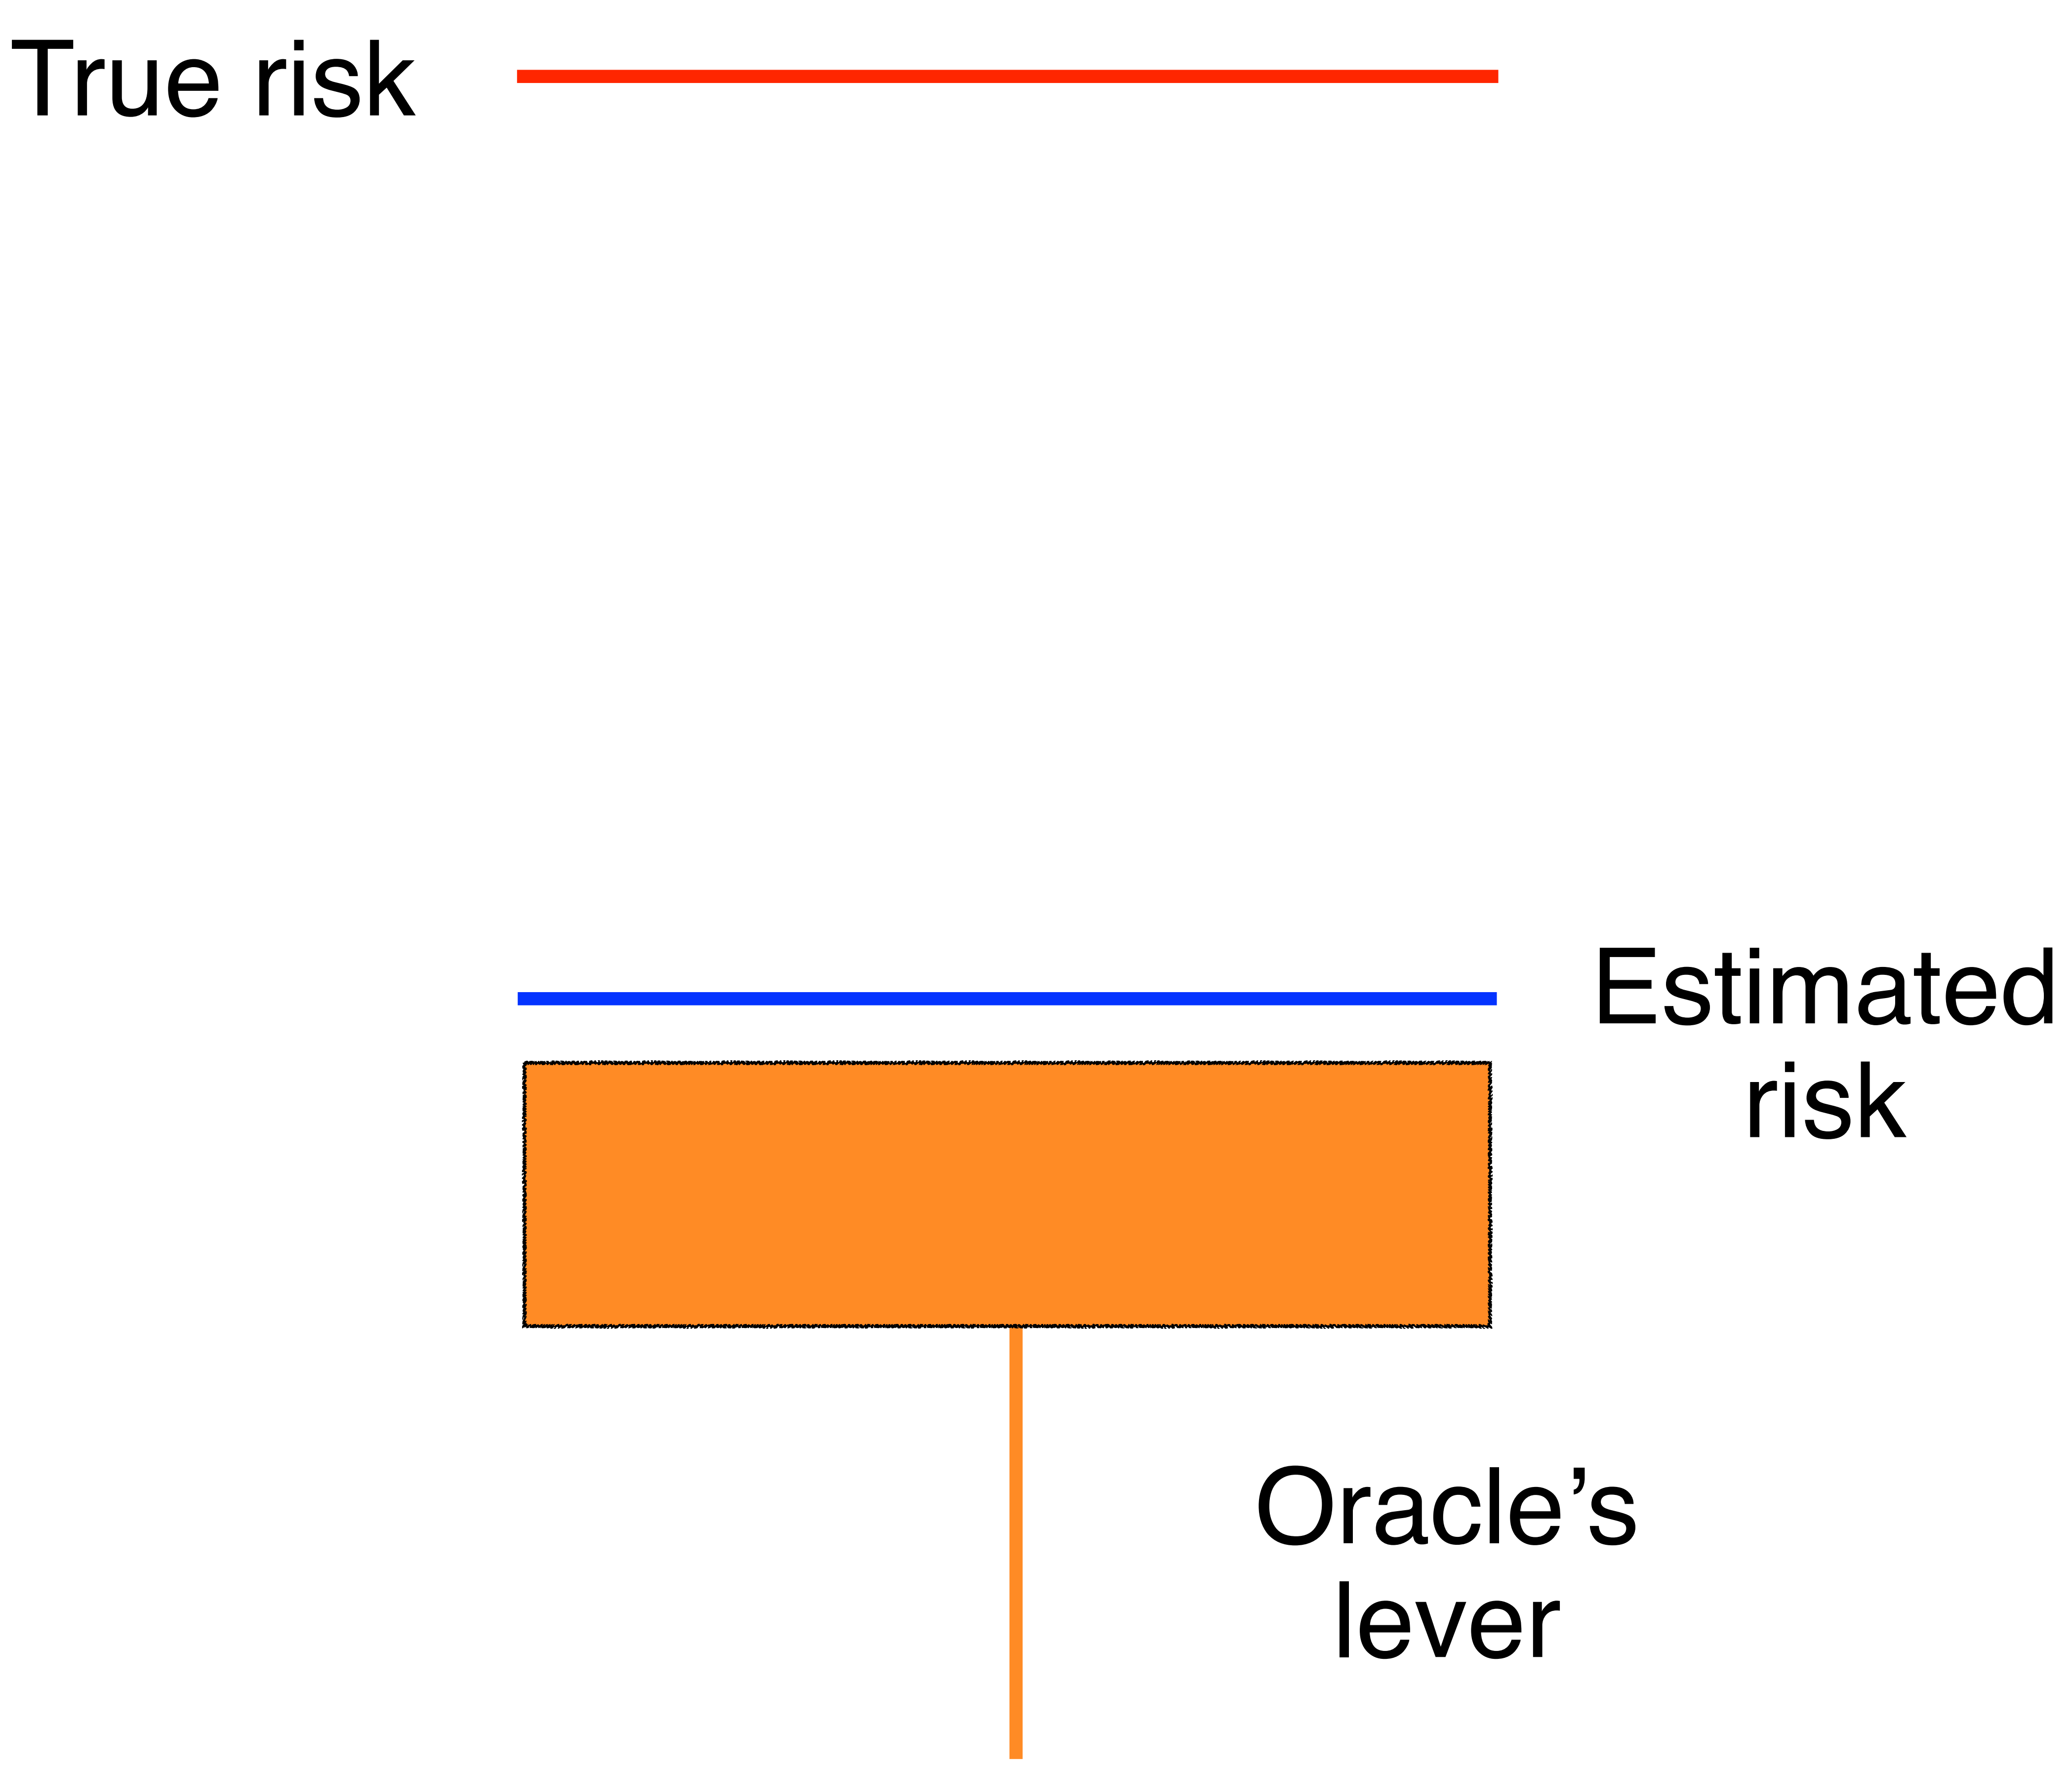

What this talk is about? Tails and Bias correction

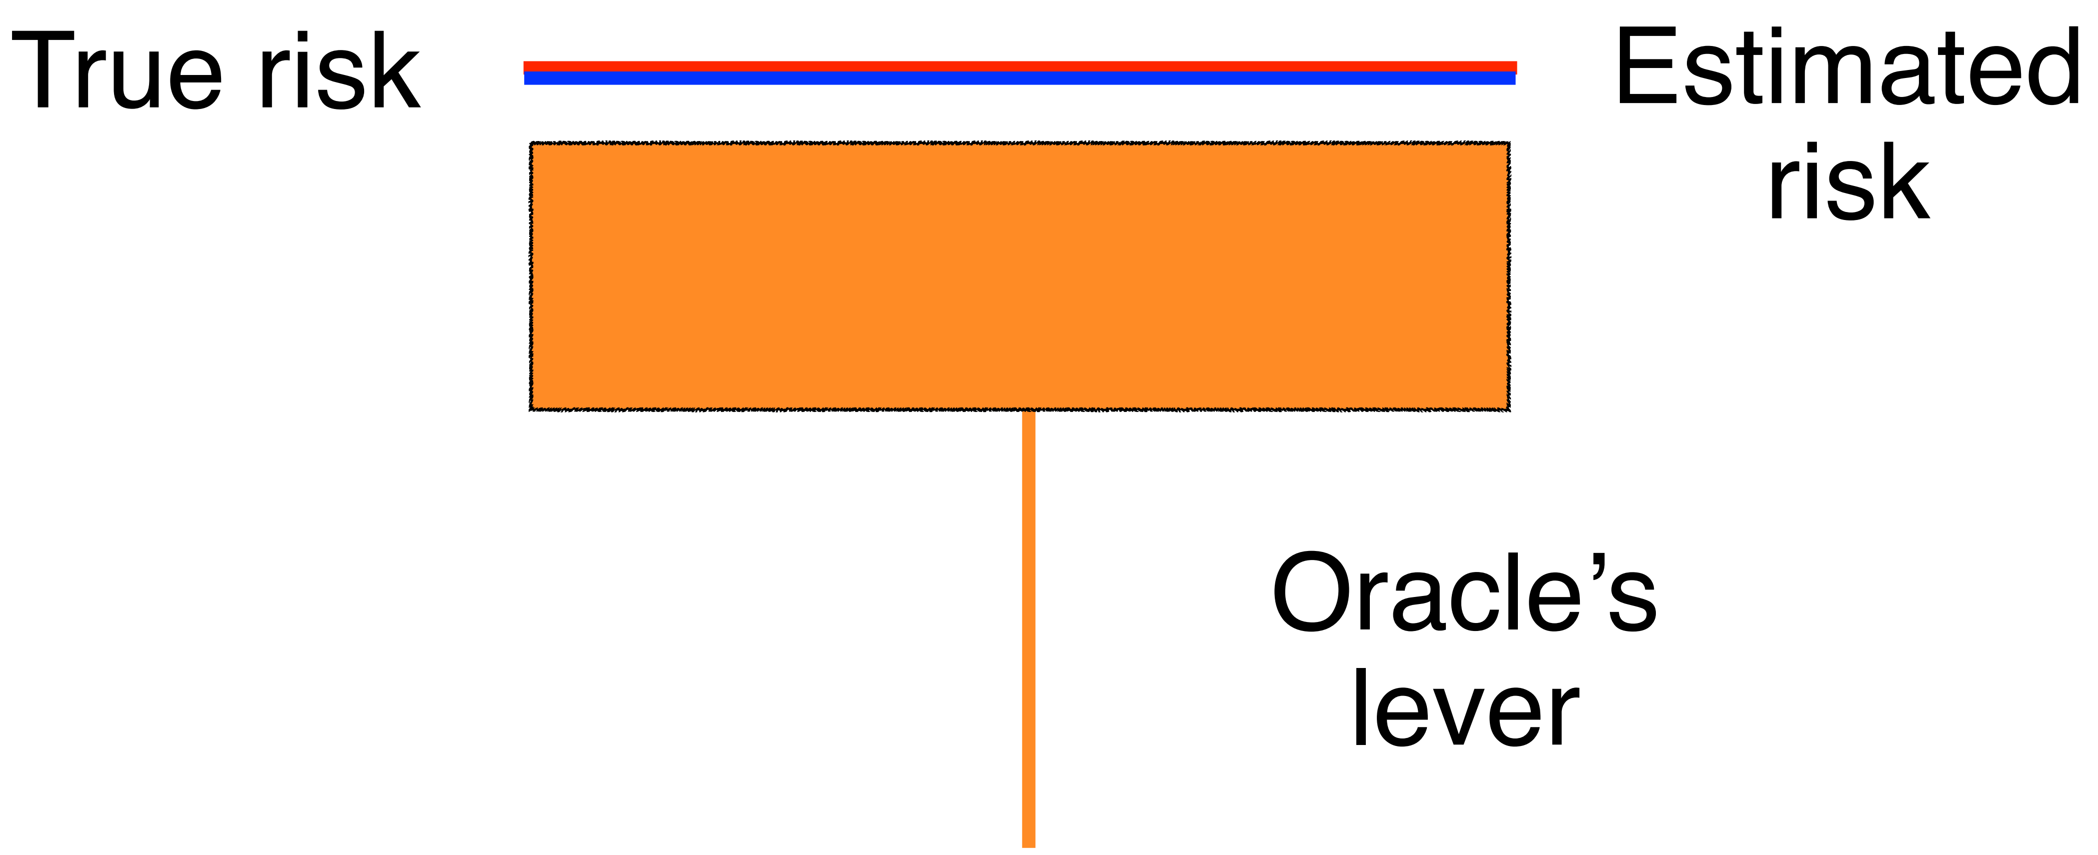

What this talk is **really** about? Tails and bias **mitigation**

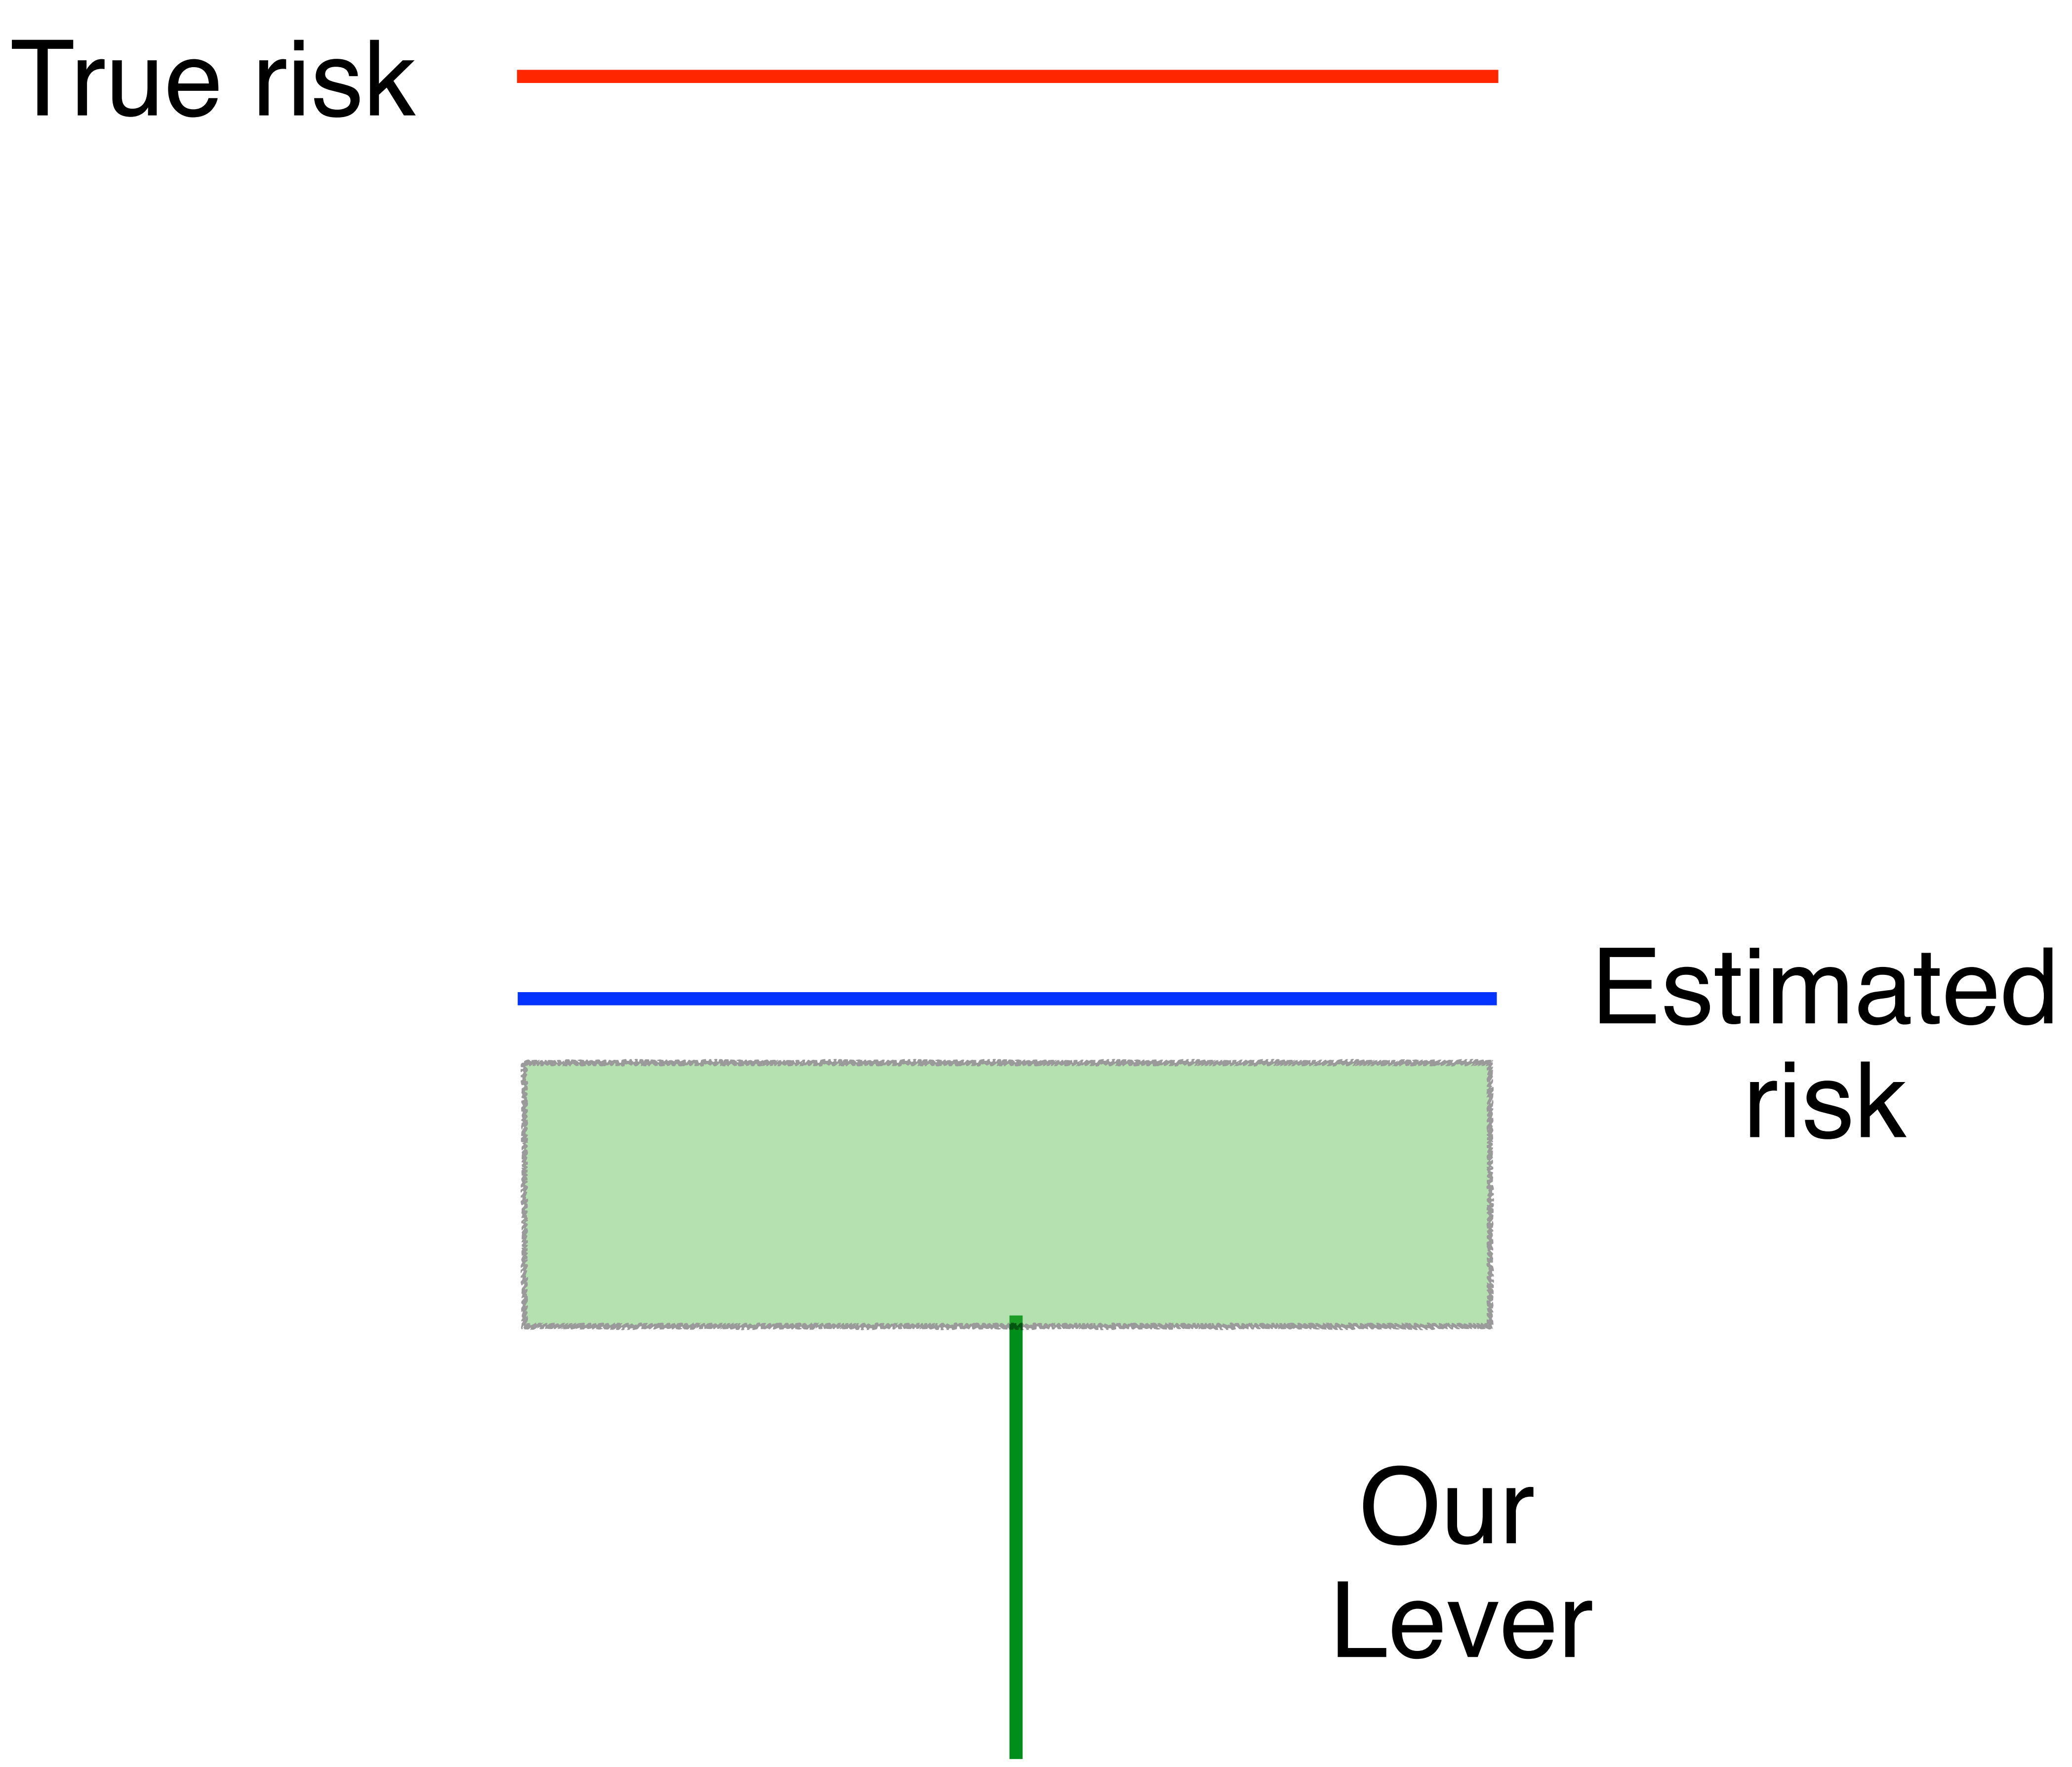

What this talk is really about? Tails and bias mitigation

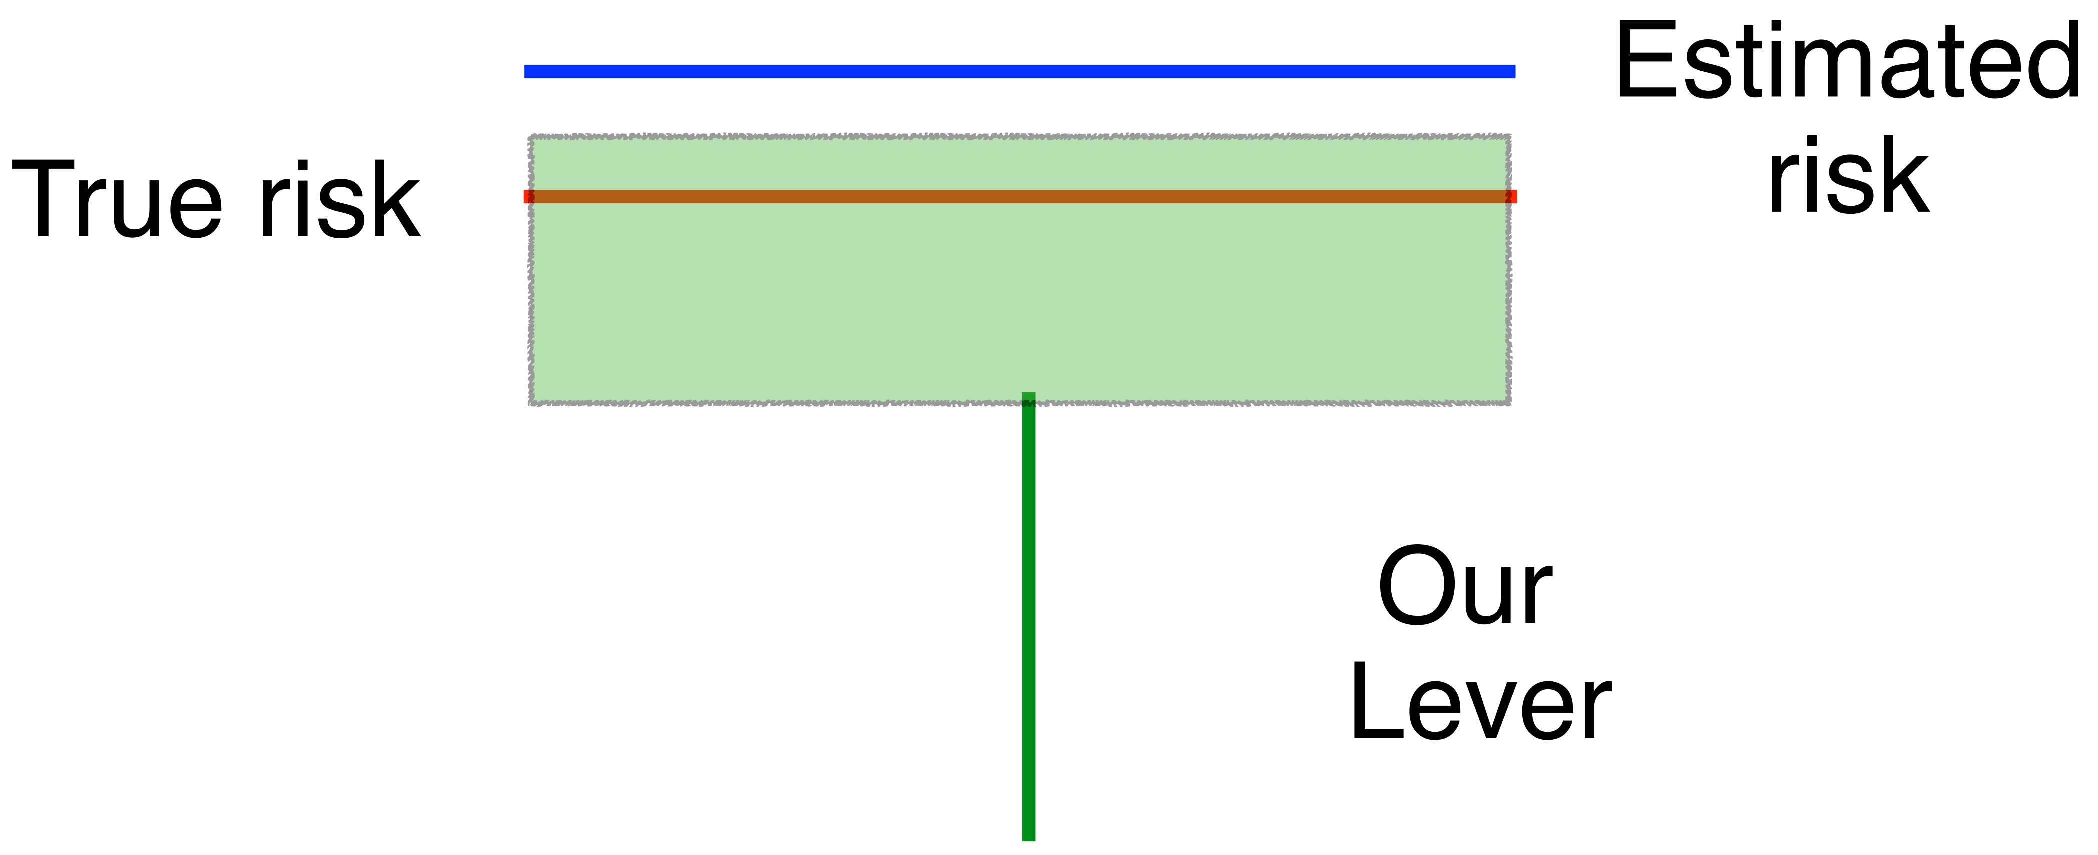

Beyond risk neutrality

Beyond risk neutrality

Beyond risk neutrality

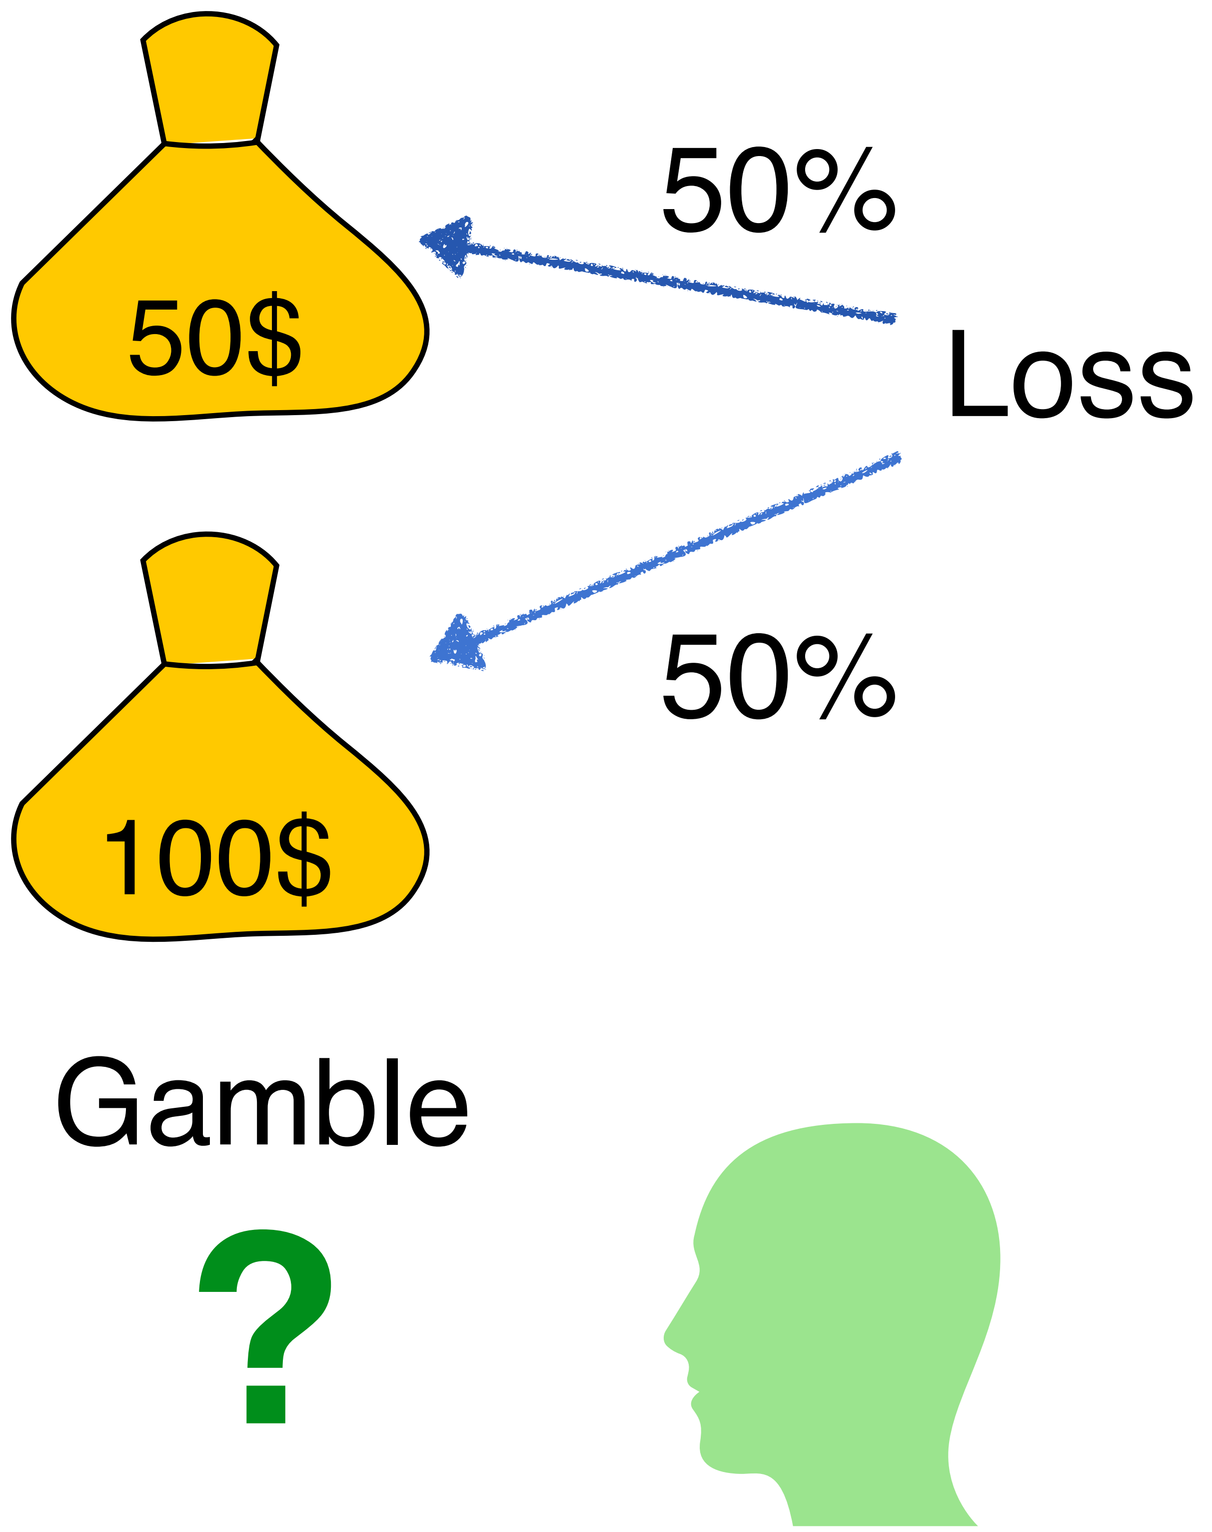

Beyond risk neutrality

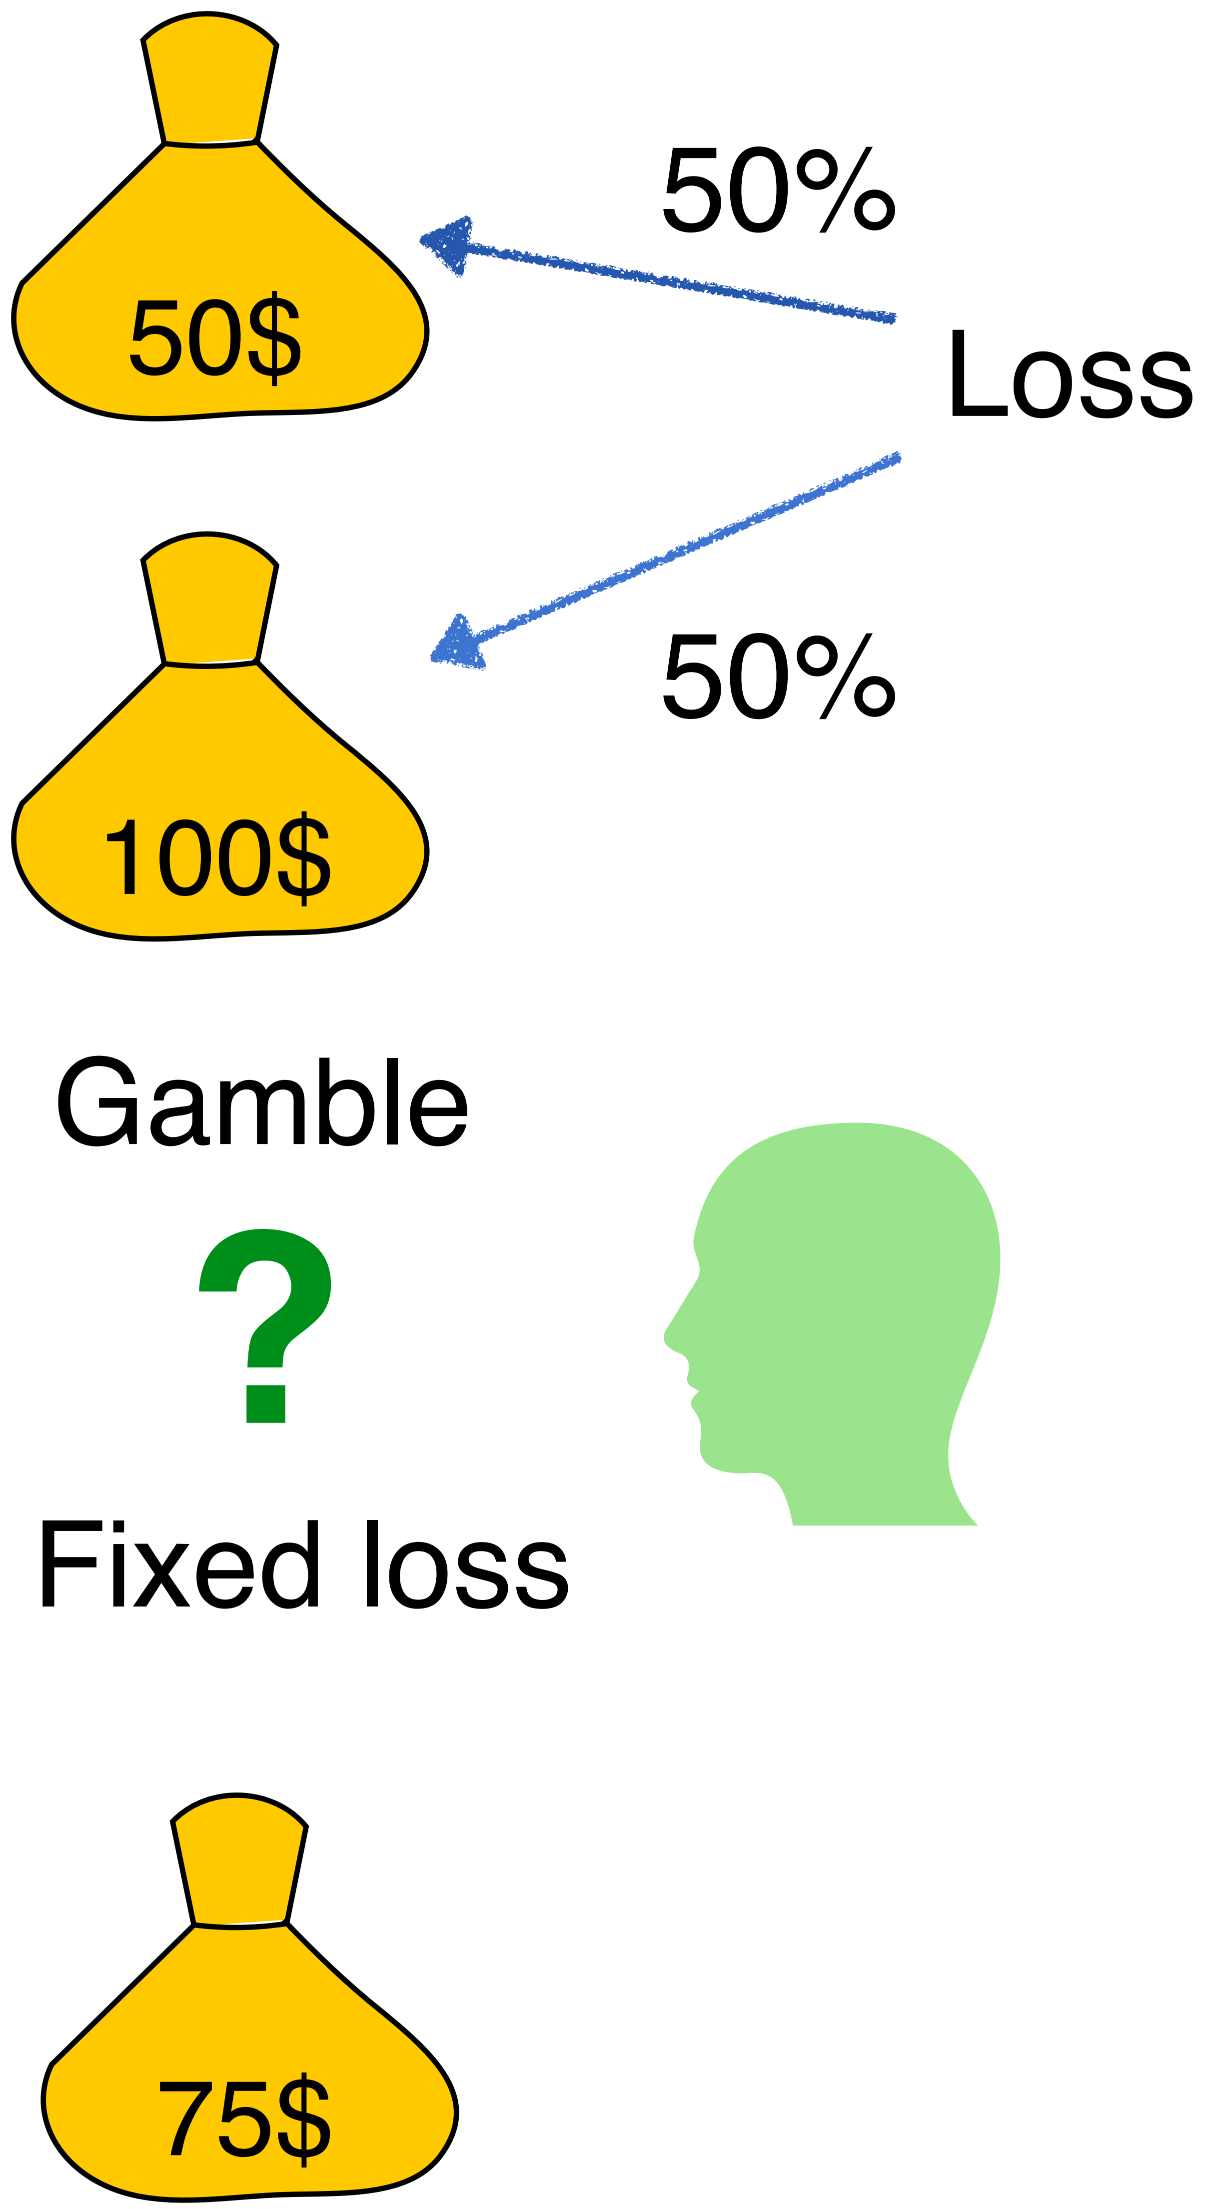

Beyond risk neutrality

Indifference between the two options

- [x] Risk neutral
- [ ] Experiments
- [ ] Entropic risk measure

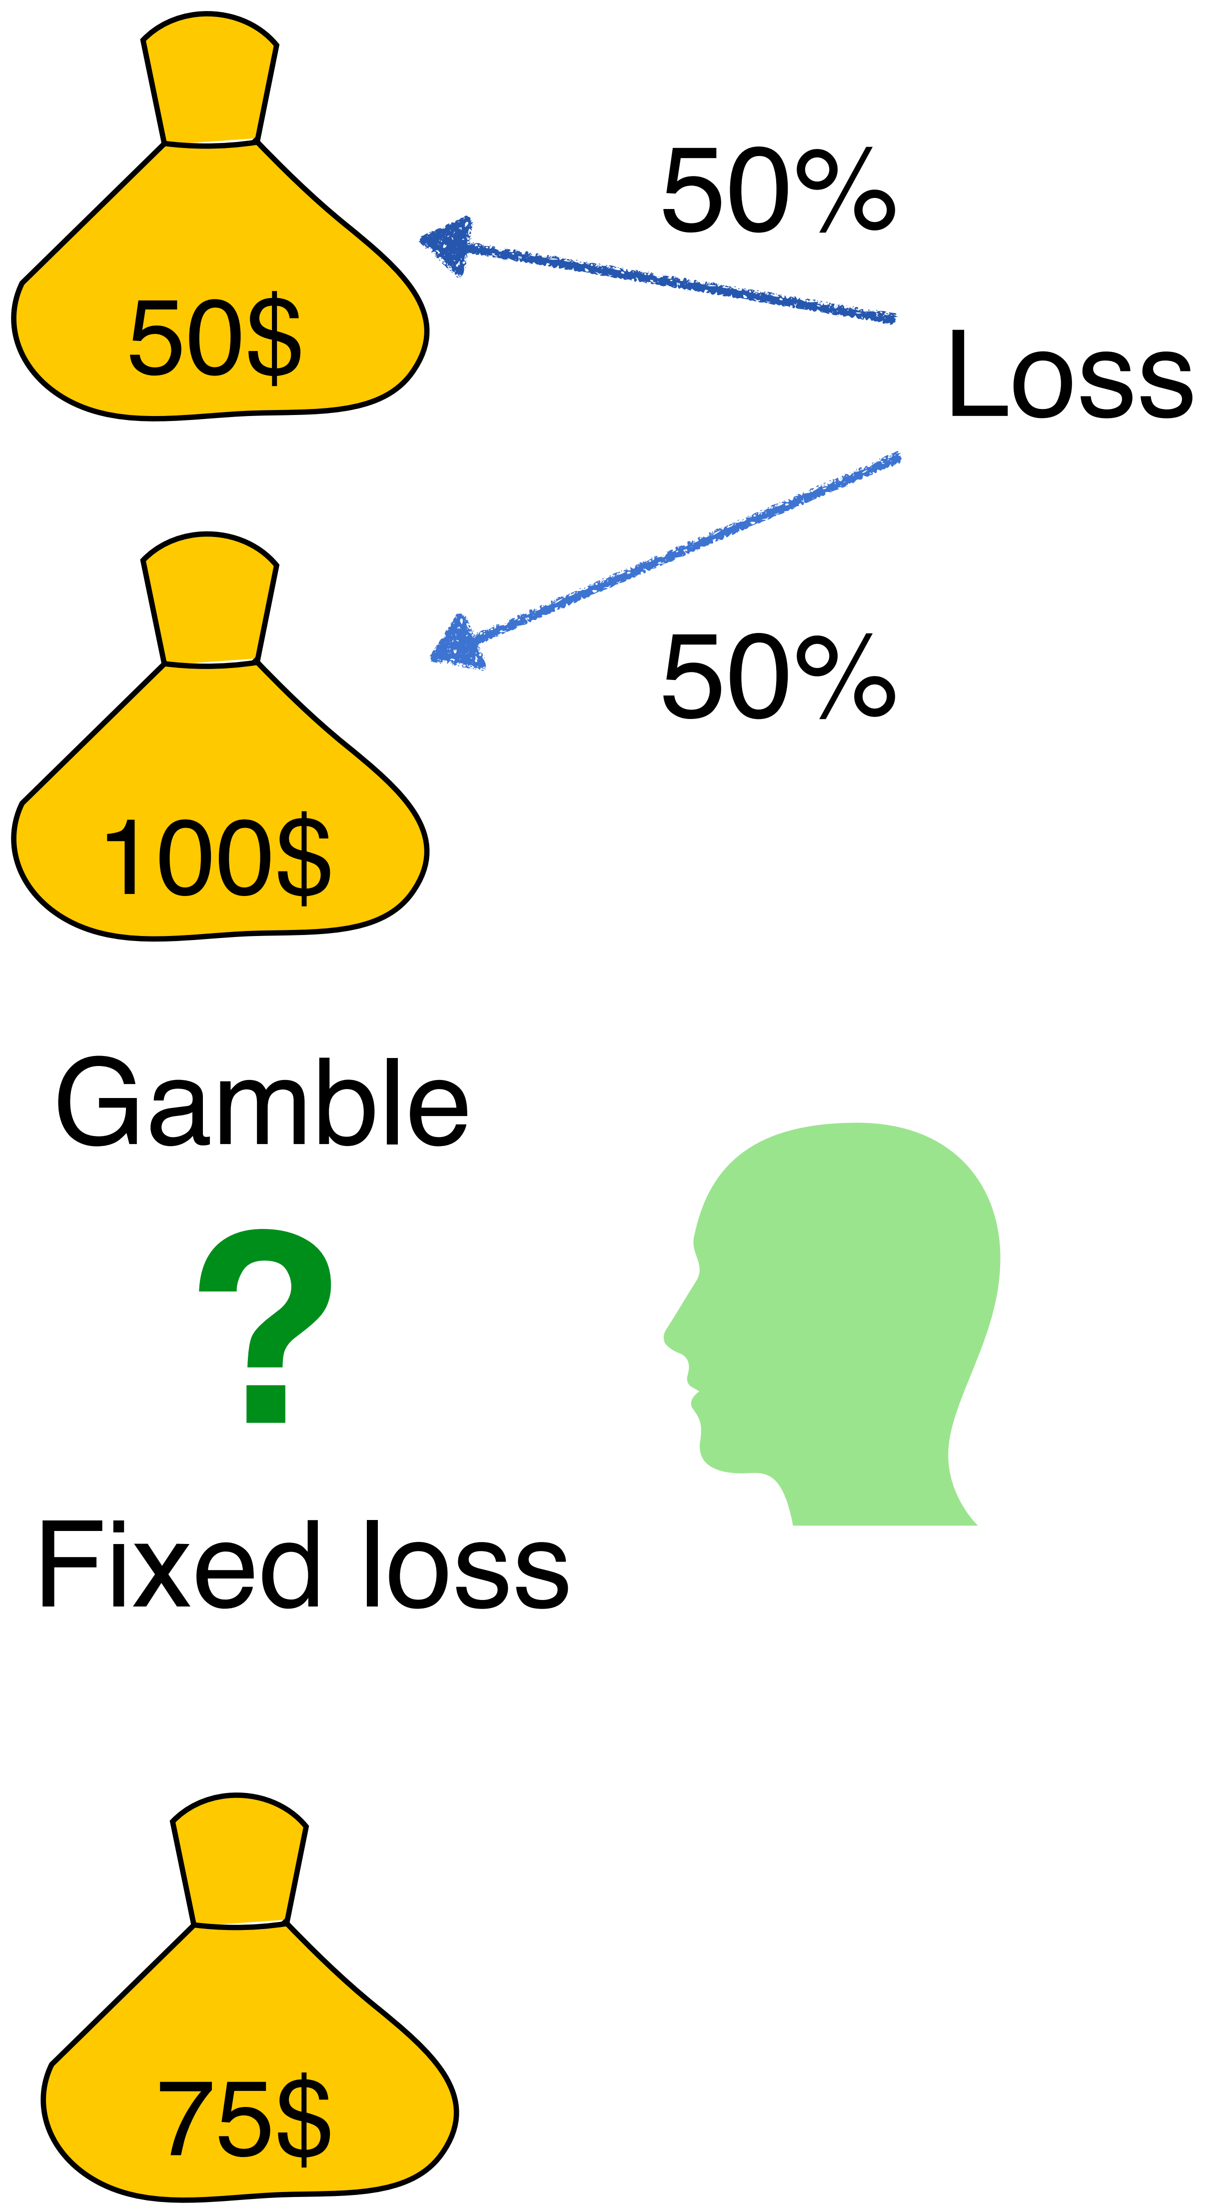

Beyond risk neutrality

Indifference between the two options

- Risk neutral
- Experiments
- Entropic risk measure

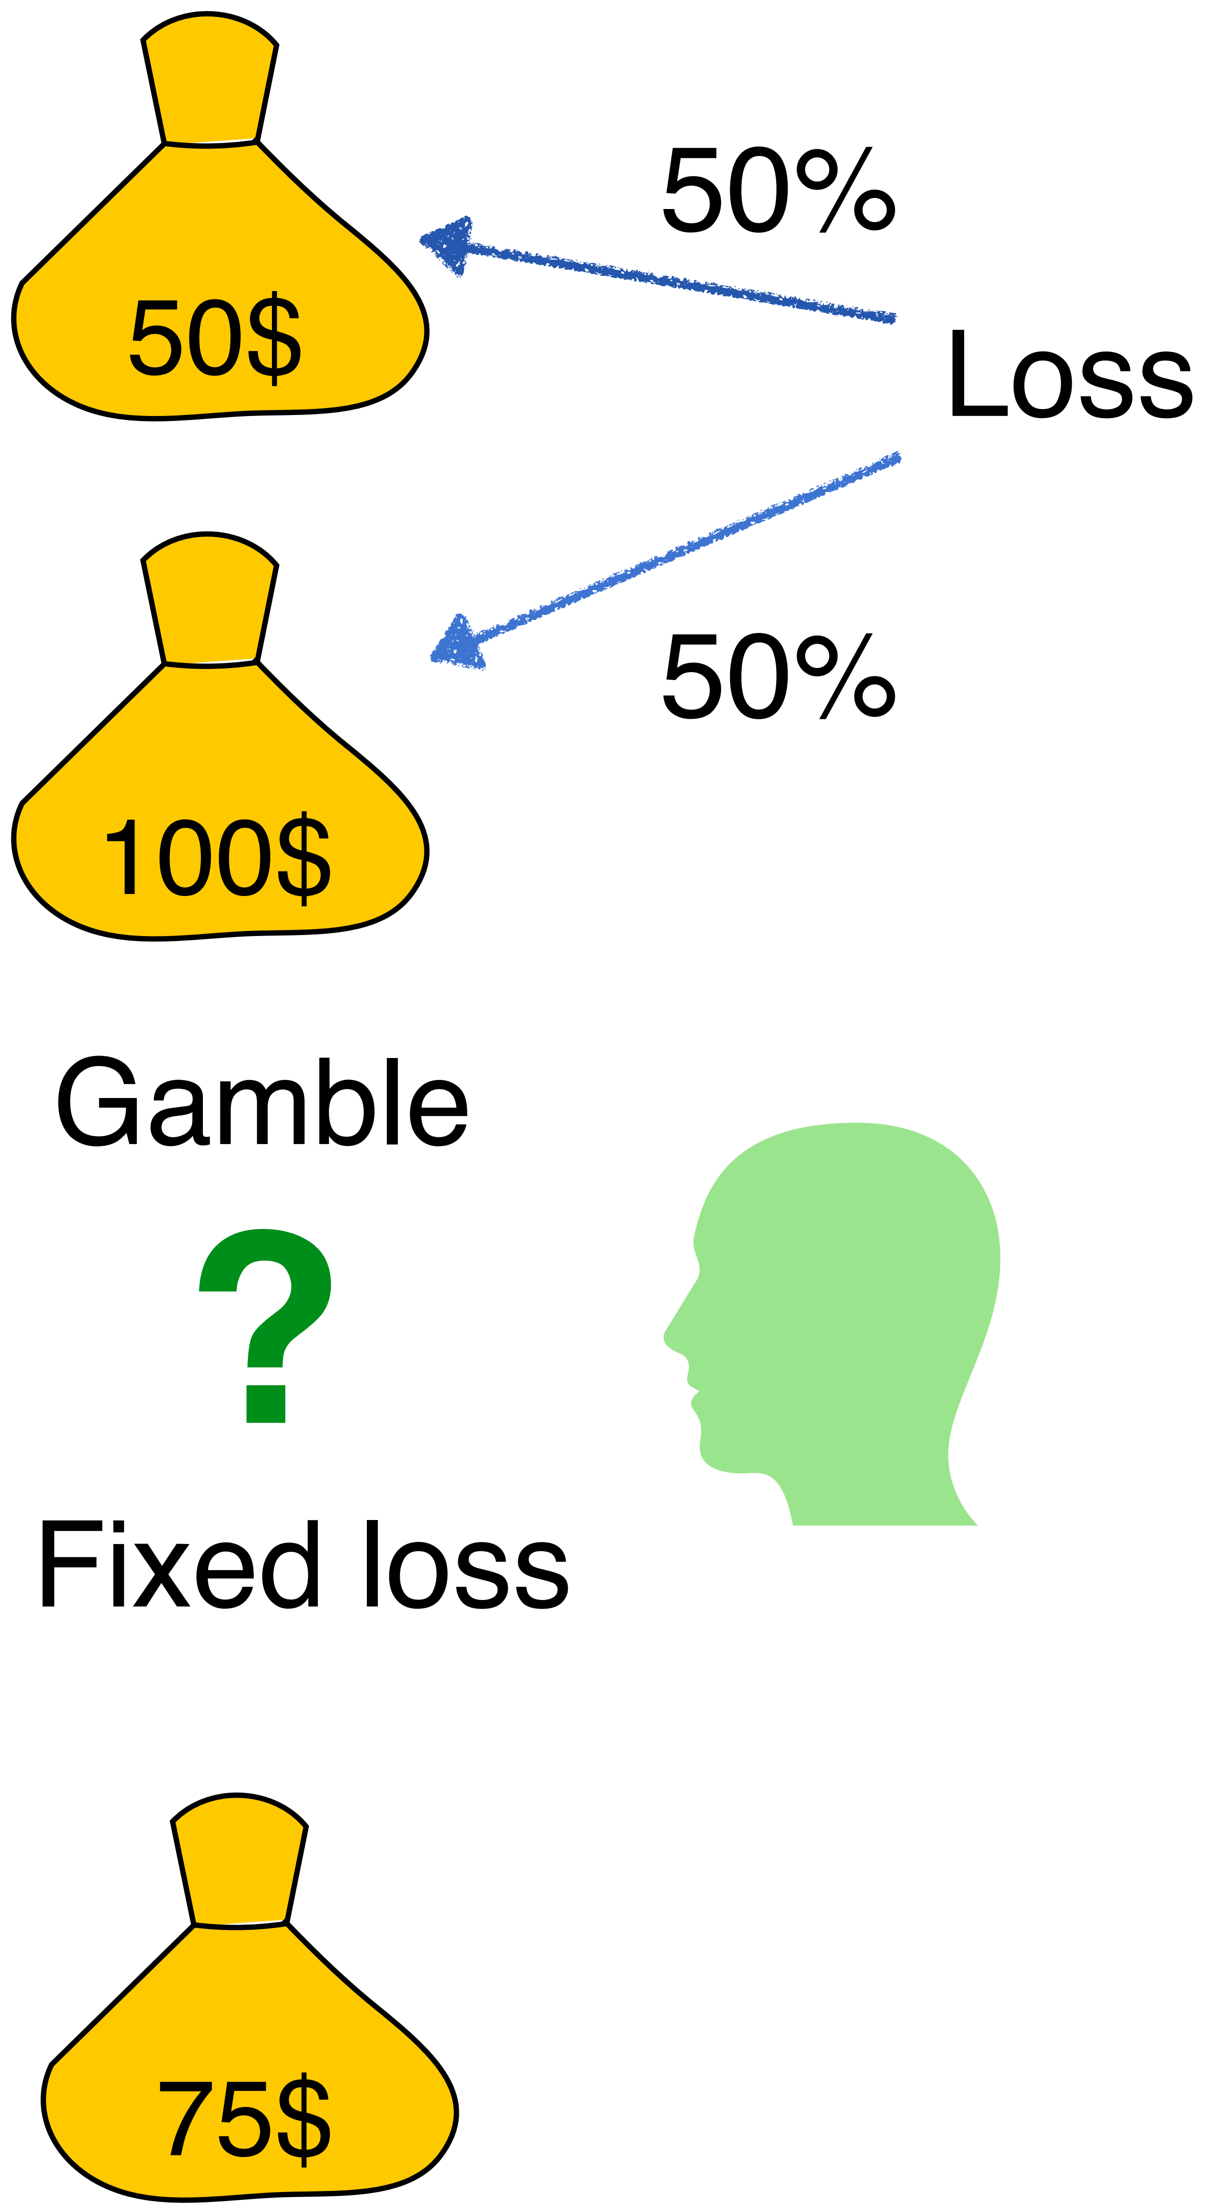

Beyond risk neutrality

Indifference between the two options

- ☑ Risk neutral
- ☒ Experiments
- ☒ Entropic risk measure

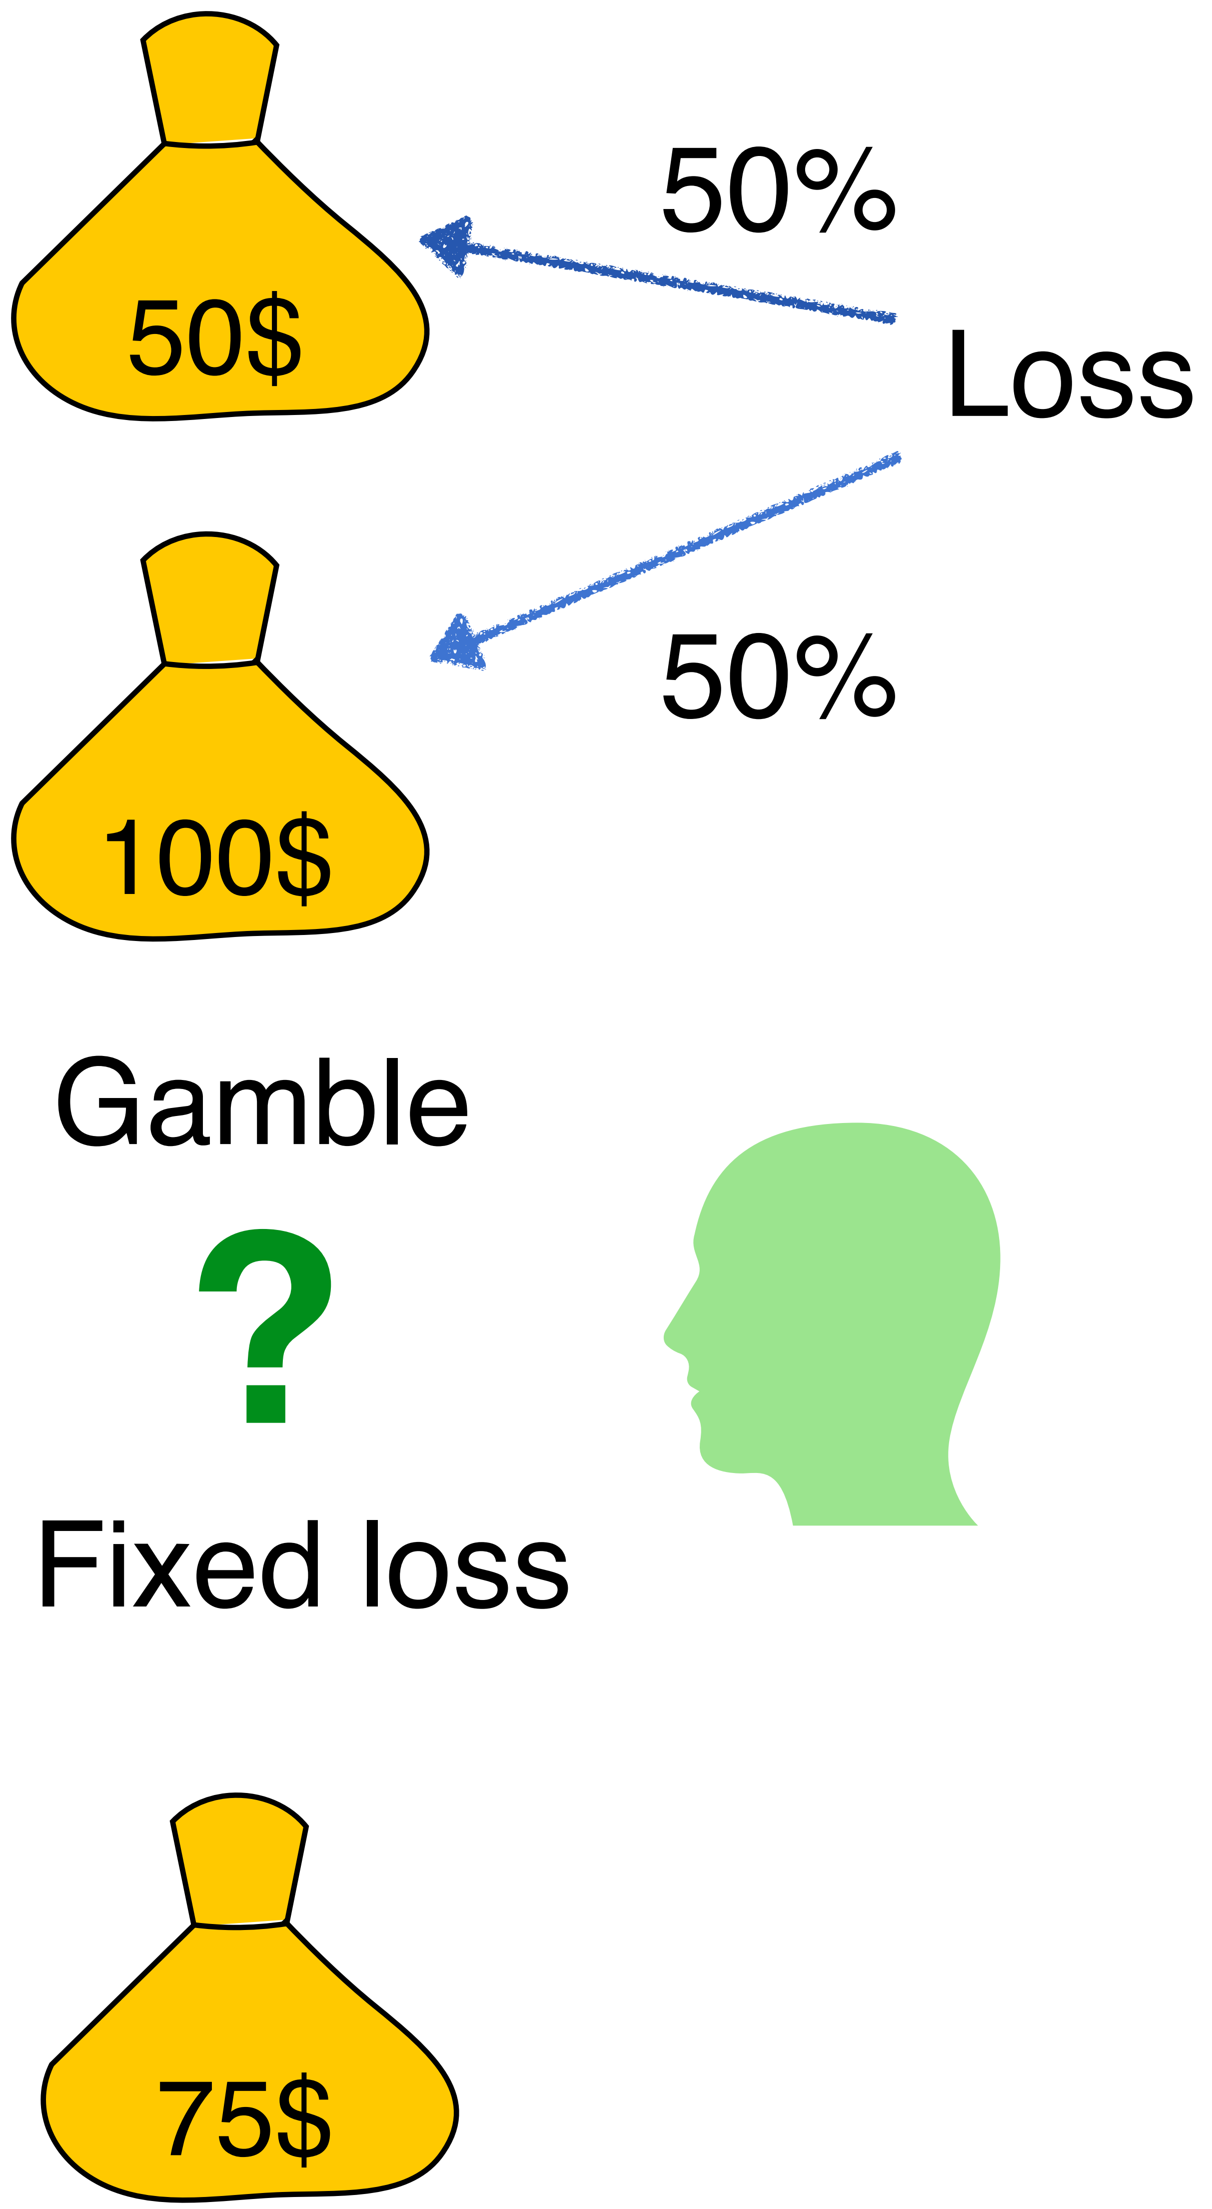

Entropic risk measure

Uncertainty

$$\boldsymbol{\xi} \sim \mathbb{P}$$

Loss

$$\ell(\boldsymbol{\xi})$$

Entropic risk

$$\rho_{\mathbb{P}}(\ell(\boldsymbol{\xi})) := \frac{1}{\alpha} \log \left(\mathbb{E}_{\mathbb{P}} \left[\exp(\alpha \ell(\boldsymbol{\xi})) \right] \right)$$

- α is the decision maker's risk aversion
- $\mathbb{P}$ is not known

Empirical entropic risk

Data ($\mathscr{D}_N$)

$\boldsymbol{\xi}_1, \boldsymbol{\xi}_2, \ldots, \boldsymbol{\xi}_N$ → Loss: $\ell(\boldsymbol{\xi}_1), \ell(\boldsymbol{\xi}_2), \ldots, \ell(\boldsymbol{\xi}_N)$ → Empirical entropic risk: $\rho_{\hat{\mathbb{P}}_N}(\ell(\boldsymbol{\xi})) := \frac{1}{\alpha} \log\left(\frac{1}{N}\sum_{i=1}^N \exp(\alpha \ell(\boldsymbol{\xi}_i))\right)$

Empirical entropic risk underestimates true entropic risk:

✓ Jensen's inequality: $\mathbb{E}$[empirical risk] < True risk

✓ Optimized certainty equivalent (OCE) measure

$$\rho_{\mathbb{P}}(\ell(\boldsymbol{\xi})) = \inf_t \mathbb{E}_{\mathbb{P}}\left(t + \frac{1}{\alpha}\exp(\alpha(\ell(\boldsymbol{\xi}) - t)) - \frac{1}{\alpha}\right)$$

→ replace with $\hat{\mathbb{P}}_N$ (optimizer's curse)

Ex 1: pricing insurance

- Loss $\xi \sim \Gamma(10, 0.24)$
- **Insurer covers the risk**:

$$\text{Premium} = \frac{1}{\alpha} \log\left(\mathbb{E}_{\mathbb{P}}\left[\exp(\alpha \ell(\boldsymbol{\xi}))\right]\right)$$

- Sample mean → true mean slowly:

Gaussian $\xi \implies \exp(\alpha\xi)$ is log-normal

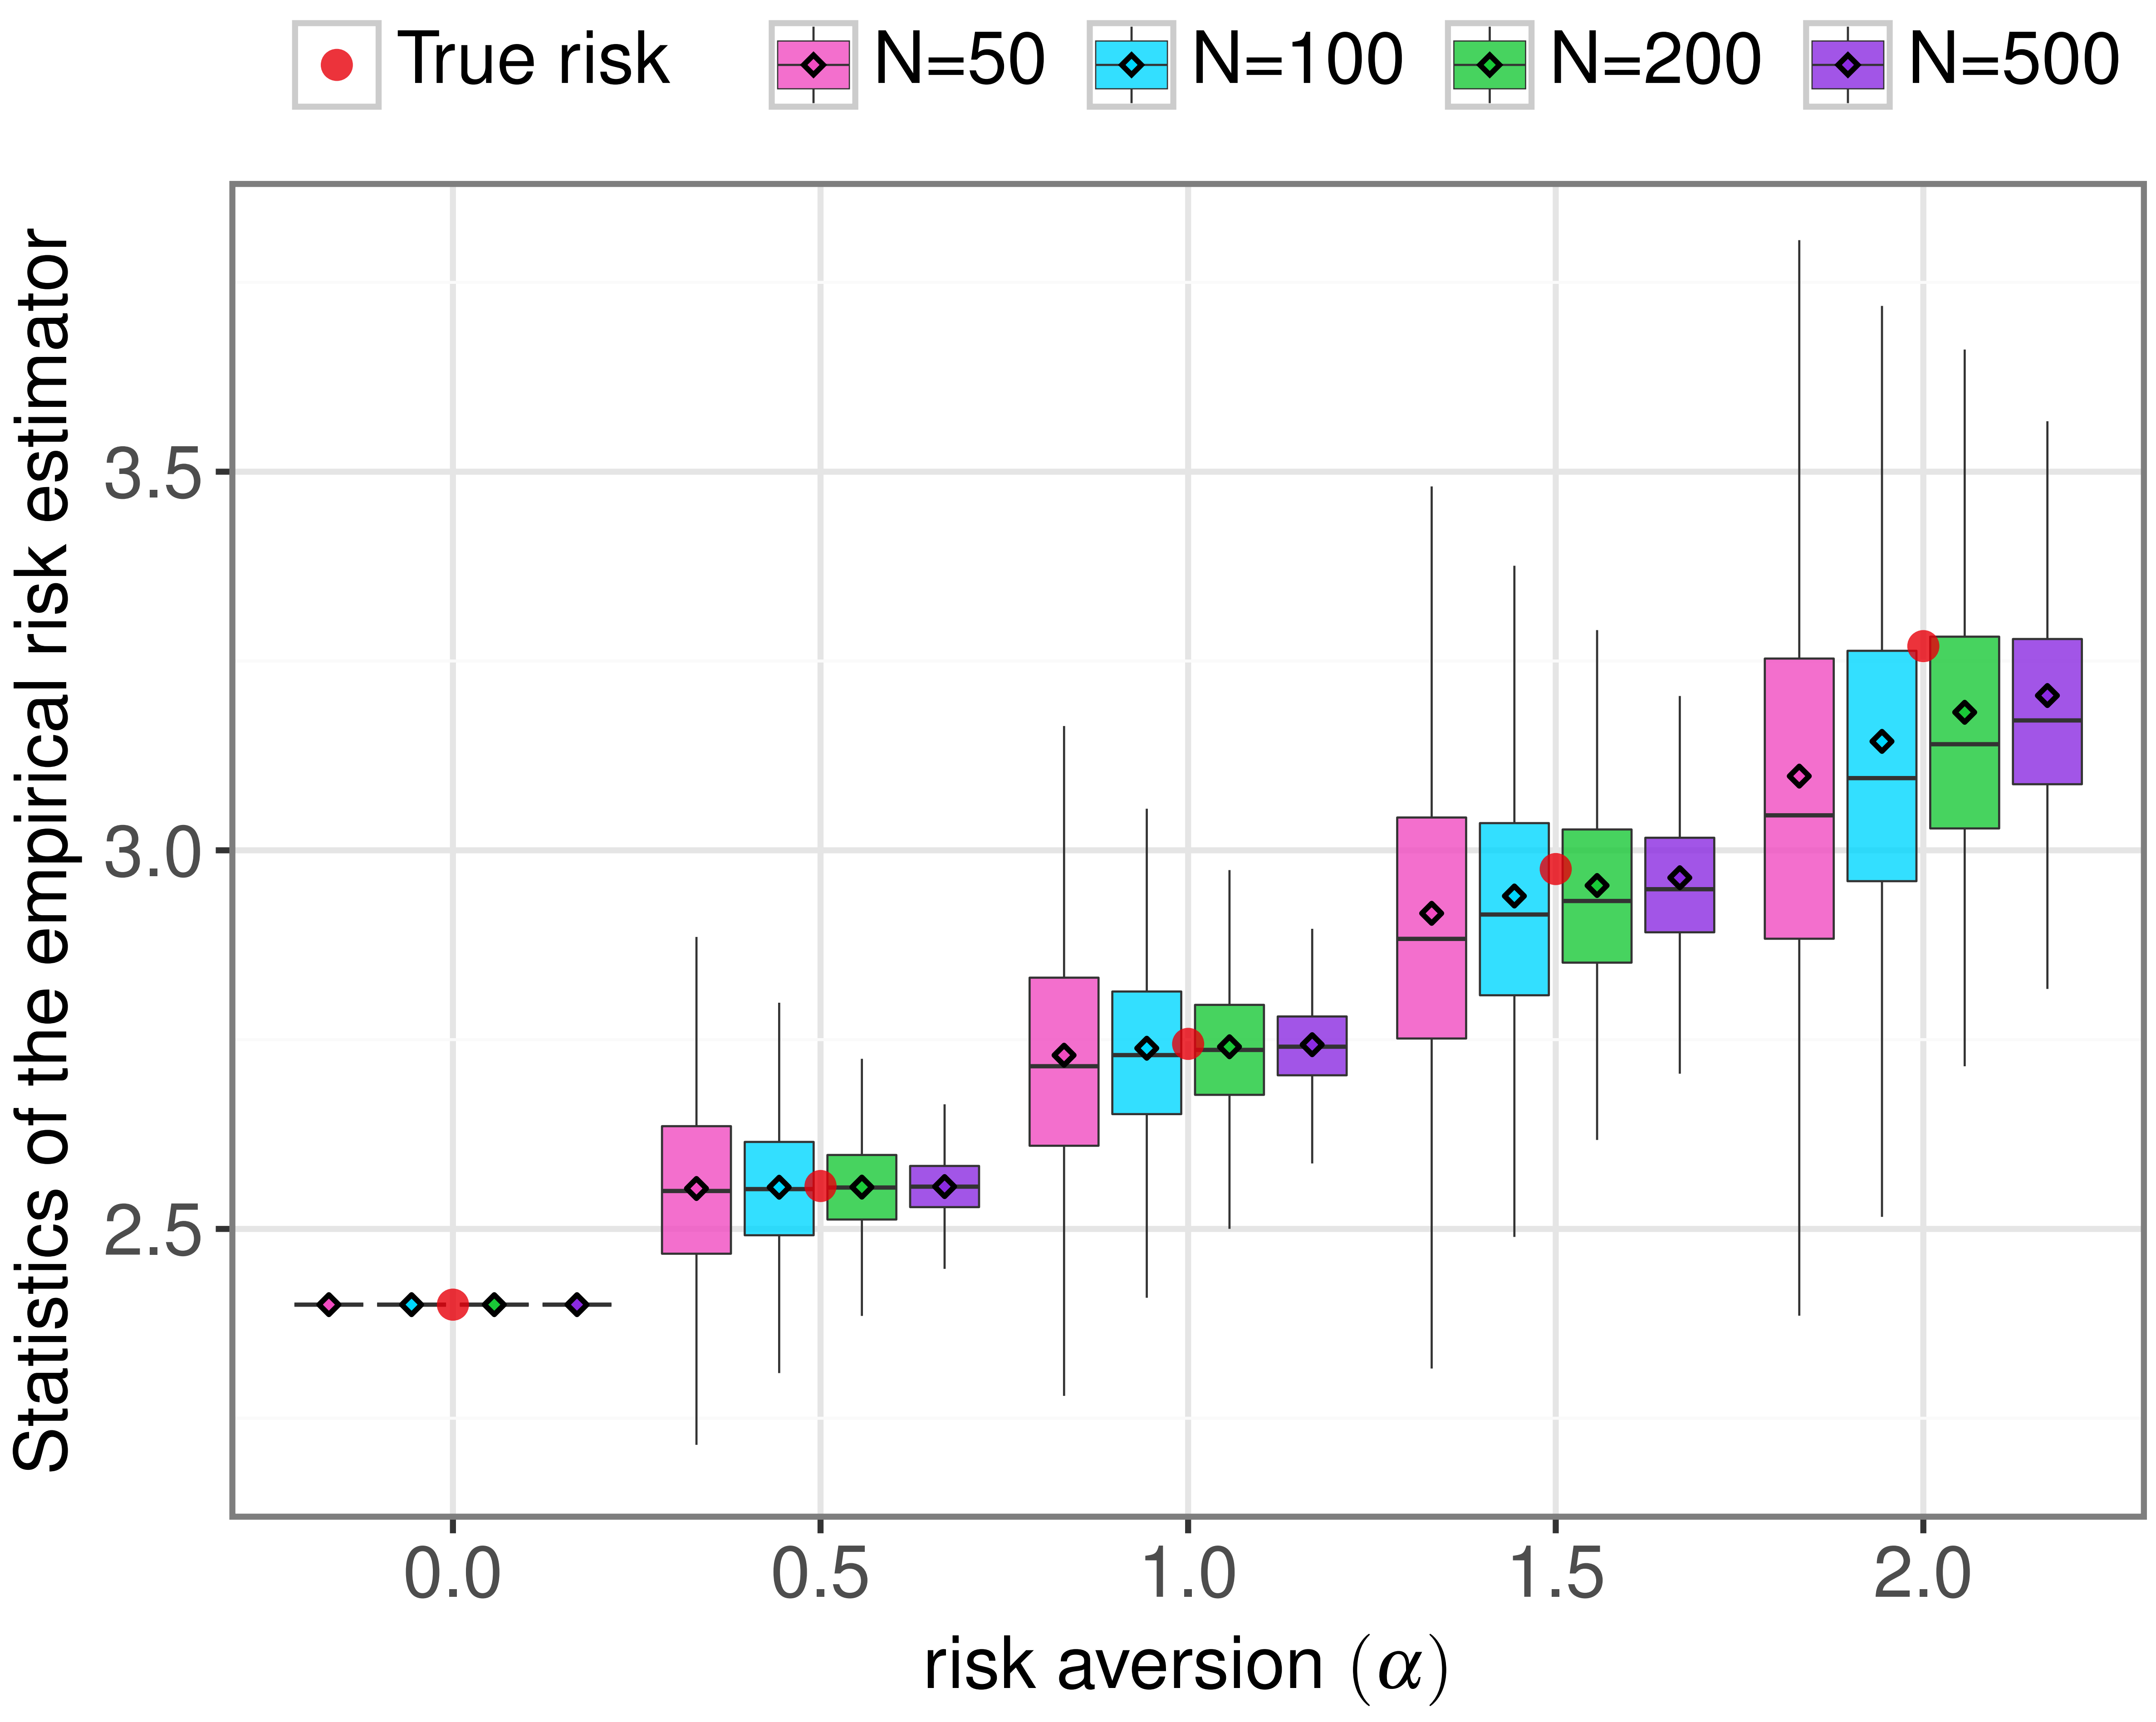

Influence function (IF)

Influence function (IF) - impact of data removal on risk

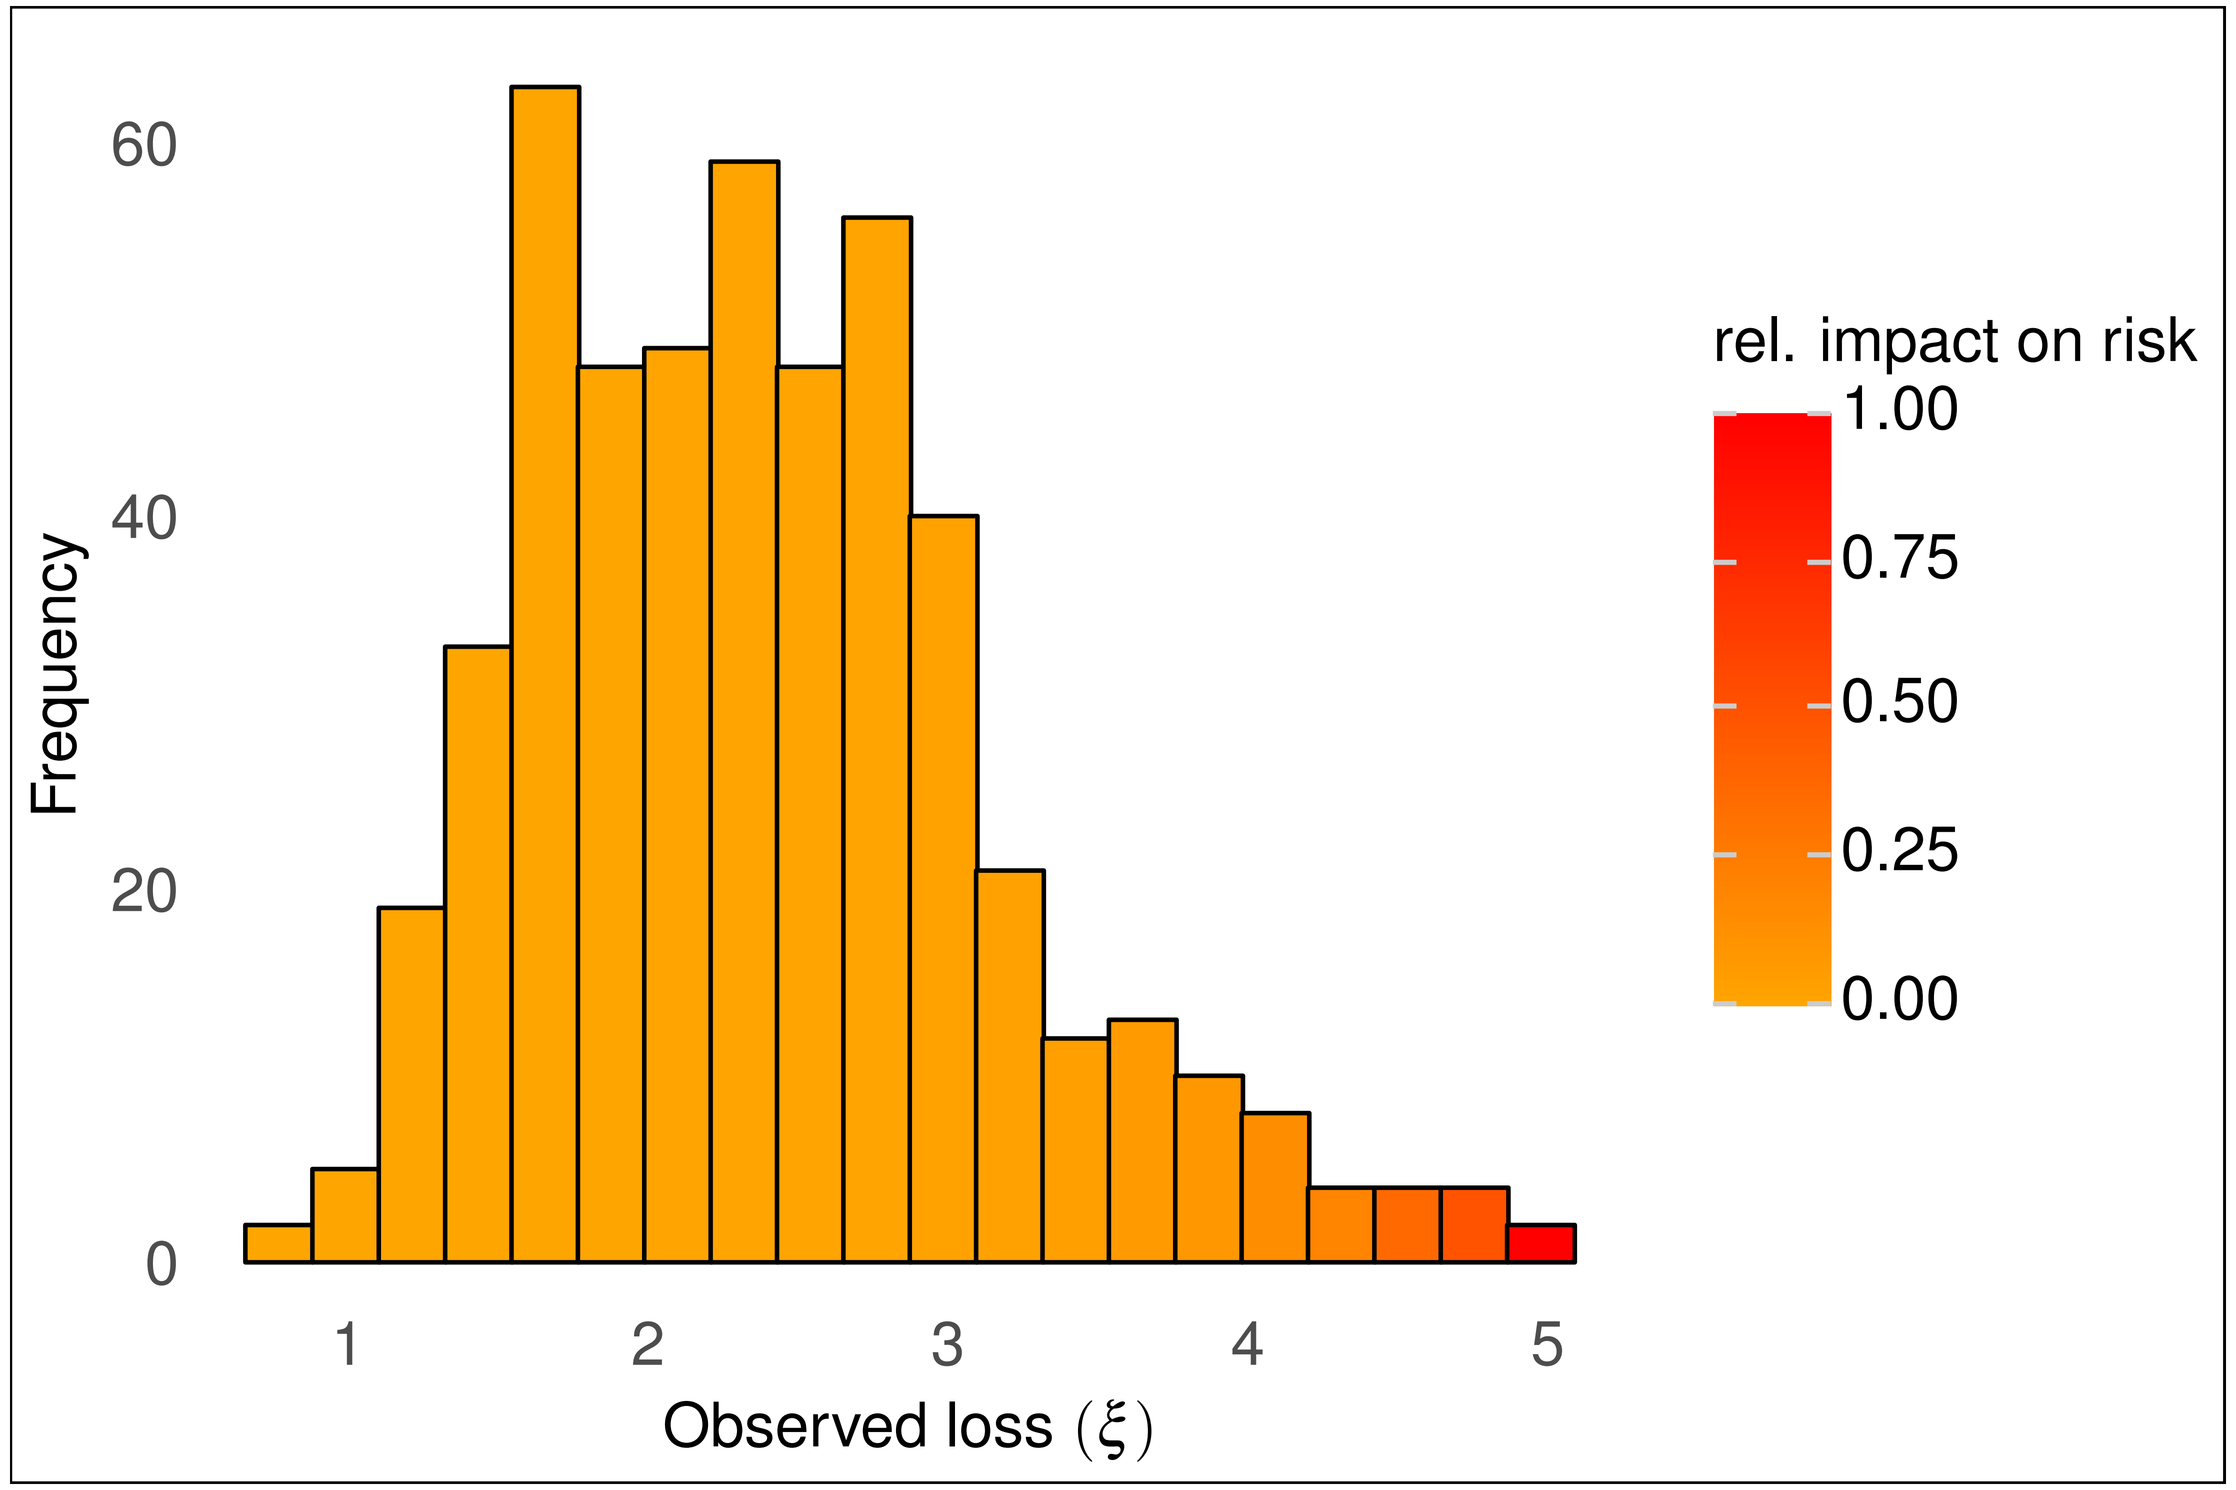

Bias mitigation with bootstrapping

Bias mitigation with bootstrapping

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

Bias mitigation with bootstrapping

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bootstrap

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios
- Draw N samples from $\mathbb{Q}$, compute risk ρ_n and repeat M times

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios
- Draw N samples from $\mathbb{Q}$, compute risk ρ_n and repeat M times
- Bias: $\delta_N(\mathbb{Q}) = \text{median}[\{\rho_{\mathbb{Q}}(\zeta) - \rho_n\}_{i=1}^{M}]$

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios
- Draw N samples from $\mathbb{Q}$, compute risk ρ_n and repeat M times
- Bias: $\delta_N(\mathbb{Q}) = \text{median}[\{\rho_{\mathbb{Q}}(\zeta) - \rho_n\}_{i=1}^{M}]$

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Bias mitigation with bootstrapping

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios
- Draw N samples from $\mathbb{Q}$, compute risk ρ_n and repeat M times
- Bias: $\delta_N(\mathbb{Q}) = \text{median}[\{\rho_{\mathbb{Q}}(\zeta) - \rho_n\}_{i=1}^{M}]$

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Theorem: Under some assumptions on tails of ζ:

$\rho_{\hat{\mathbb{P}}_N}(\zeta) + \delta_N(\mathbb{Q})$ **almost surely converges to true entropic risk**

Bias mitigation with bootstrapping

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios
- Draw N samples from $\mathbb{Q}$, compute risk ρ_n and repeat M times
- Bias: $\delta_N(\mathbb{Q}) = \text{median}[\{\rho_{\mathbb{Q}}(\zeta) - \rho_n\}_{i=1}^{M}]$

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Theorem: Under some assumptions on tails of ζ:

$\rho_{\hat{\mathbb{P}}_N}(\zeta) + \delta_N(\mathbb{Q})$ **almost surely converges to true entropic risk**

Bias mitigation with bootstrapping

Efficiently computable risk under $\mathbb{Q}$
Gaussian mixture models are universal function approximators

$$\rho_{\mathbb{Q}}(\zeta) = (1/\alpha)\log\left(\sum_y \pi_y \exp(\alpha\mu_y + \alpha^2\sigma_y^2/2)\right)$$

Bootstrap

- Fit a distribution $\mathbb{Q}$ to the loss scenarios
- Draw N samples from $\mathbb{Q}$, compute risk ρ_n and repeat M times
- Bias: $\delta_N(\mathbb{Q}) = \text{median}[\{\rho_{\mathbb{Q}}(\zeta) - \rho_n\}_{i=1}^{M}]$

Bias: $\delta = \mathbb{E}[\rho_{\mathbb{P}}(\zeta) - \rho_{\hat{\mathbb{P}}_N}(\zeta)]$

$\mathbb{P}$ is unknown

Theorem: Under some assumptions on tails of ζ:

$\rho_{\hat{\mathbb{P}}_N}(\zeta) + \delta_N(\mathbb{Q})$ **almost surely converges to true entropic risk**

Model 1: Fit using maximum likelihood (BS-MLE)

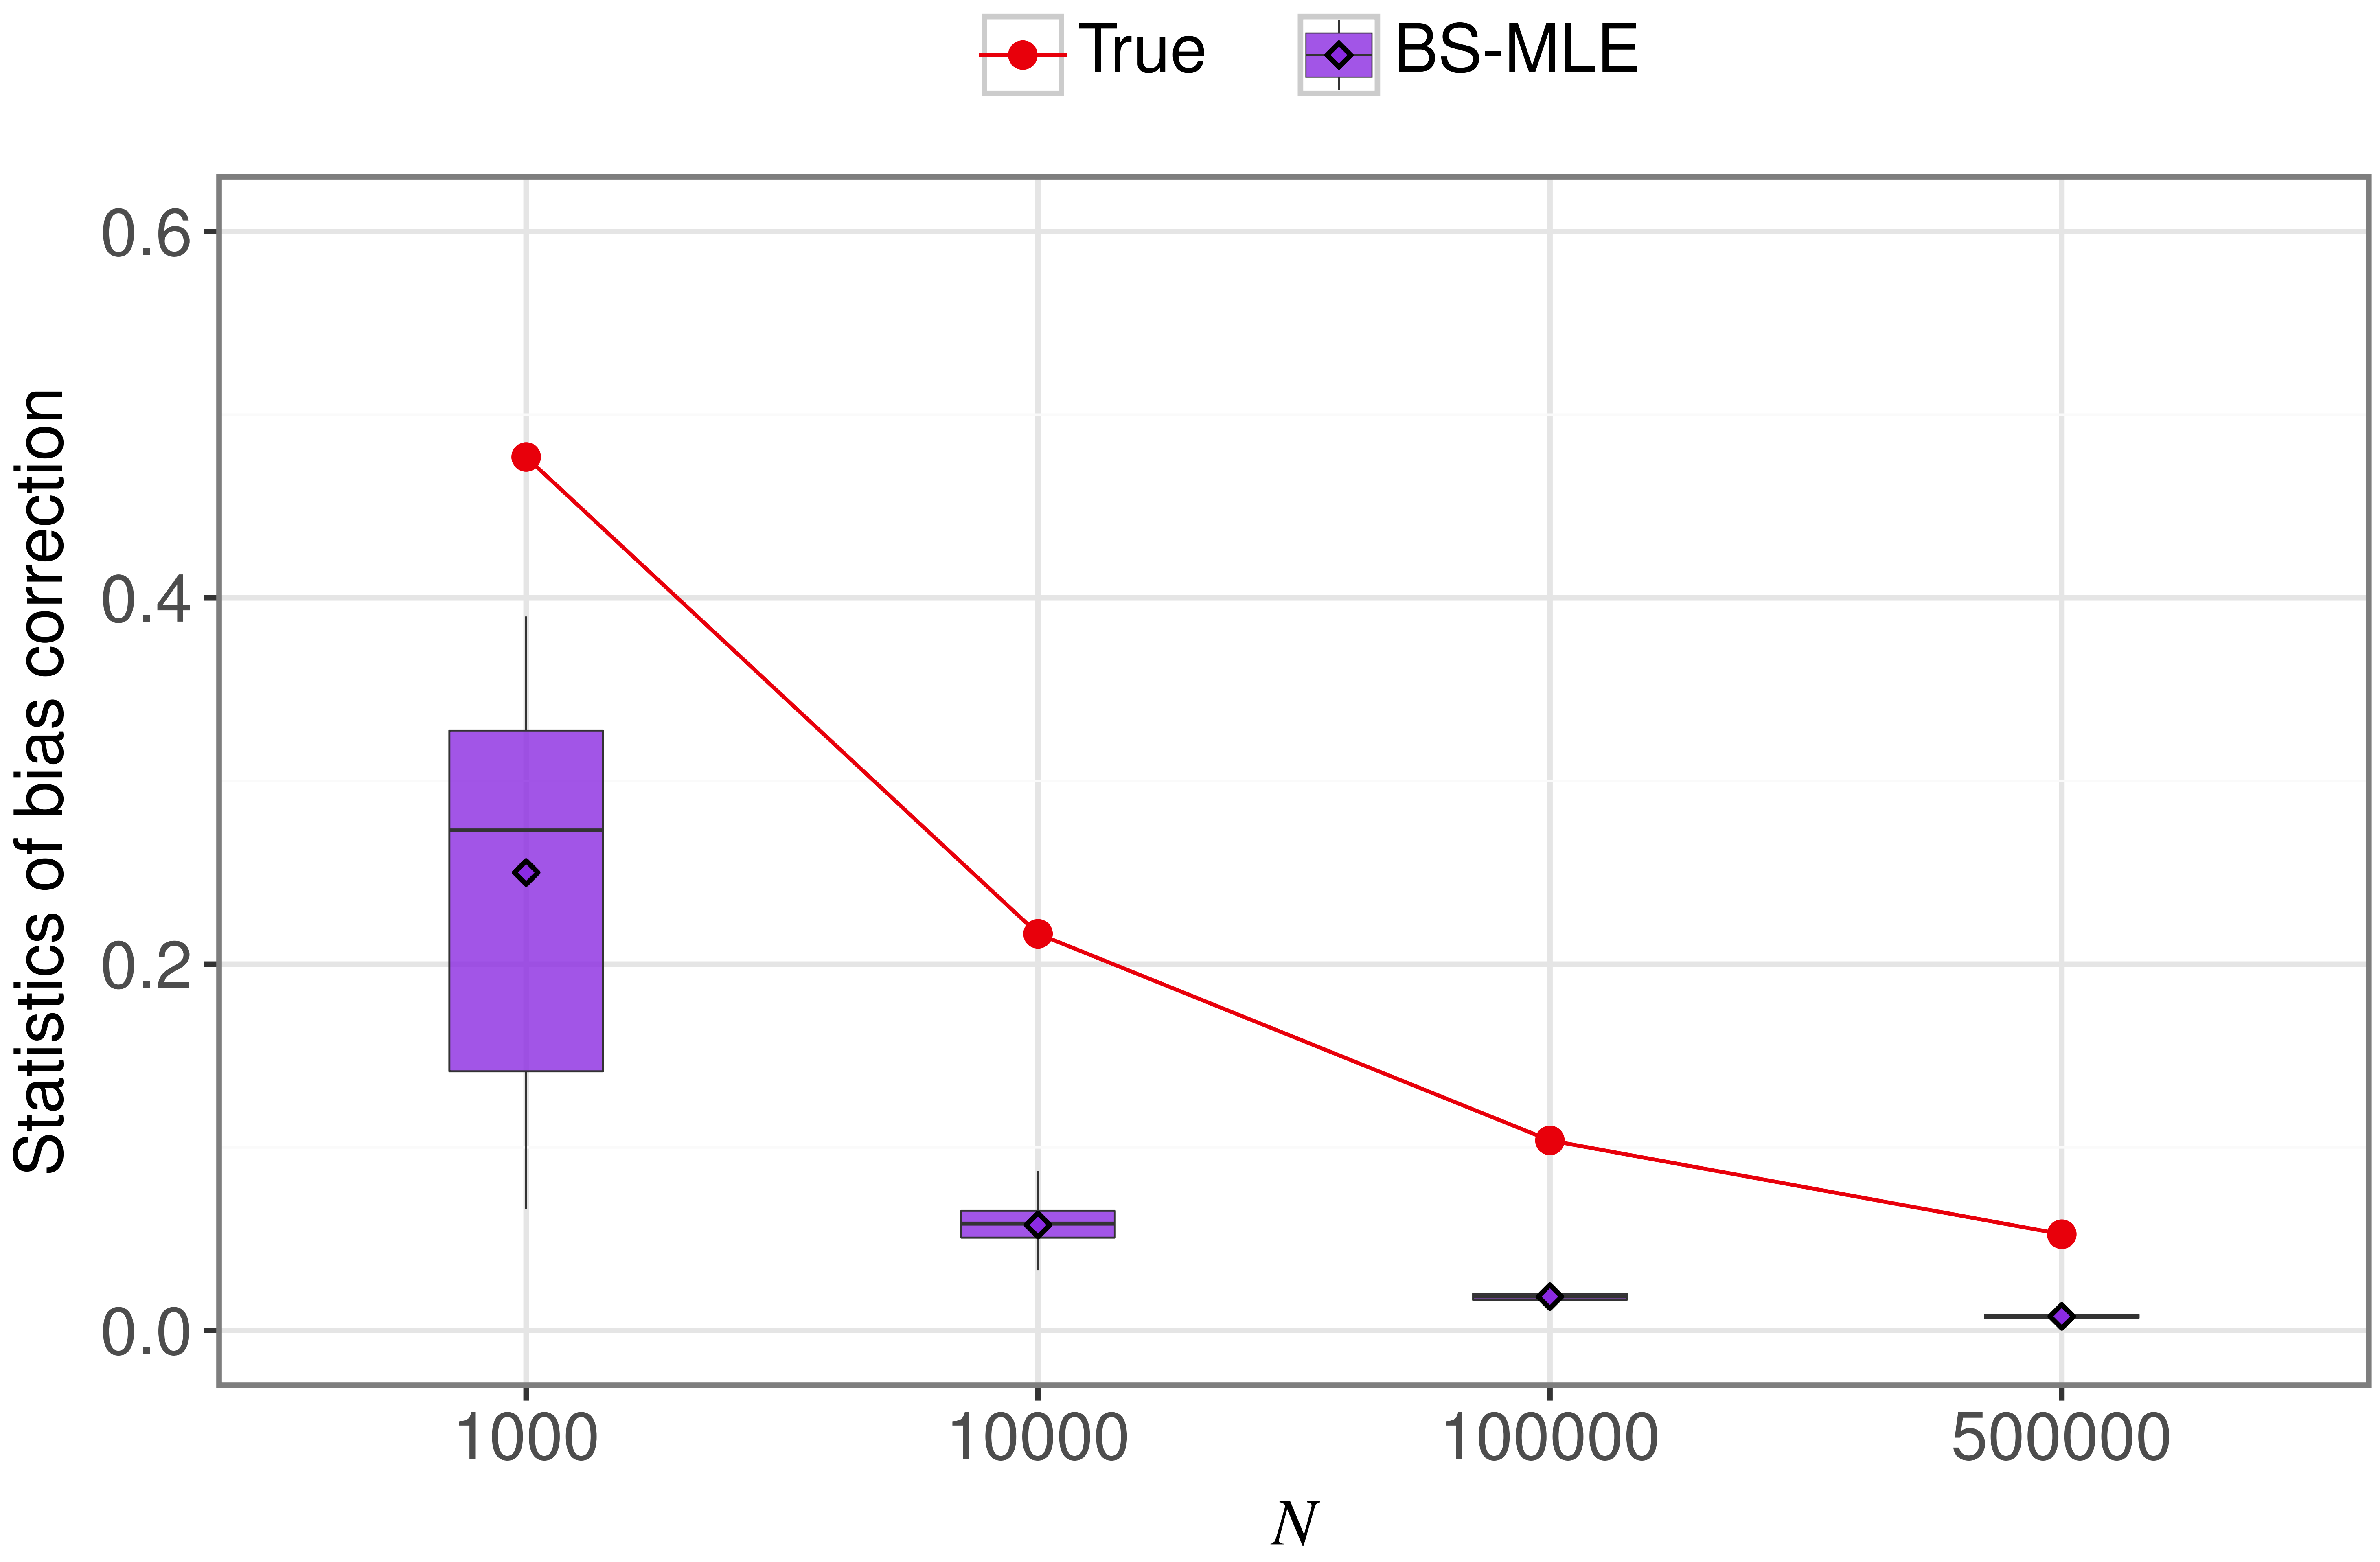

- Ex: Compute entropic risk
- $\boldsymbol{\xi} \sim$ GMM$(\boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\sigma})$,
  $\boldsymbol{\pi} = [0.7\ 0.3], \boldsymbol{\mu} = [0.5\ 1]$,
  $\boldsymbol{\sigma} = [2\ 1]$
- **BS-MLE - Fit $\mathbb{Q}$ using MLE**
- **Underestimation persists**

Model 1: Fit using maximum likelihood (BS-MLE)

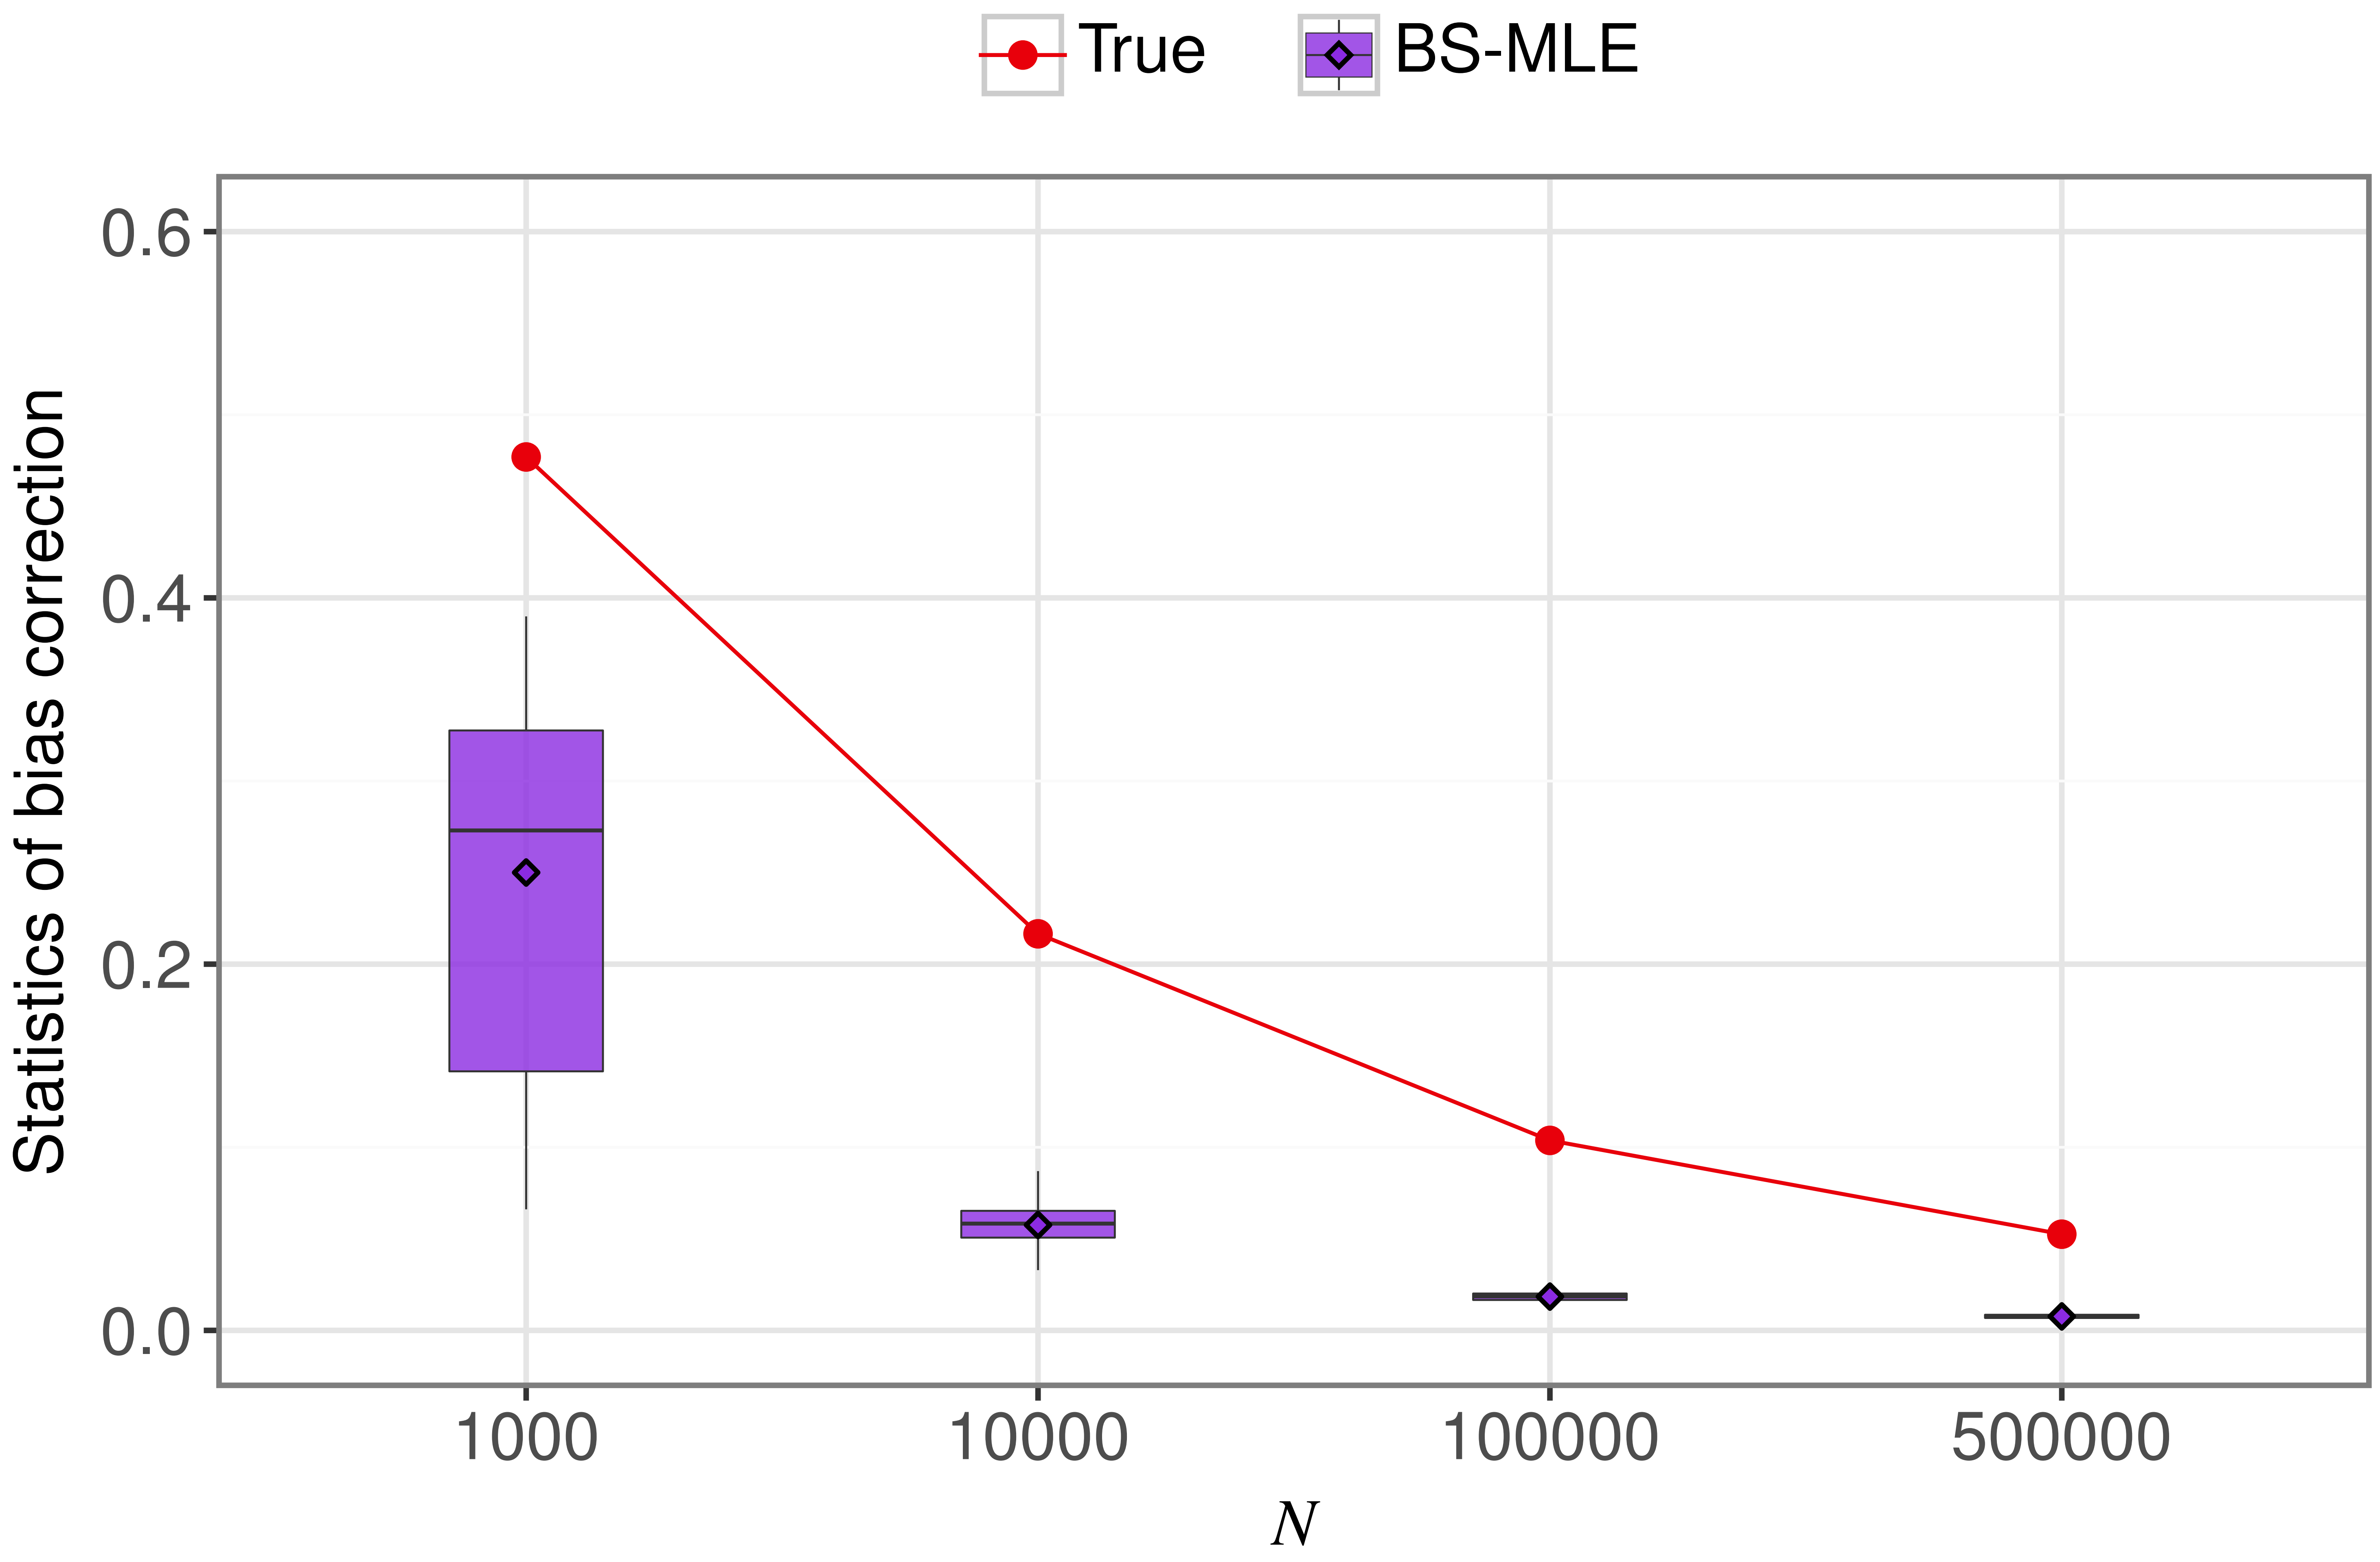

- Ex: Compute entropic risk
- $\boldsymbol{\xi} \sim \text{GMM}(\boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\sigma})$, $\boldsymbol{\pi} = [0.7\ 0.3], \boldsymbol{\mu} = [0.5\ \ 1]$, $\boldsymbol{\sigma} = [2\ 1]$
- **BS-MLE - Fit $\mathbb{Q}$ using MLE**
- **Underestimation persists**

Bias mitigation using Bias-aware bootstrapping

Wish:

Fit distribution $\mathbb{Q}$ whose samples have the same bias as the bias in the data

Bias mitigation using Bias-aware bootstrapping

Model 2: Entropic risk matching (BS-Match)

Idea: Match distributions of the entropic risk over the samples

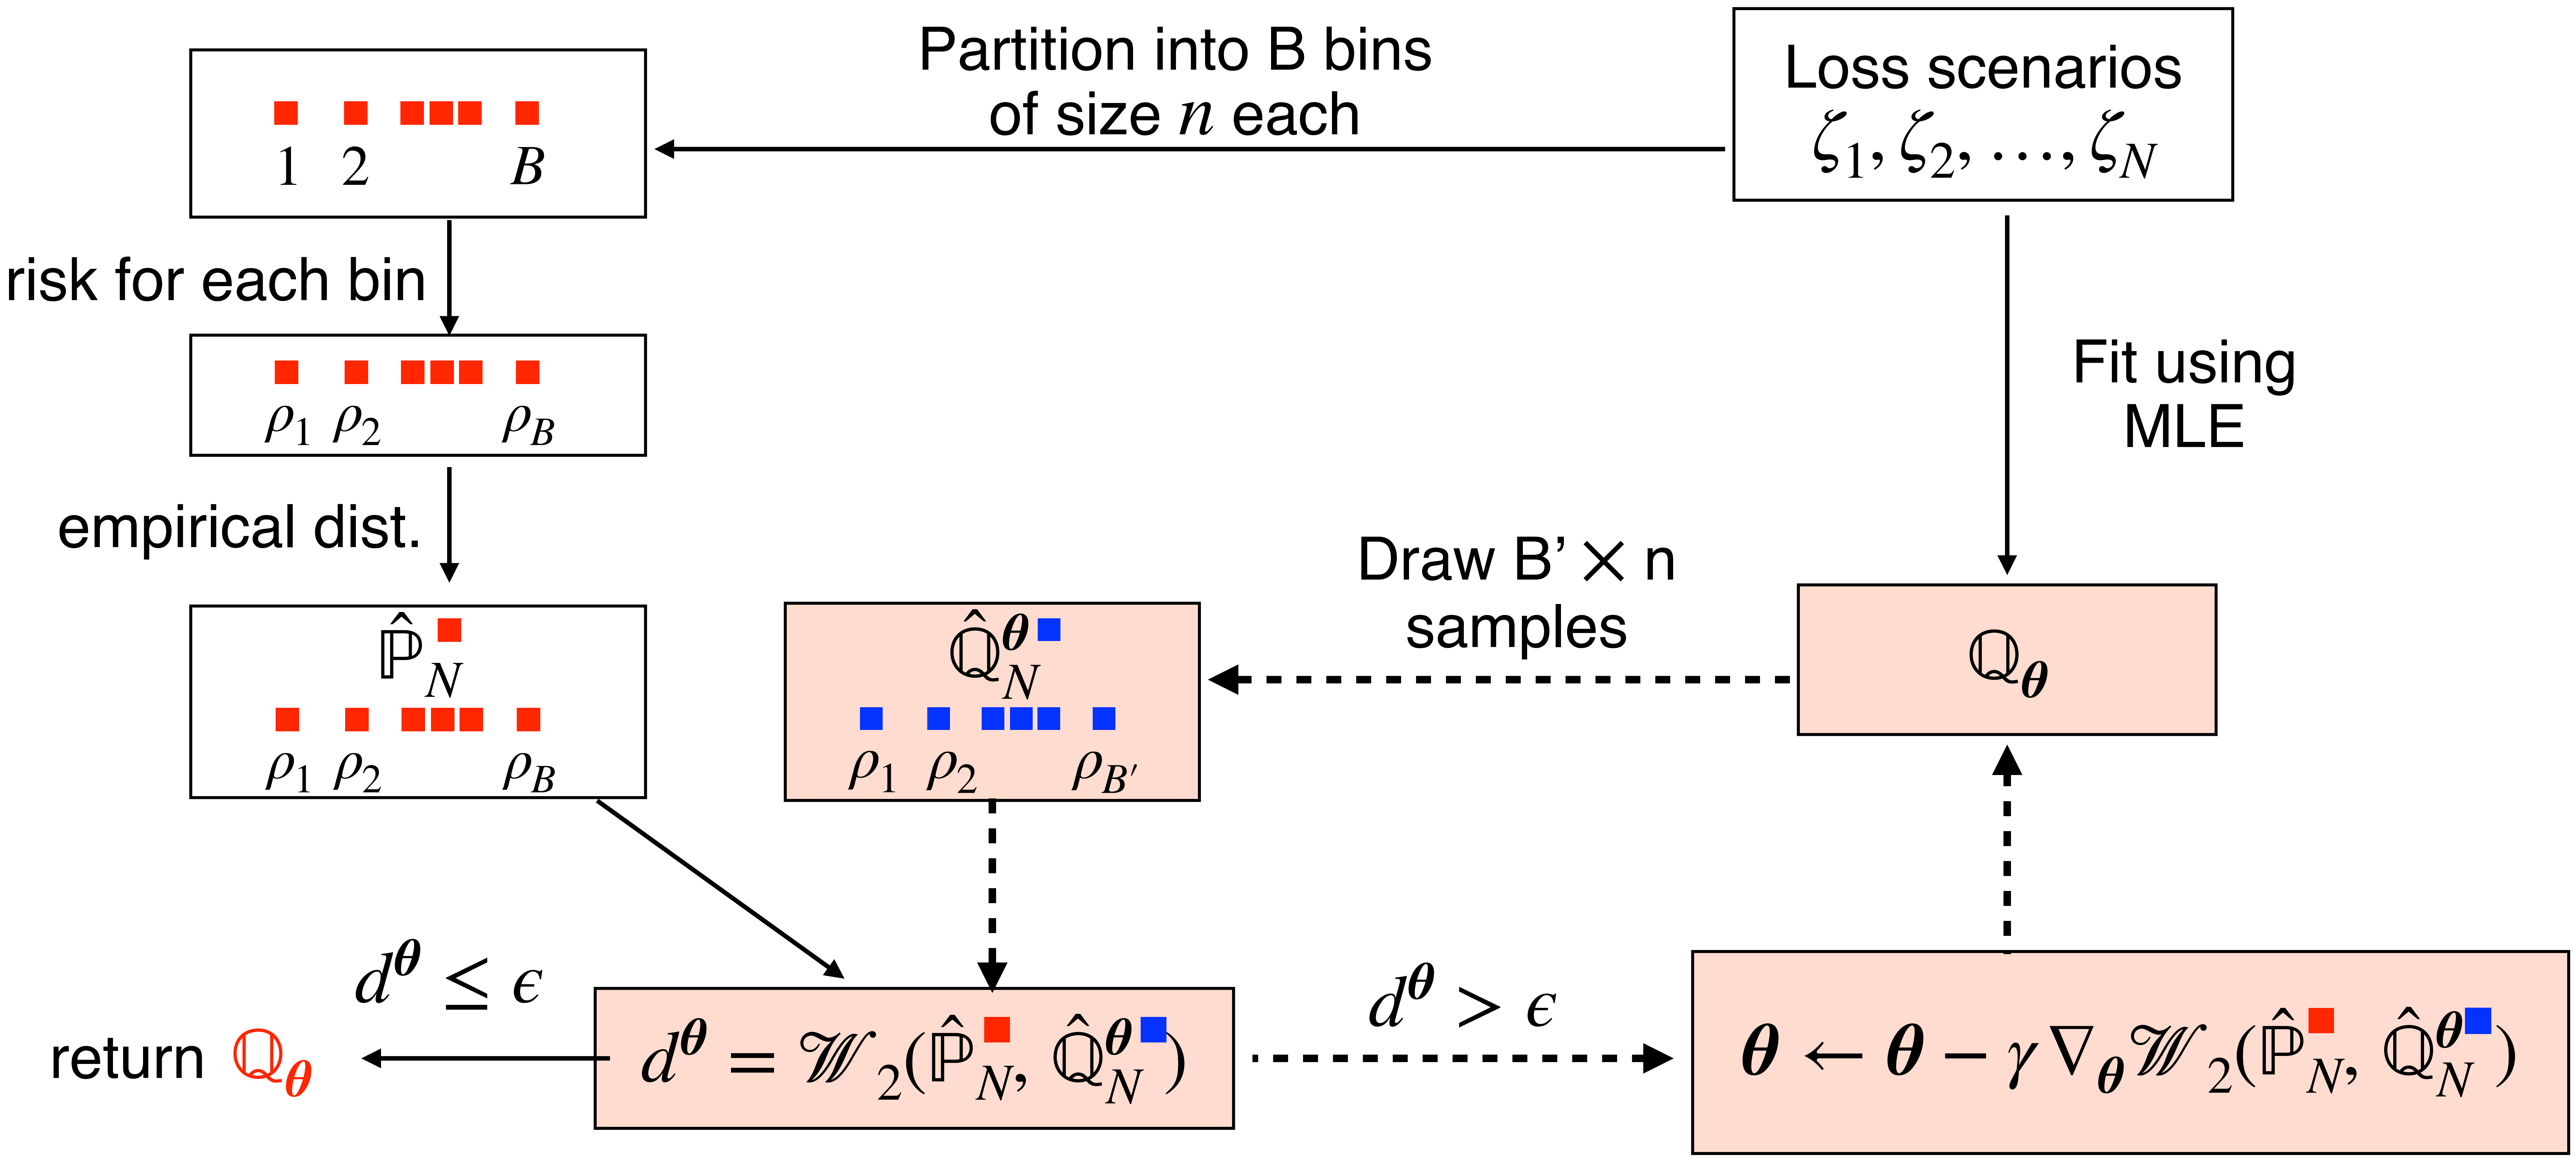

Model 3: Extreme value theory (BS-Match)

Model 3: Extreme value theory (BS-Match)

- Loss scenarios $\zeta_1, \zeta_2, \cdots, \zeta_n$ iid

Model 3: Extreme value theory (BS-Match)

- Loss scenarios $\zeta_1, \zeta_2, \cdots, \zeta_n$ iid
- $M_n = \max\{\zeta_1, \zeta_2, \cdots, \zeta_n\}$

Model 3: Extreme value theory (BS-Match)

- Loss scenarios $\zeta_1, \zeta_2, \cdots, \zeta_n$ iid
- $M_n = \max\{\zeta_1, \zeta_2, \cdots, \zeta_n\}$

Model 3: Extreme value theory (BS-Match)

- Loss scenarios $\zeta_1, \zeta_2, \cdots, \zeta_n$ iid
- $M_n = \max\{\zeta_1, \zeta_2, \cdots, \zeta_n\}$

Fisher–Tippett–Gnedenko theorem:
As $n \to \infty$, distribution of M_n converges to either Weibull, Fréchet or Gumbel

-Fit using MLE

Model 3: Extreme value theory (BS-Match)

- Loss scenarios $\zeta_1, \zeta_2, \cdots, \zeta_n$ iid
- $M_n = \max\{\zeta_1, \zeta_2, \cdots, \zeta_n\}$

Our approach:

- cdf normal rv - $\Phi(\mu, \sigma)$
- Fit $\Phi^n(\mu, \sigma)$ to $m_1, m_2, \cdots, m_B$

Fisher–Tippett–Gnedenko theorem:
As $n \to \infty$, distribution of M_n converges to either Weibull, Fréchet or Gumbel

-Fit using MLE

Model 3: Extreme value theory (BS-Match)

- Loss scenarios $\zeta_1, \zeta_2, \cdots, \zeta_n$ iid
- $M_n = \max\{\zeta_1, \zeta_2, \cdots, \zeta_n\}$

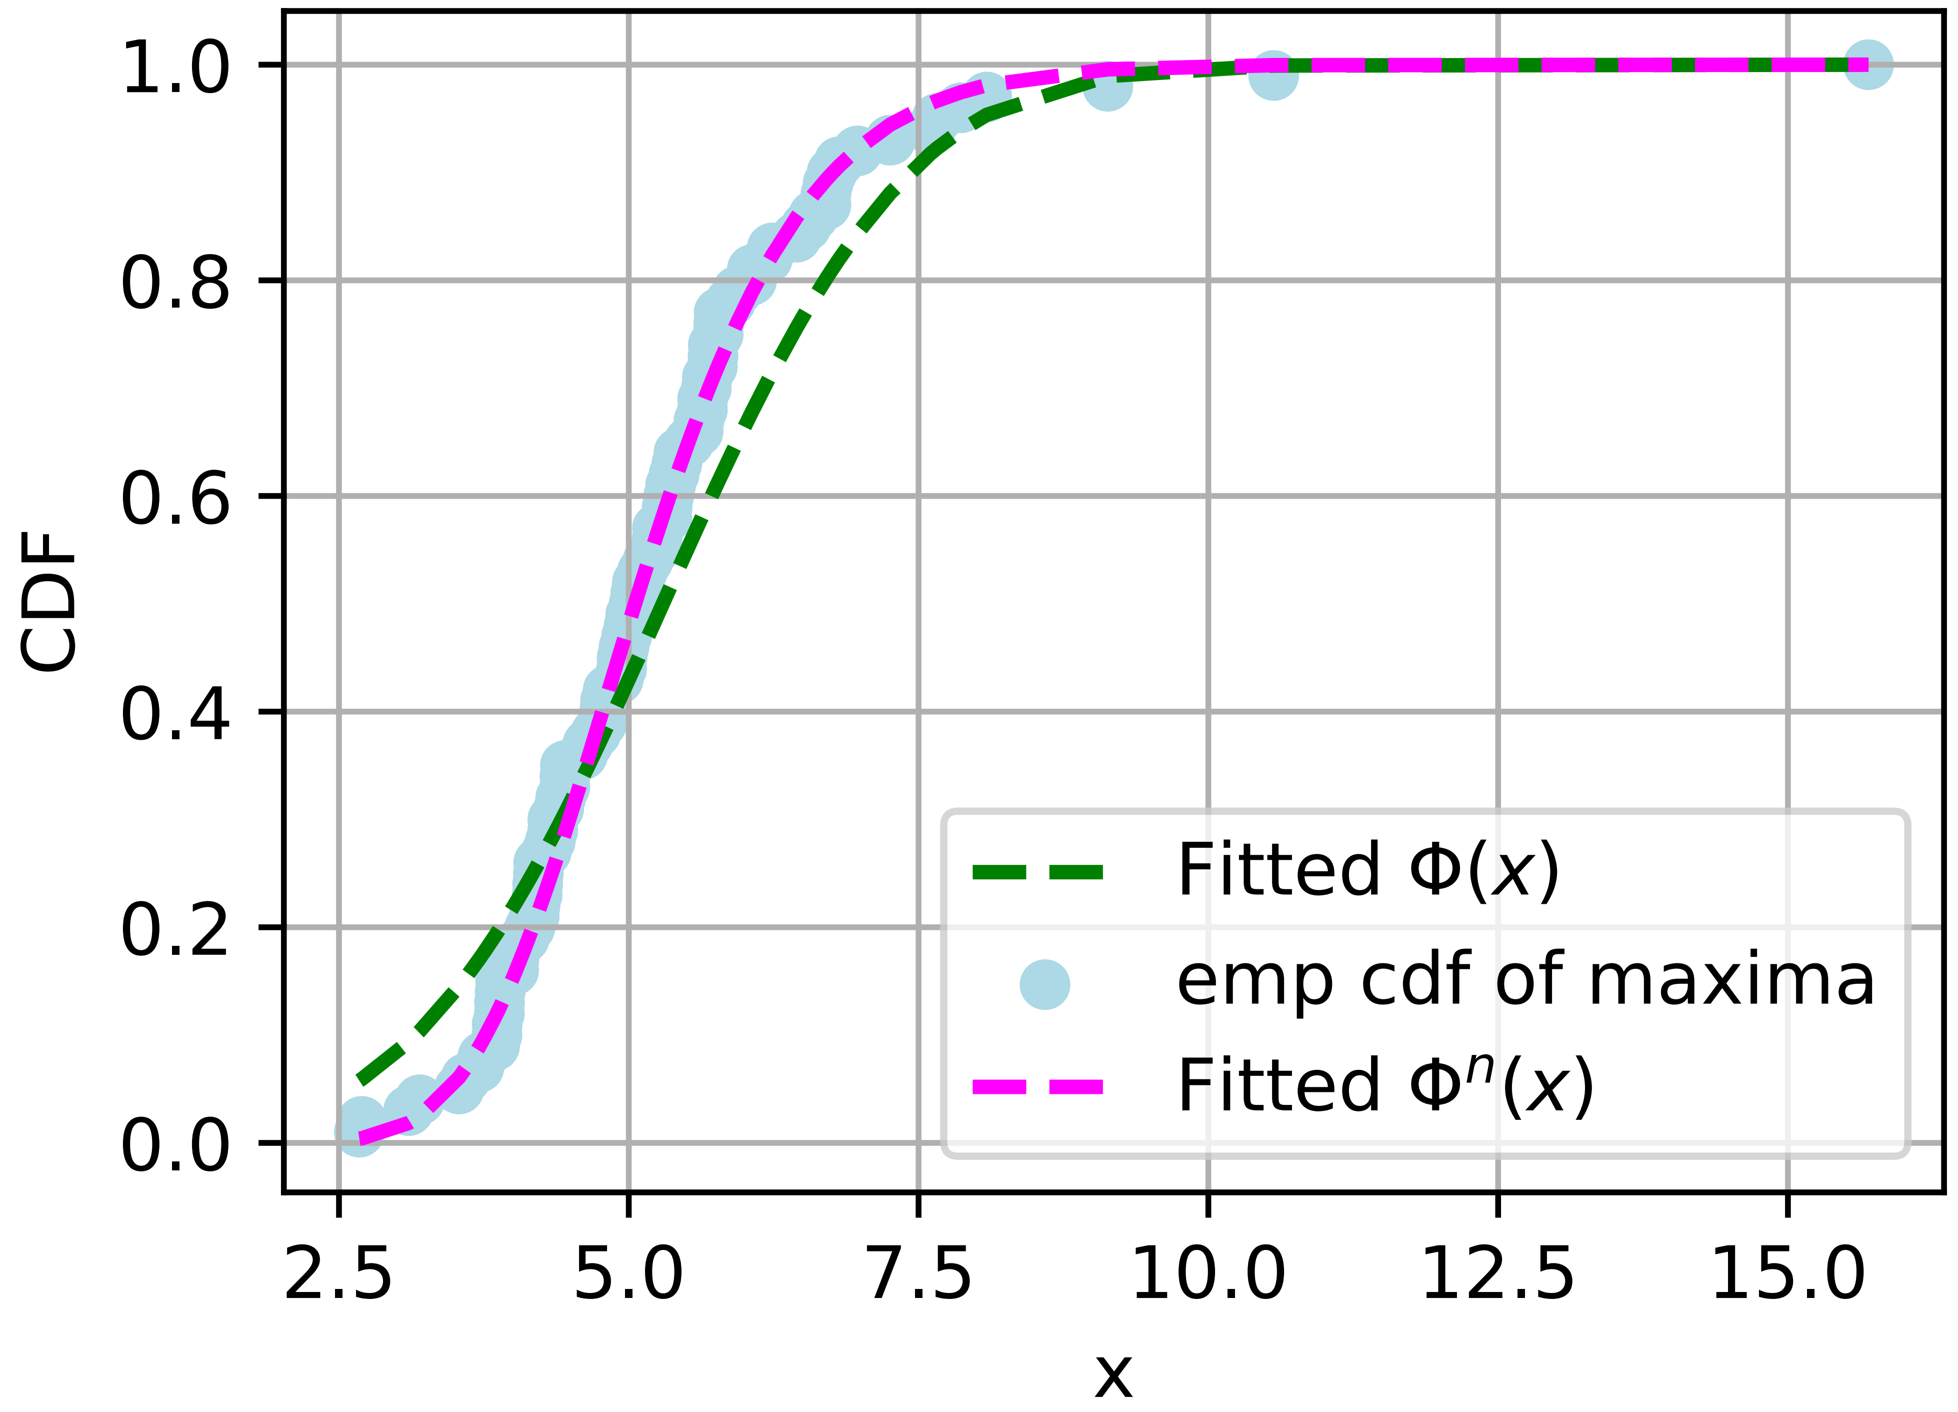

Our approach:

- cdf normal rv - $\Phi(\mu, \sigma)$
- Fit $\Phi^n(\mu, \sigma)$ to $m_1, m_2, \cdots, m_B$

Fisher–Tippett–Gnedenko theorem:
As $n \to \infty$, distribution of M_n converges to either Weibull, Fréchet or Gumbel

-Fit using MLE

Model 3: Extreme value theory (BS-Match)

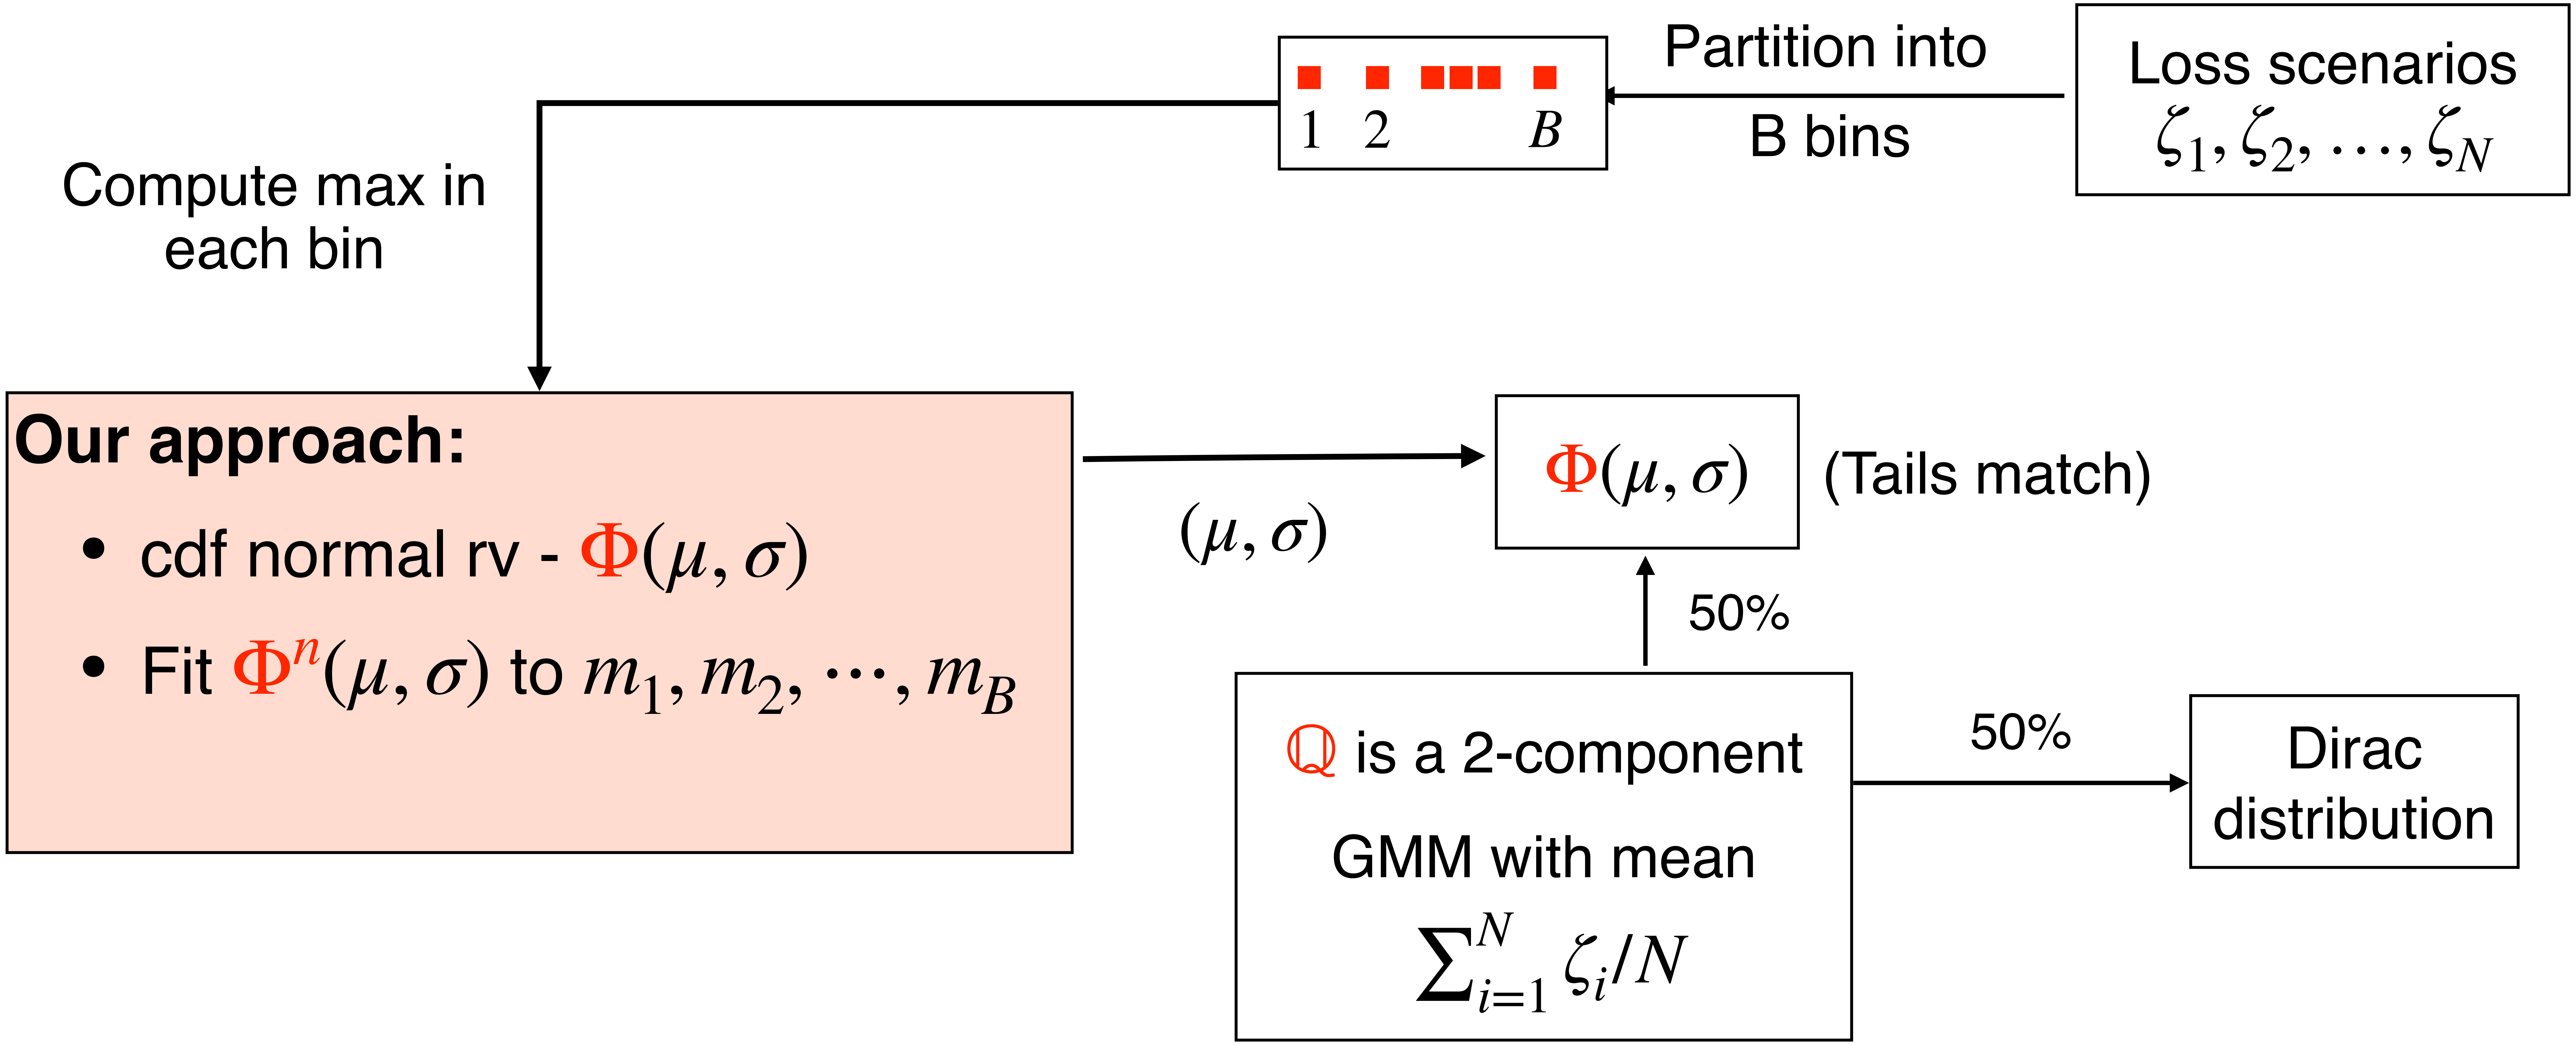

Ex 2: Bias mitigation

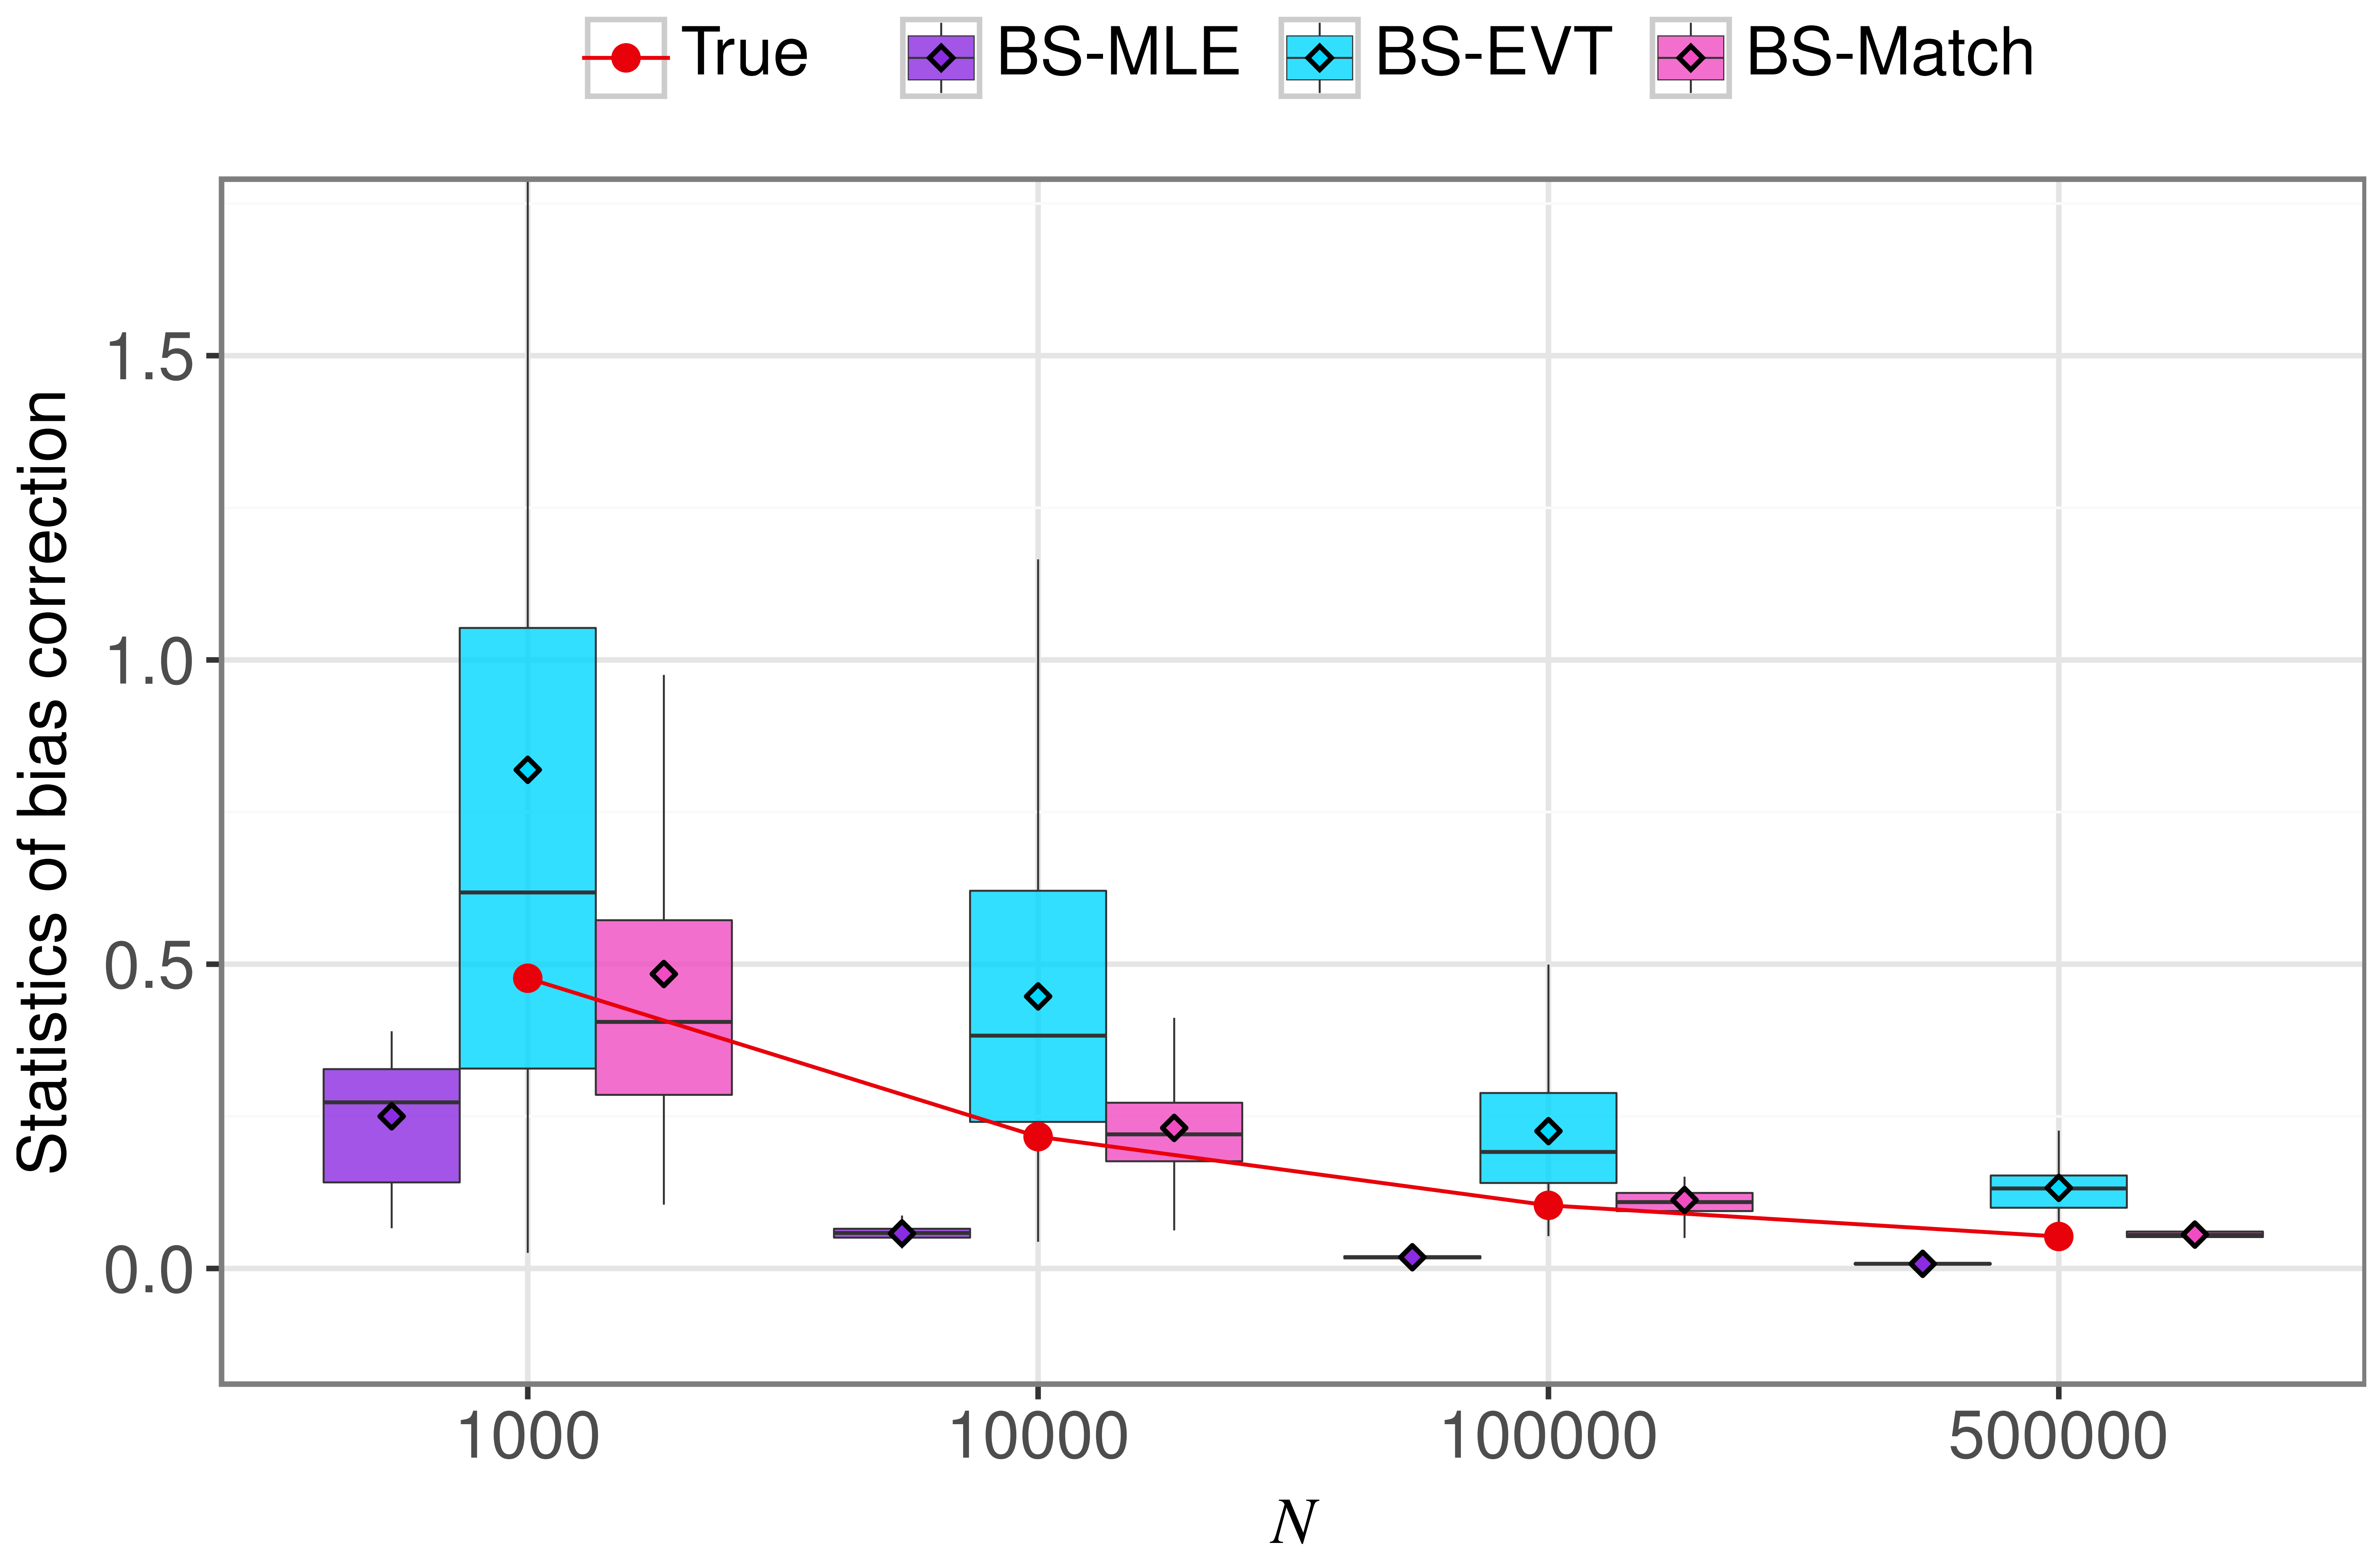

- Ex: Compute entropic risk
- $\boldsymbol{\xi} \sim \text{GMM}(\boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\sigma})$, $\boldsymbol{\pi} = [0.7\ 0.3]$, $\boldsymbol{\mu} = [0.5\ 1]$, $\boldsymbol{\sigma} = [2\ 1]$
- BS-MLE - Fit $\mathbb{Q}$ using MLE
- Underestimation persists

- **BS-EVT - Fit $\mathbb{Q}$ by matching tails**
- **BS-Match - Fit $\mathbb{Q}$ by entropic risk matching**

Ex3: Compare with estimators from literature

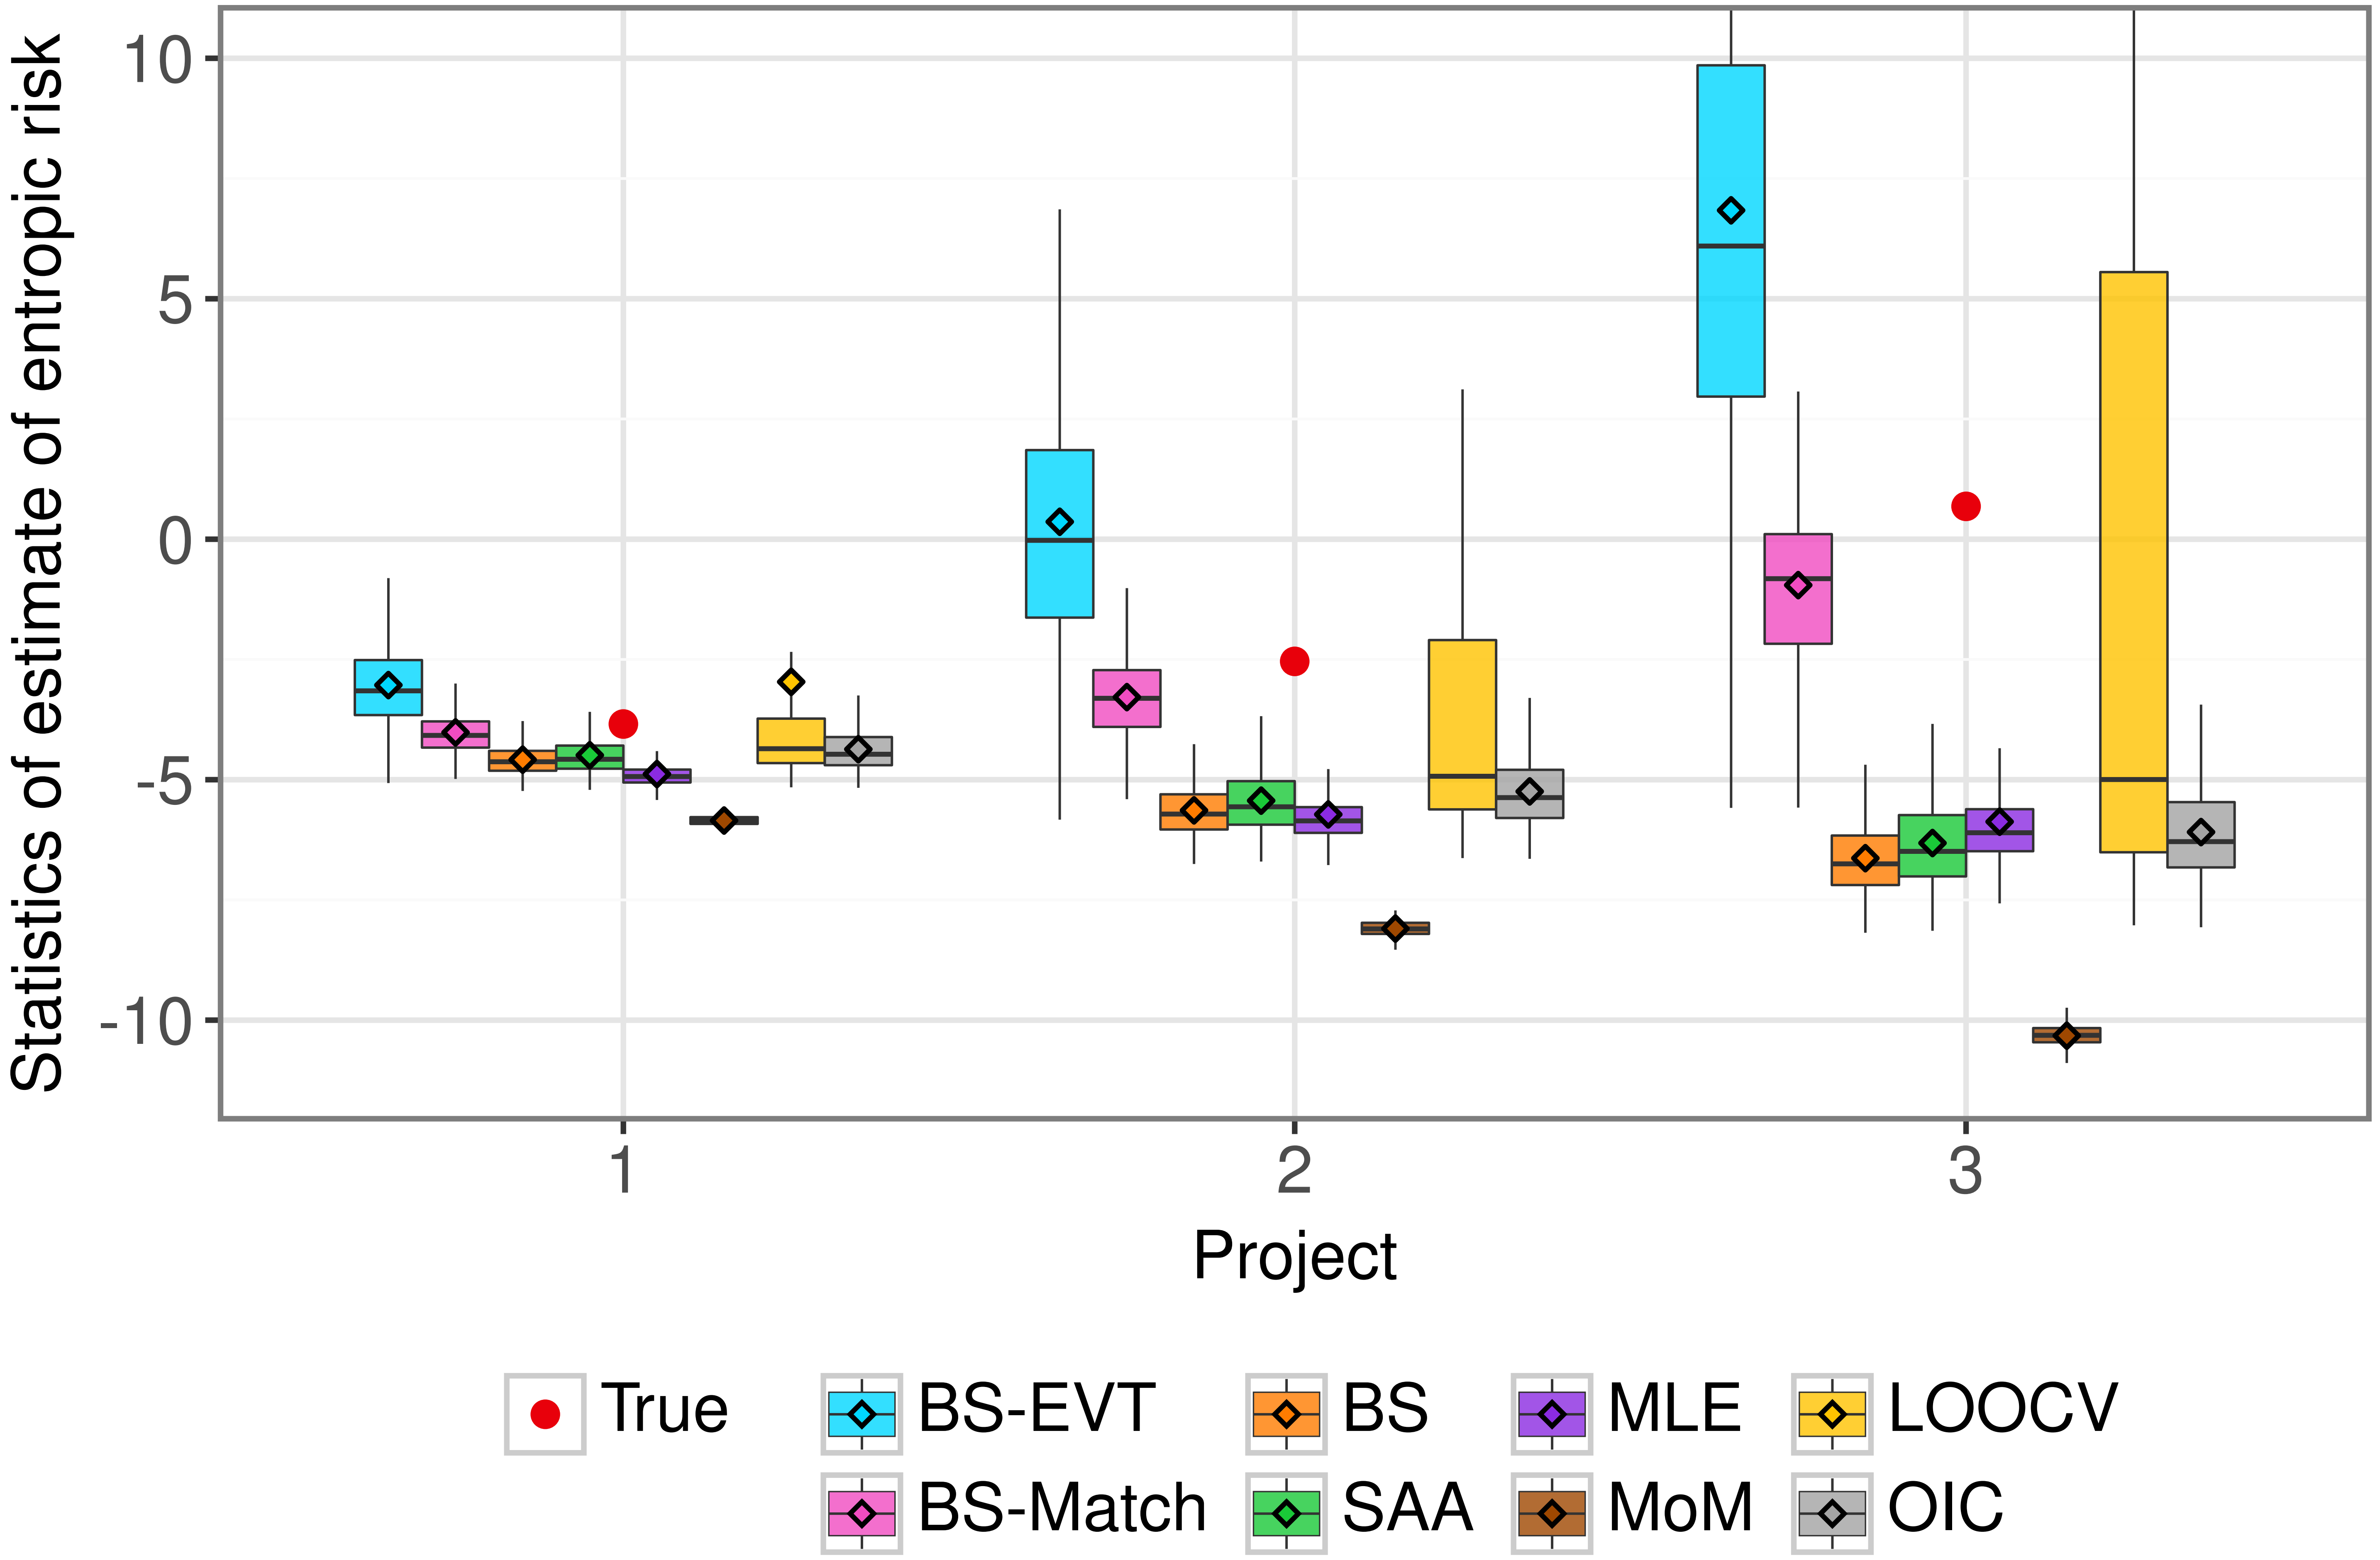

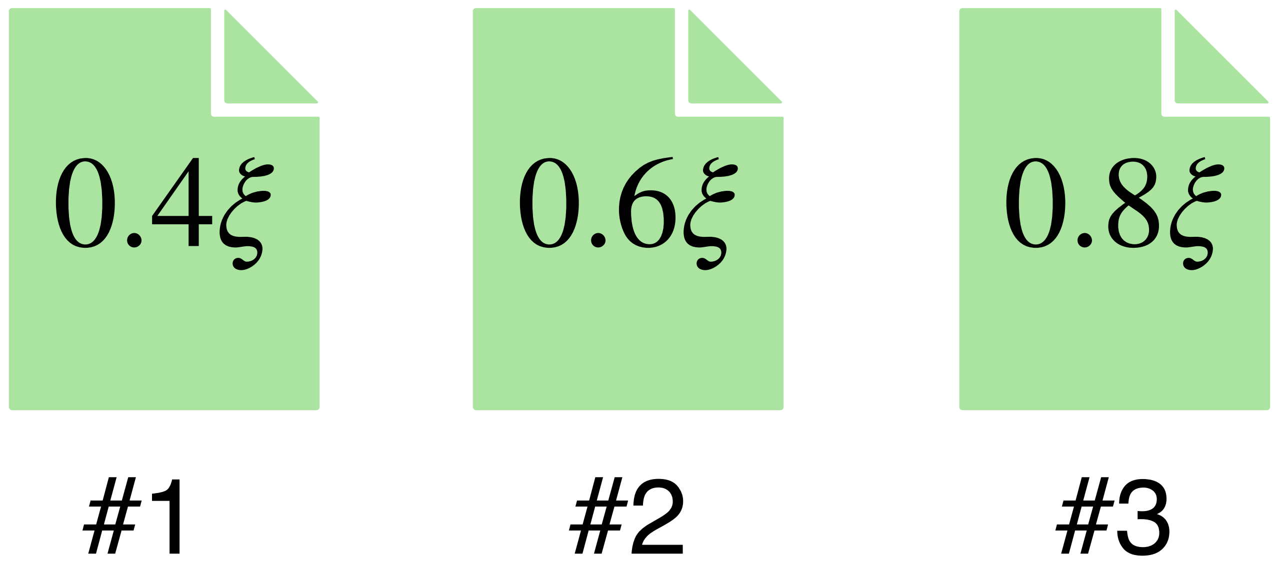

- $\xi \sim$ GMM$(\boldsymbol{\pi}, \boldsymbol{\mu}, \Sigma)$ with 5 components
- across components - $\mu_\xi = -18.6$ $\sigma_\xi = 2.9$
- Which project has lowest entropic risk based on 100 sets of 10000 samples with $\alpha = 3$?

Ex3: Compare with estimators from literature

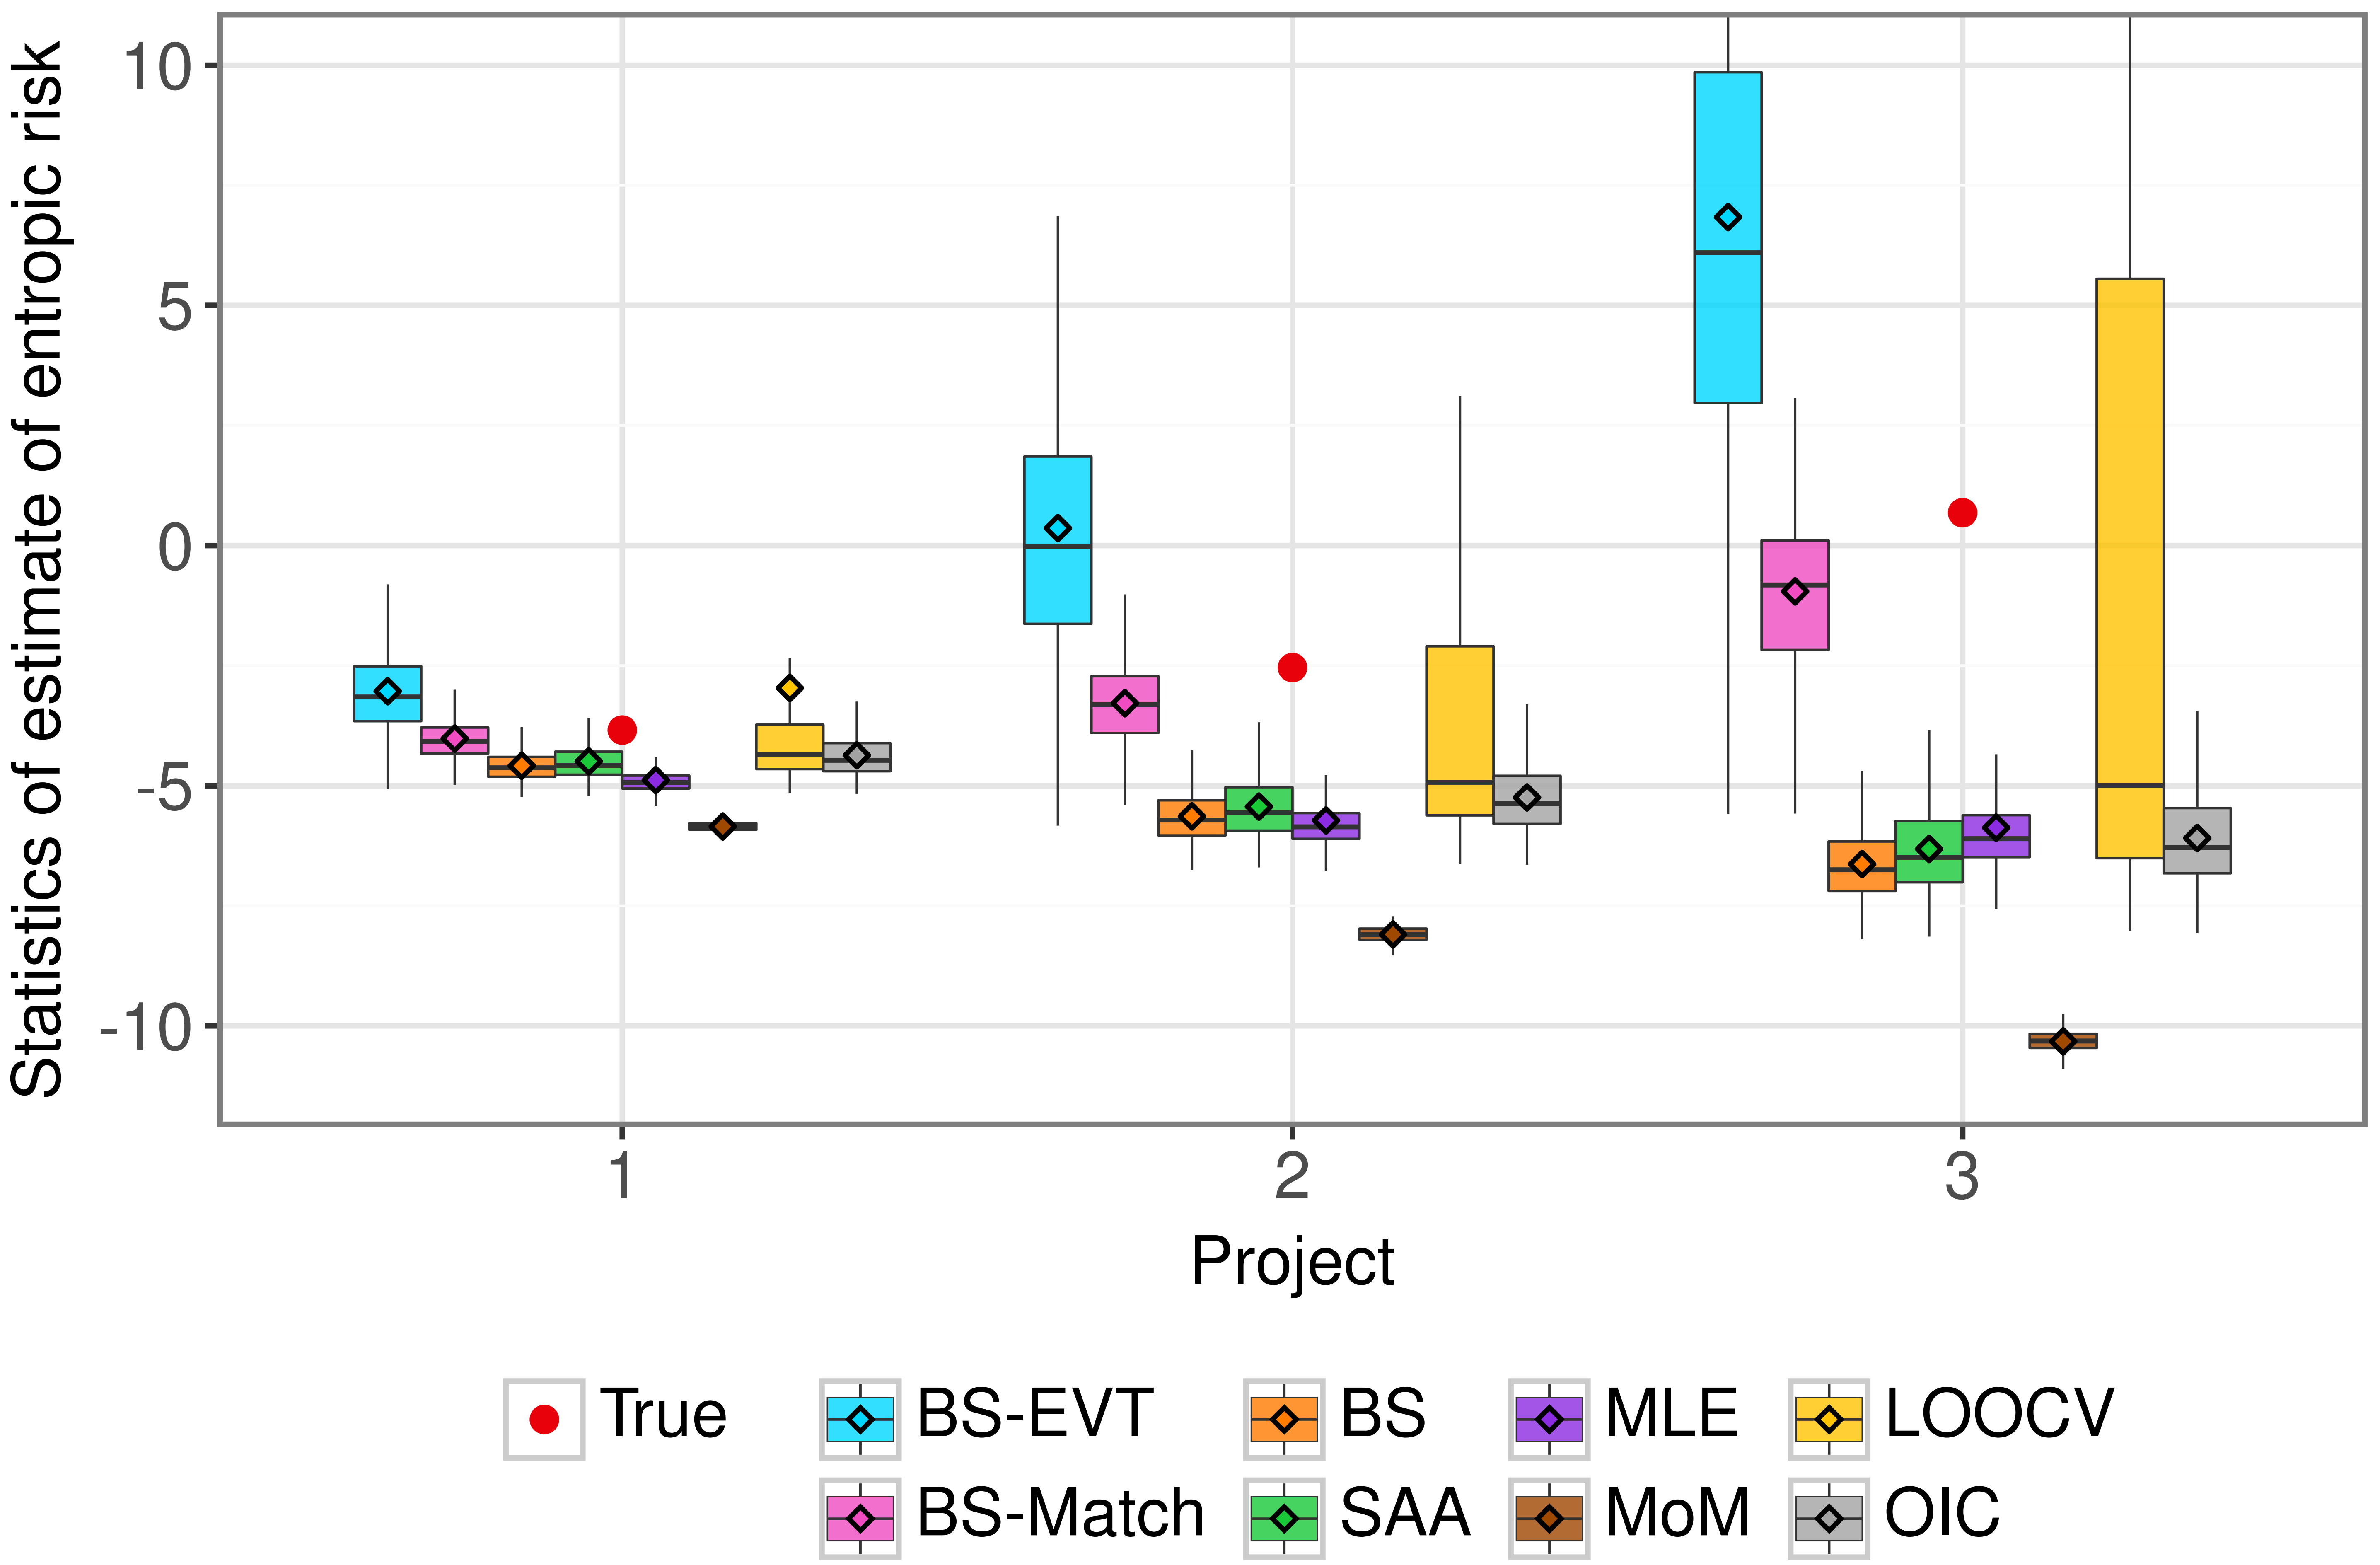

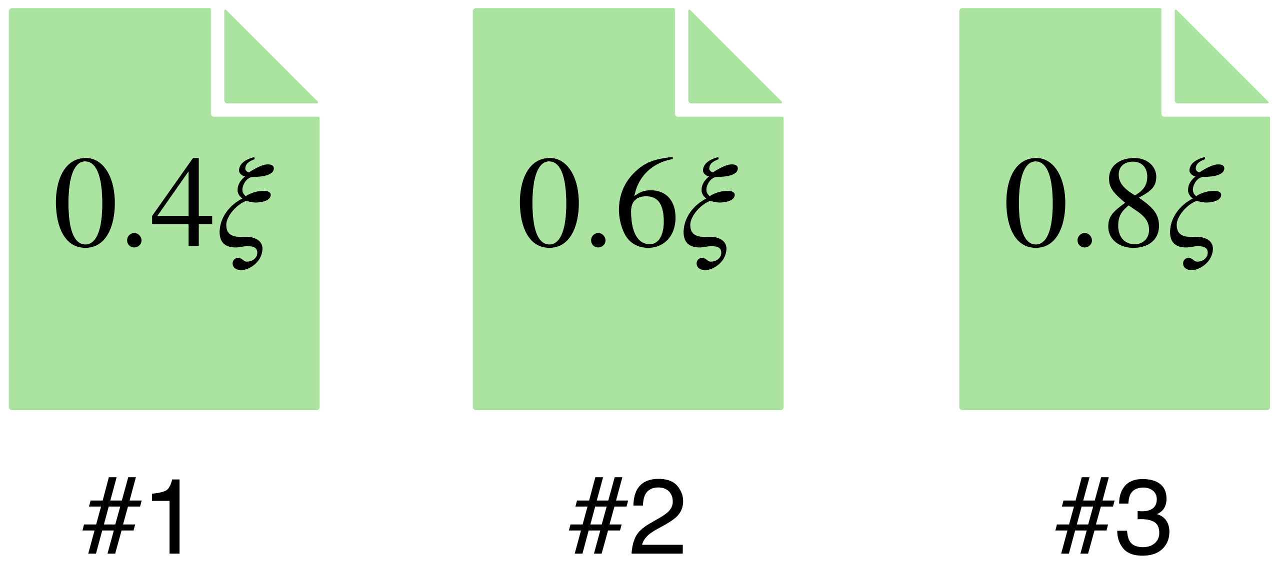

- $\xi \sim$ GMM$(\boldsymbol{\pi}, \boldsymbol{\mu}, \Sigma)$ with 5 components
- across components - $\mu_\xi = -18.6$ $\sigma_\xi = 2.9$
- Which project has lowest entropic risk based on 100 sets of 10000 samples with $\alpha = 3$?

Going from estimation to optimization

Distributionally robust optimization

- Loss depends on $z \in \mathcal{Z}$:

$$\rho^* = \min_{z \in \mathcal{Z}} \rho_{\mathbb{P}}(\ell(z, \boldsymbol{\xi}))$$

- Sample average approximation

$$\rho_{SAA} = \min_{z \in \mathcal{Z}} \rho_{\hat{\mathbb{P}}_N}(\ell(z, \boldsymbol{\xi}))$$

- DRO accounts for distributional ambiguity:

$$\rho_{DRO} = \min_{z \in \mathcal{Z}} \sup_{\mathbb{Q} \in \mathscr{B}_p(\epsilon)} \rho_{\mathbb{Q}}(\ell(z, \boldsymbol{\xi}))$$

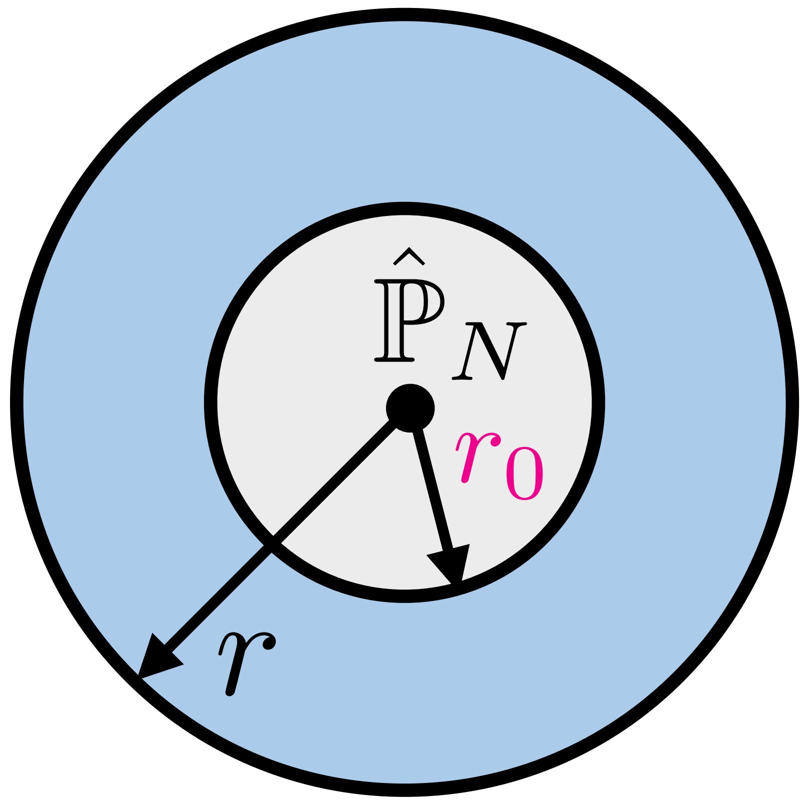

$\mathscr{B}_p(\epsilon)$

Distributionally robust optimization

Distributionally robust optimization

- KL divergence and Type-p Wasserstein ($p < \infty$): unbounded worst-case loss

Distributionally robust optimization

- KL divergence and Type-p Wasserstein ($p < \infty$): unbounded worst-case loss
- Type $\infty-$Wasserstein: bounded loss

Distributionally robust optimization

- ☒ KL divergence and Type-p Wasserstein ($p < \infty$): unbounded worst-case loss
- ☒ Type $\infty-$Wasserstein: bounded loss

Theorem: $\rho_{SAA} \to \rho^*$, $\rho_{DRO} \to \rho^*$ in probability at rate $\mathcal{O}(1/\sqrt{N})$

Regularized exponential cone program

Regularized exponential cone program

Theorem: With a linear loss function $\ell(z, \boldsymbol{\xi}) = z^\top \boldsymbol{\xi}$, DRO with type-$\infty$ Wasserstein ambiguity set reduces to:

$$\min_{z \in \mathscr{Z}} \frac{1}{\alpha} \log \left(\mathbb{E}_{\hat{\mathbb{P}}_N} \left[\exp(\alpha z^\top \boldsymbol{\xi}) \right] \right) + \epsilon \|z\|_*$$

Regularized exponential cone program

Theorem: With a linear loss function $\ell(z, \boldsymbol{\xi}) = z^\top \boldsymbol{\xi}$, DRO with type-$\infty$ Wasserstein ambiguity set reduces to:

$$\min_{z \in \mathcal{Z}} \frac{1}{\alpha} \log \left(\mathbb{E}_{\hat{\mathbb{P}}_N} \left[\exp(\alpha z^\top \boldsymbol{\xi}) \right] \right) + \epsilon \|z\|_*$$

- More general loss functions - refer to our paper
- How to choose the radius ϵ?
- **Validation data - underestimates the true risk**

  - suboptimal radius

  - Bias correction using bootstrapping

Distributionally robust insurance pricing

Distributionally robust insurance pricing

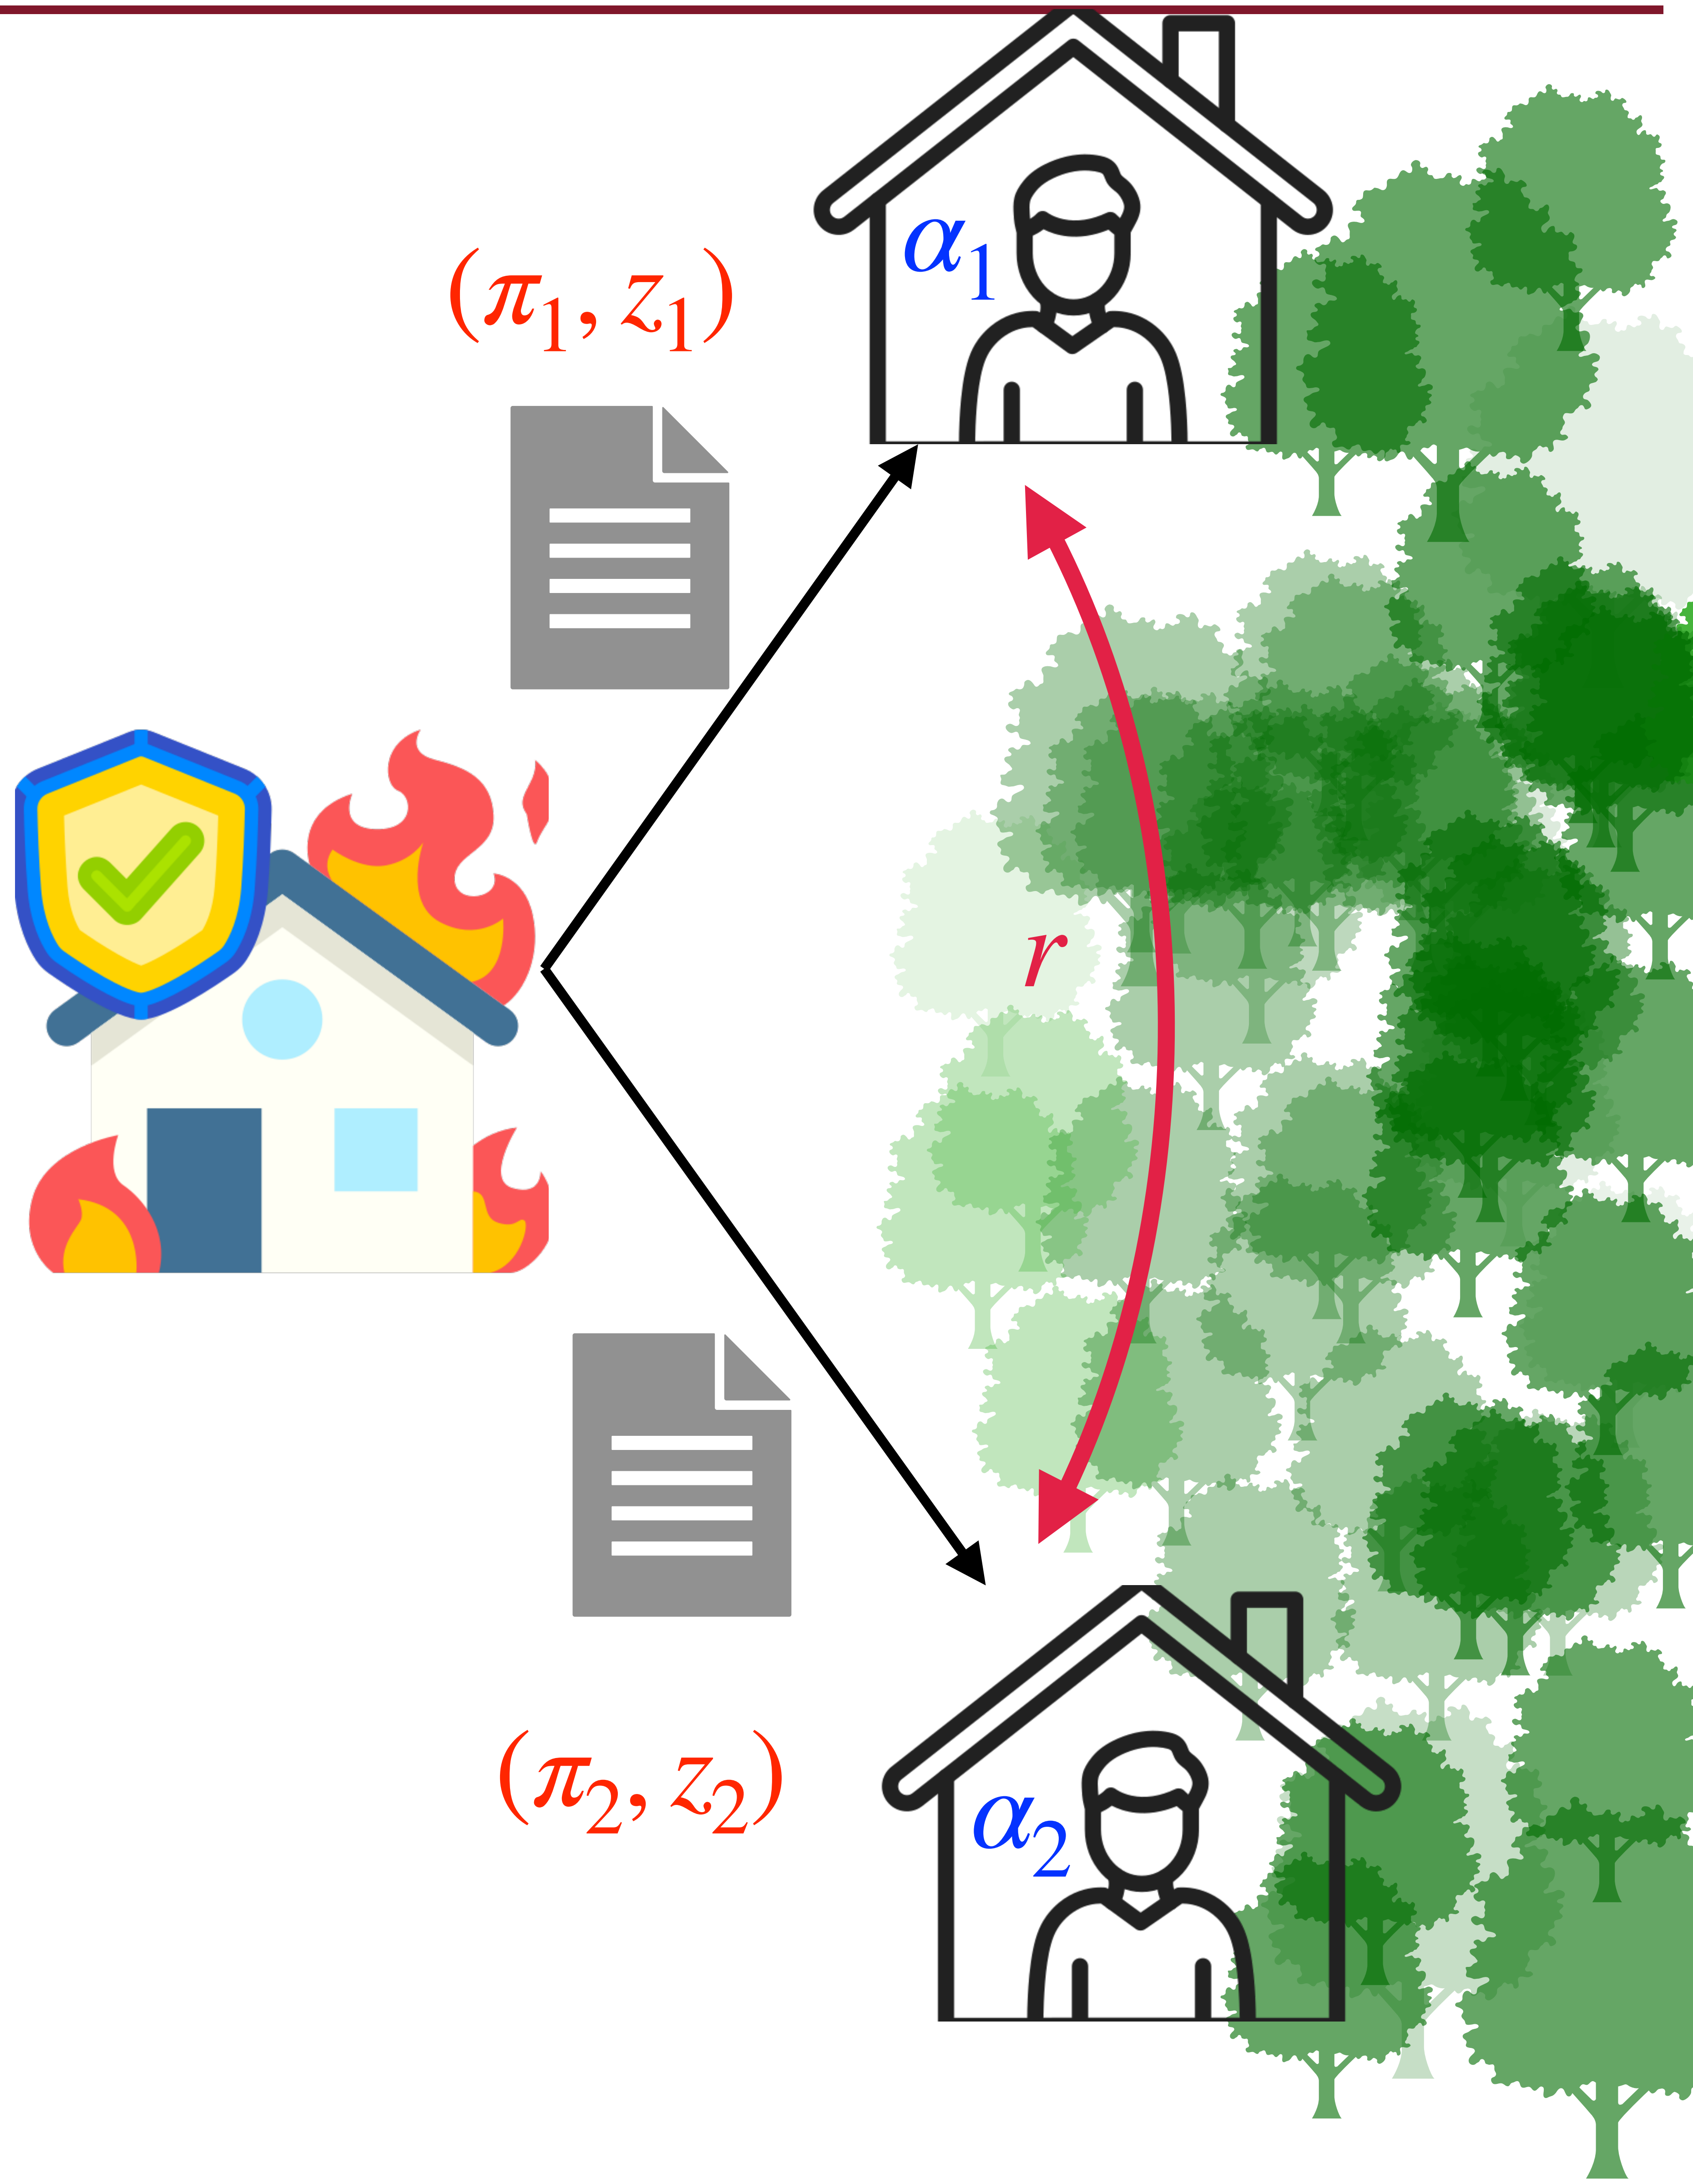

Distributionally robust insurance pricing

- Insurer offers coverage $z_h\xi$ at premium π_h
- α_h: homeowner's risk preference
- α_0: insurer's risk preference

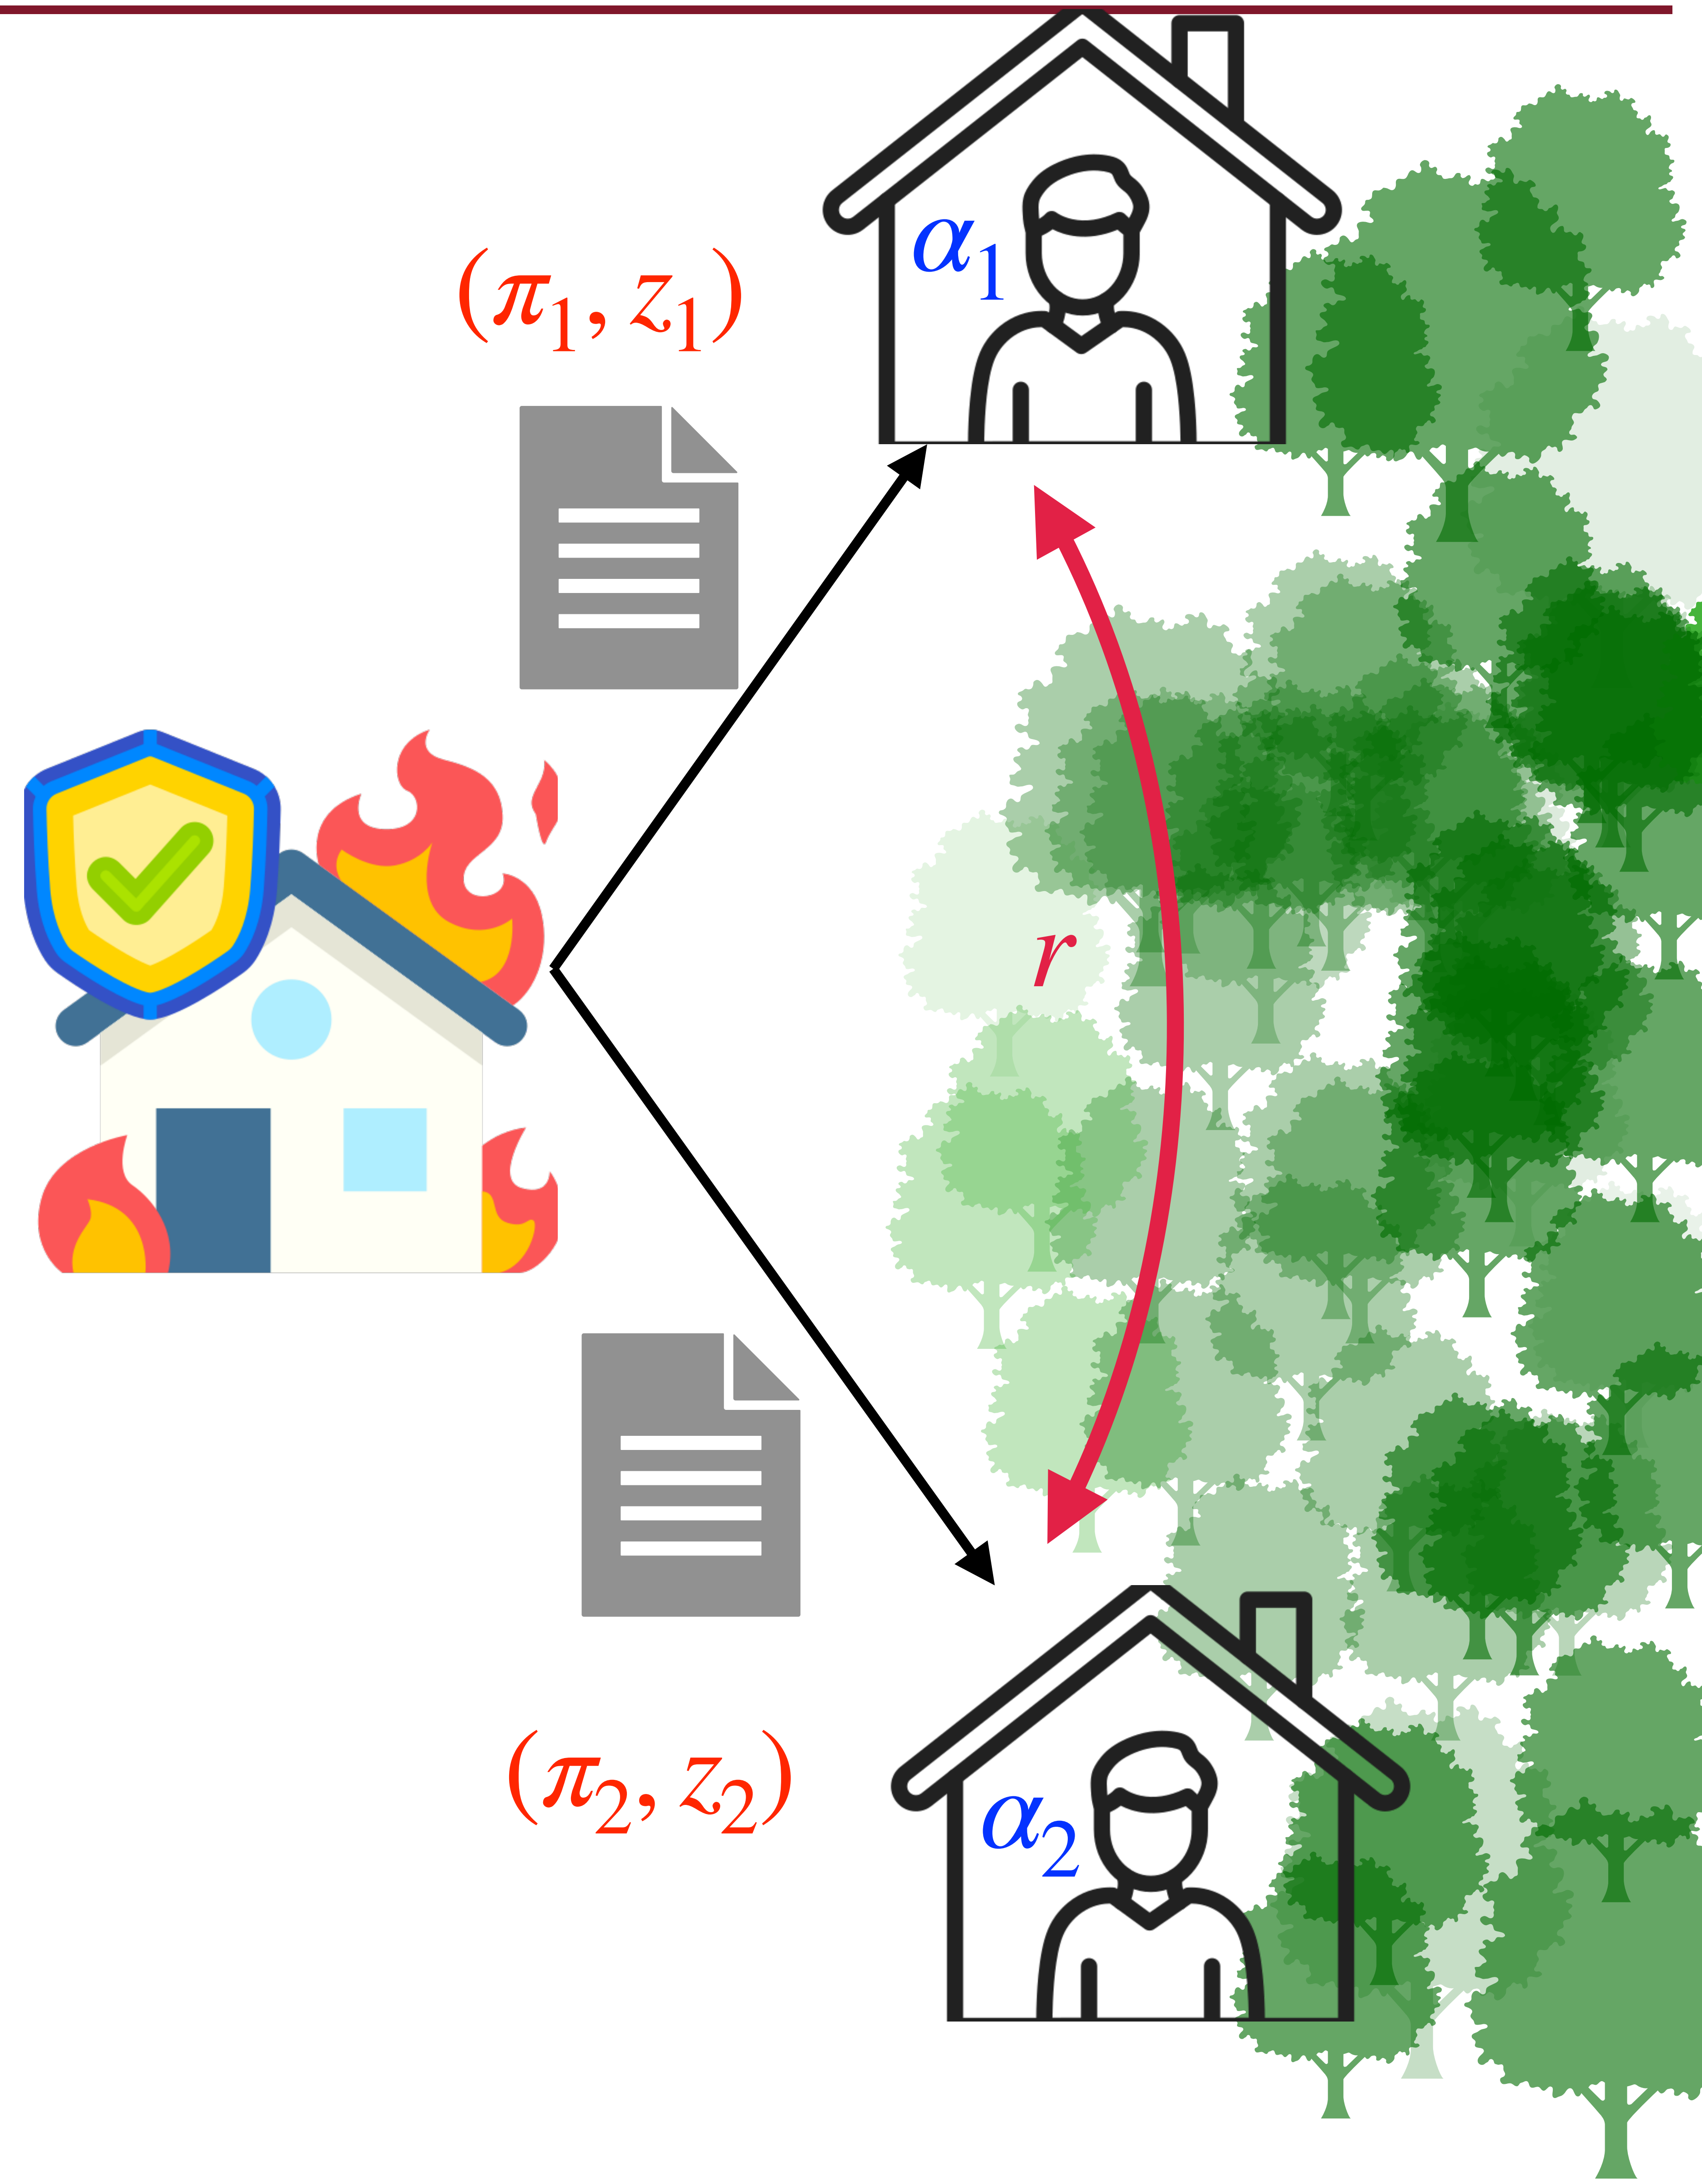

Distributionally robust insurance pricing

- Insurer offers coverage $z_h \xi$ at premium π_h
- α_h: homeowner's risk preference
- α_0: insurer's risk preference

$$\min \sup_{\mathbb{Q} \in \mathscr{B}_\infty(\epsilon)} \quad \rho_{\mathbb{Q}}^{\alpha_0}\left(\boldsymbol{z}^\top \boldsymbol{\xi} - \mathbf{1}^\top \boldsymbol{\pi}\right) + \epsilon \|z\|_*$$

$$\text{s.t.} \quad \boldsymbol{\pi} \in \mathbb{R}_+^M, z \in [0,1]^M$$

$$\boxed{\rho_{\hat{\mathbb{P}}_{h,N}}^{\alpha_h}\left(\pi_h + (1 - z_h)\xi_h\right) \leq \rho_{\hat{\mathbb{P}}_{h,N}}^{\alpha_h}\left(\xi_h\right)} \ \forall h$$

Demand response model: Household accept/reject the contract based on their estimate of empirical entropic risk

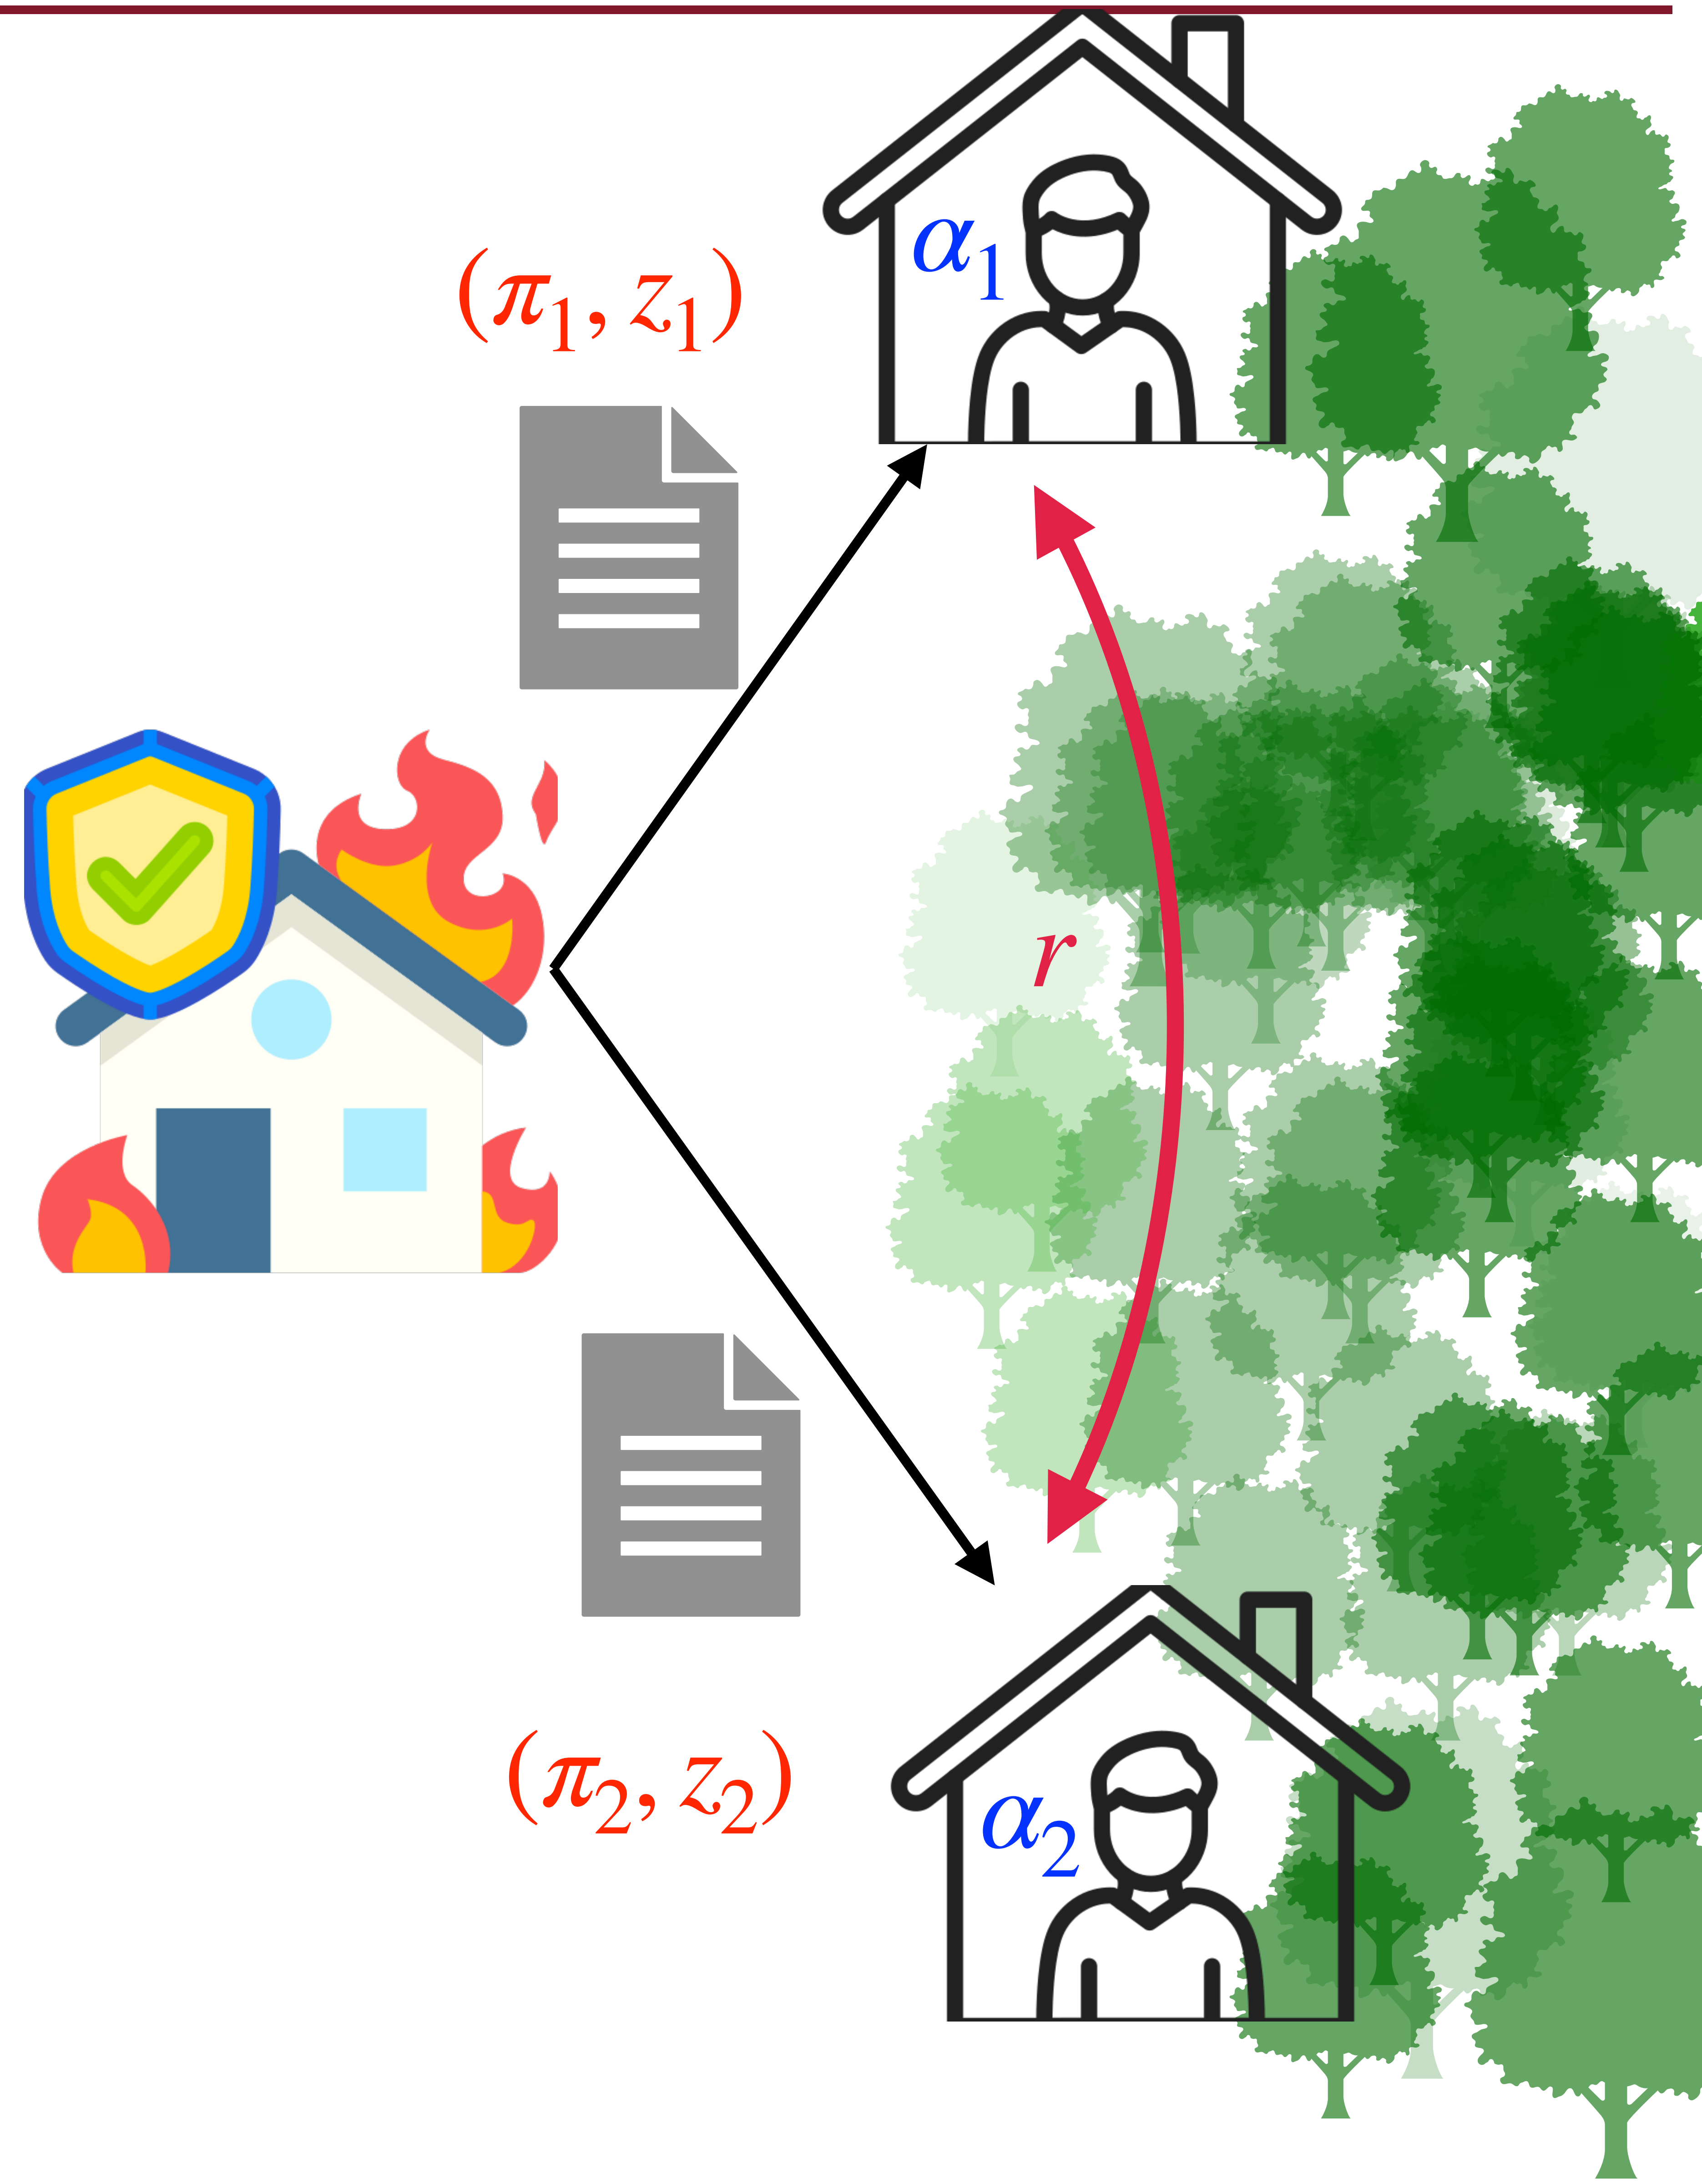

Reformulation as exponential cone

- A coverage of $z_h \xi$ is offered at premium π_h
- α_h: homeowner's risk preference
- α_0: insurer's risk preference

$$\begin{aligned} \min \quad & \rho^{\alpha_0}_{\hat{\mathbb{P}}_N}\left(\boldsymbol{z}^\top \boldsymbol{\xi} - \mathbf{1}^\top \boldsymbol{\pi}\right) + \epsilon \|\boldsymbol{z}\|_* \\ \text{s.t.} \quad & \boldsymbol{\pi} \in \mathbb{R}^M_+, \boldsymbol{z} \in [0,1]^M \\ & \rho^{\alpha_h}_{\hat{\mathbb{P}}_{h,N}}\left(\pi_h + (1 - z_h)\xi_h\right) \leq \rho^{\alpha_h}_{\hat{\mathbb{P}}_{h,N}}\left(\xi_h\right) \;\; \forall h \end{aligned}$$

Data for numerical experiments:

Loss scenarios are generated from Gaussian copula with $\Gamma(\kappa_h, \lambda_h)$ marginals

Out-of-sample risk and radius - vary N

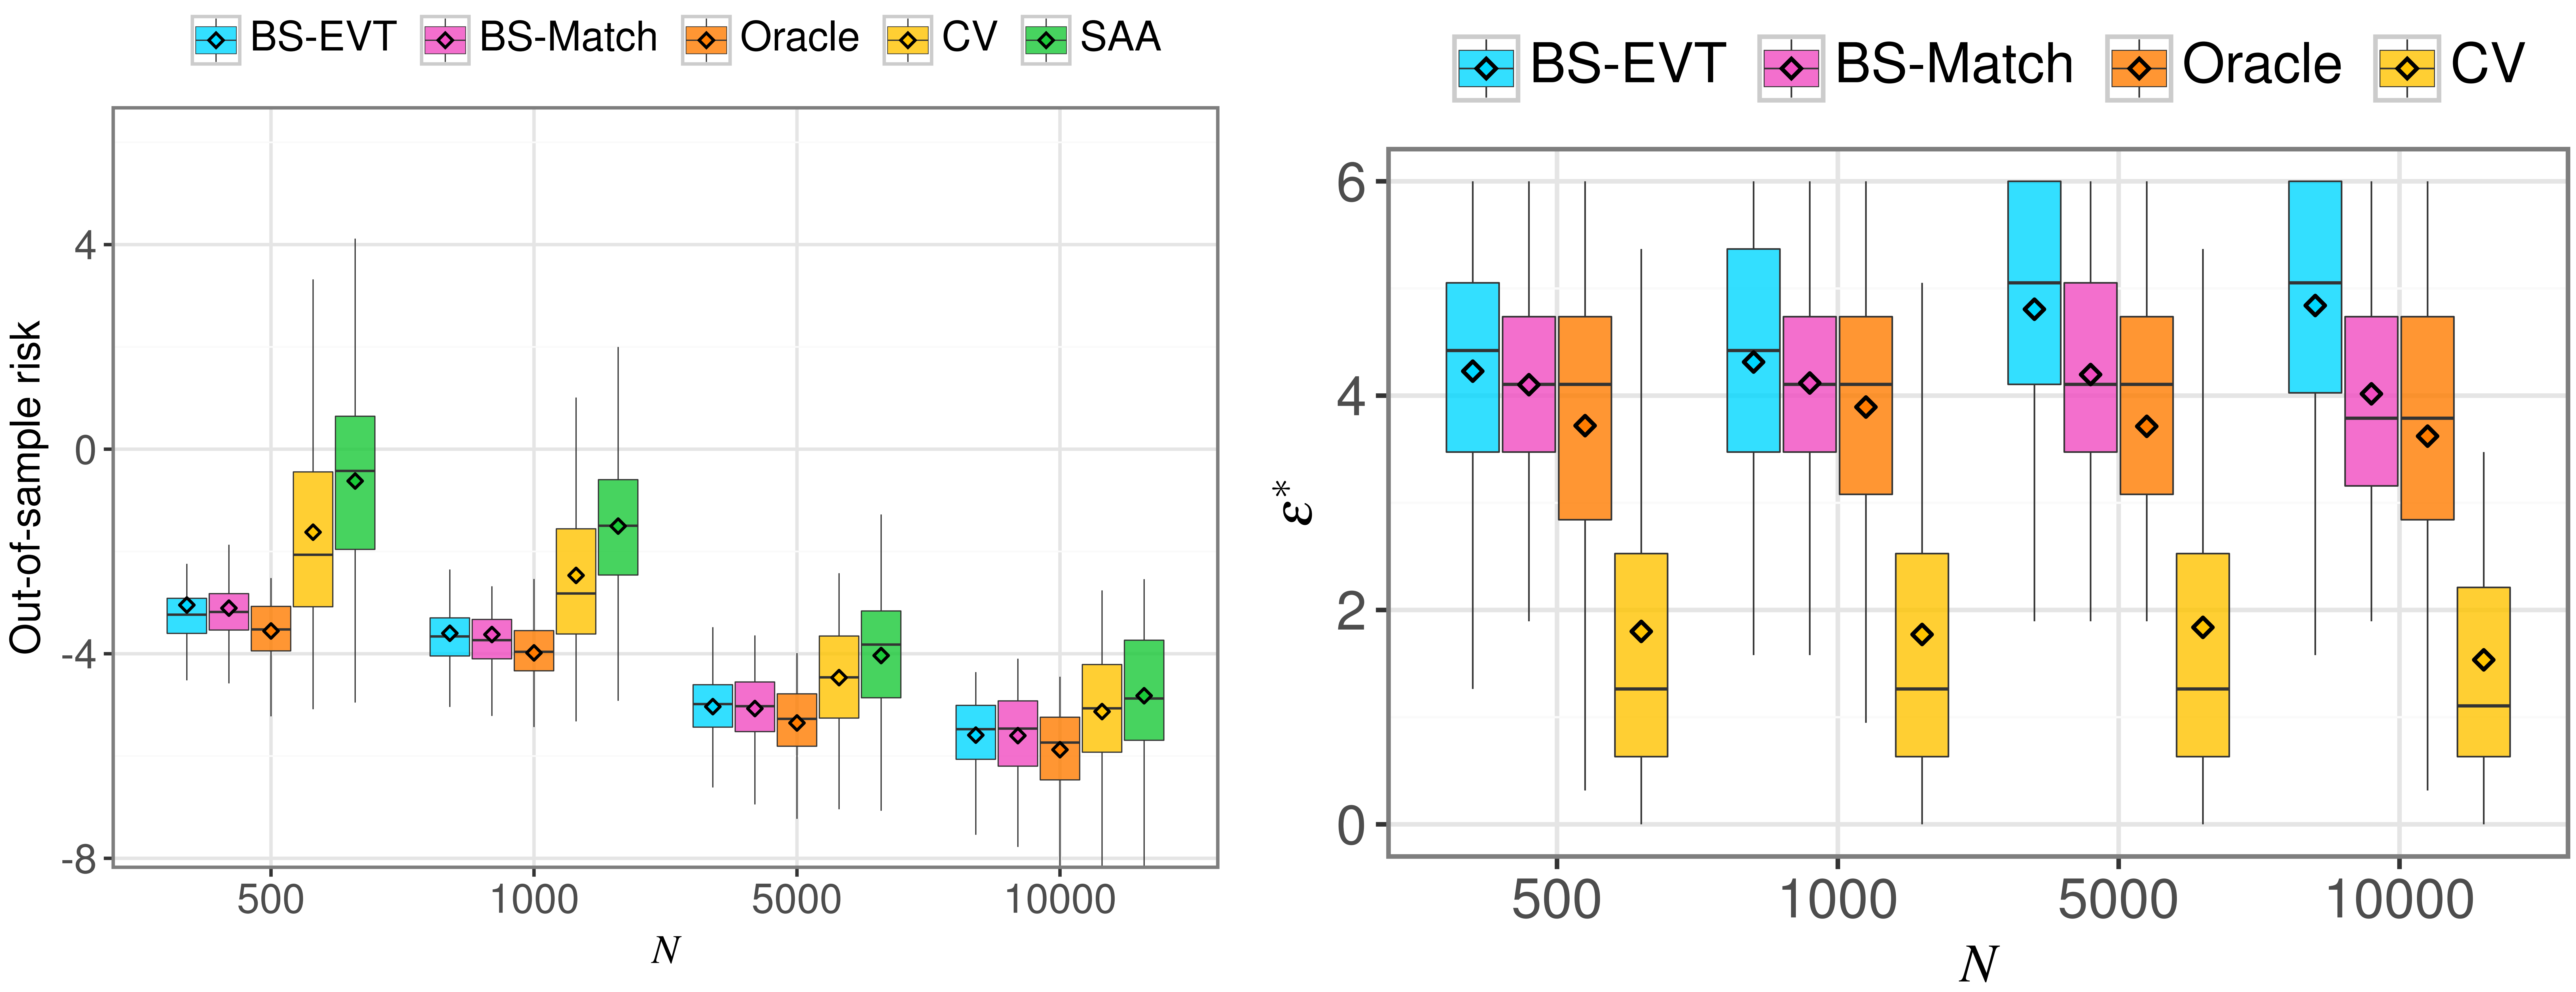

Risk decreases as training samples increase

Our models choose higher radius while traditional CV chooses lower radius

Premium per unit coverage - vary N

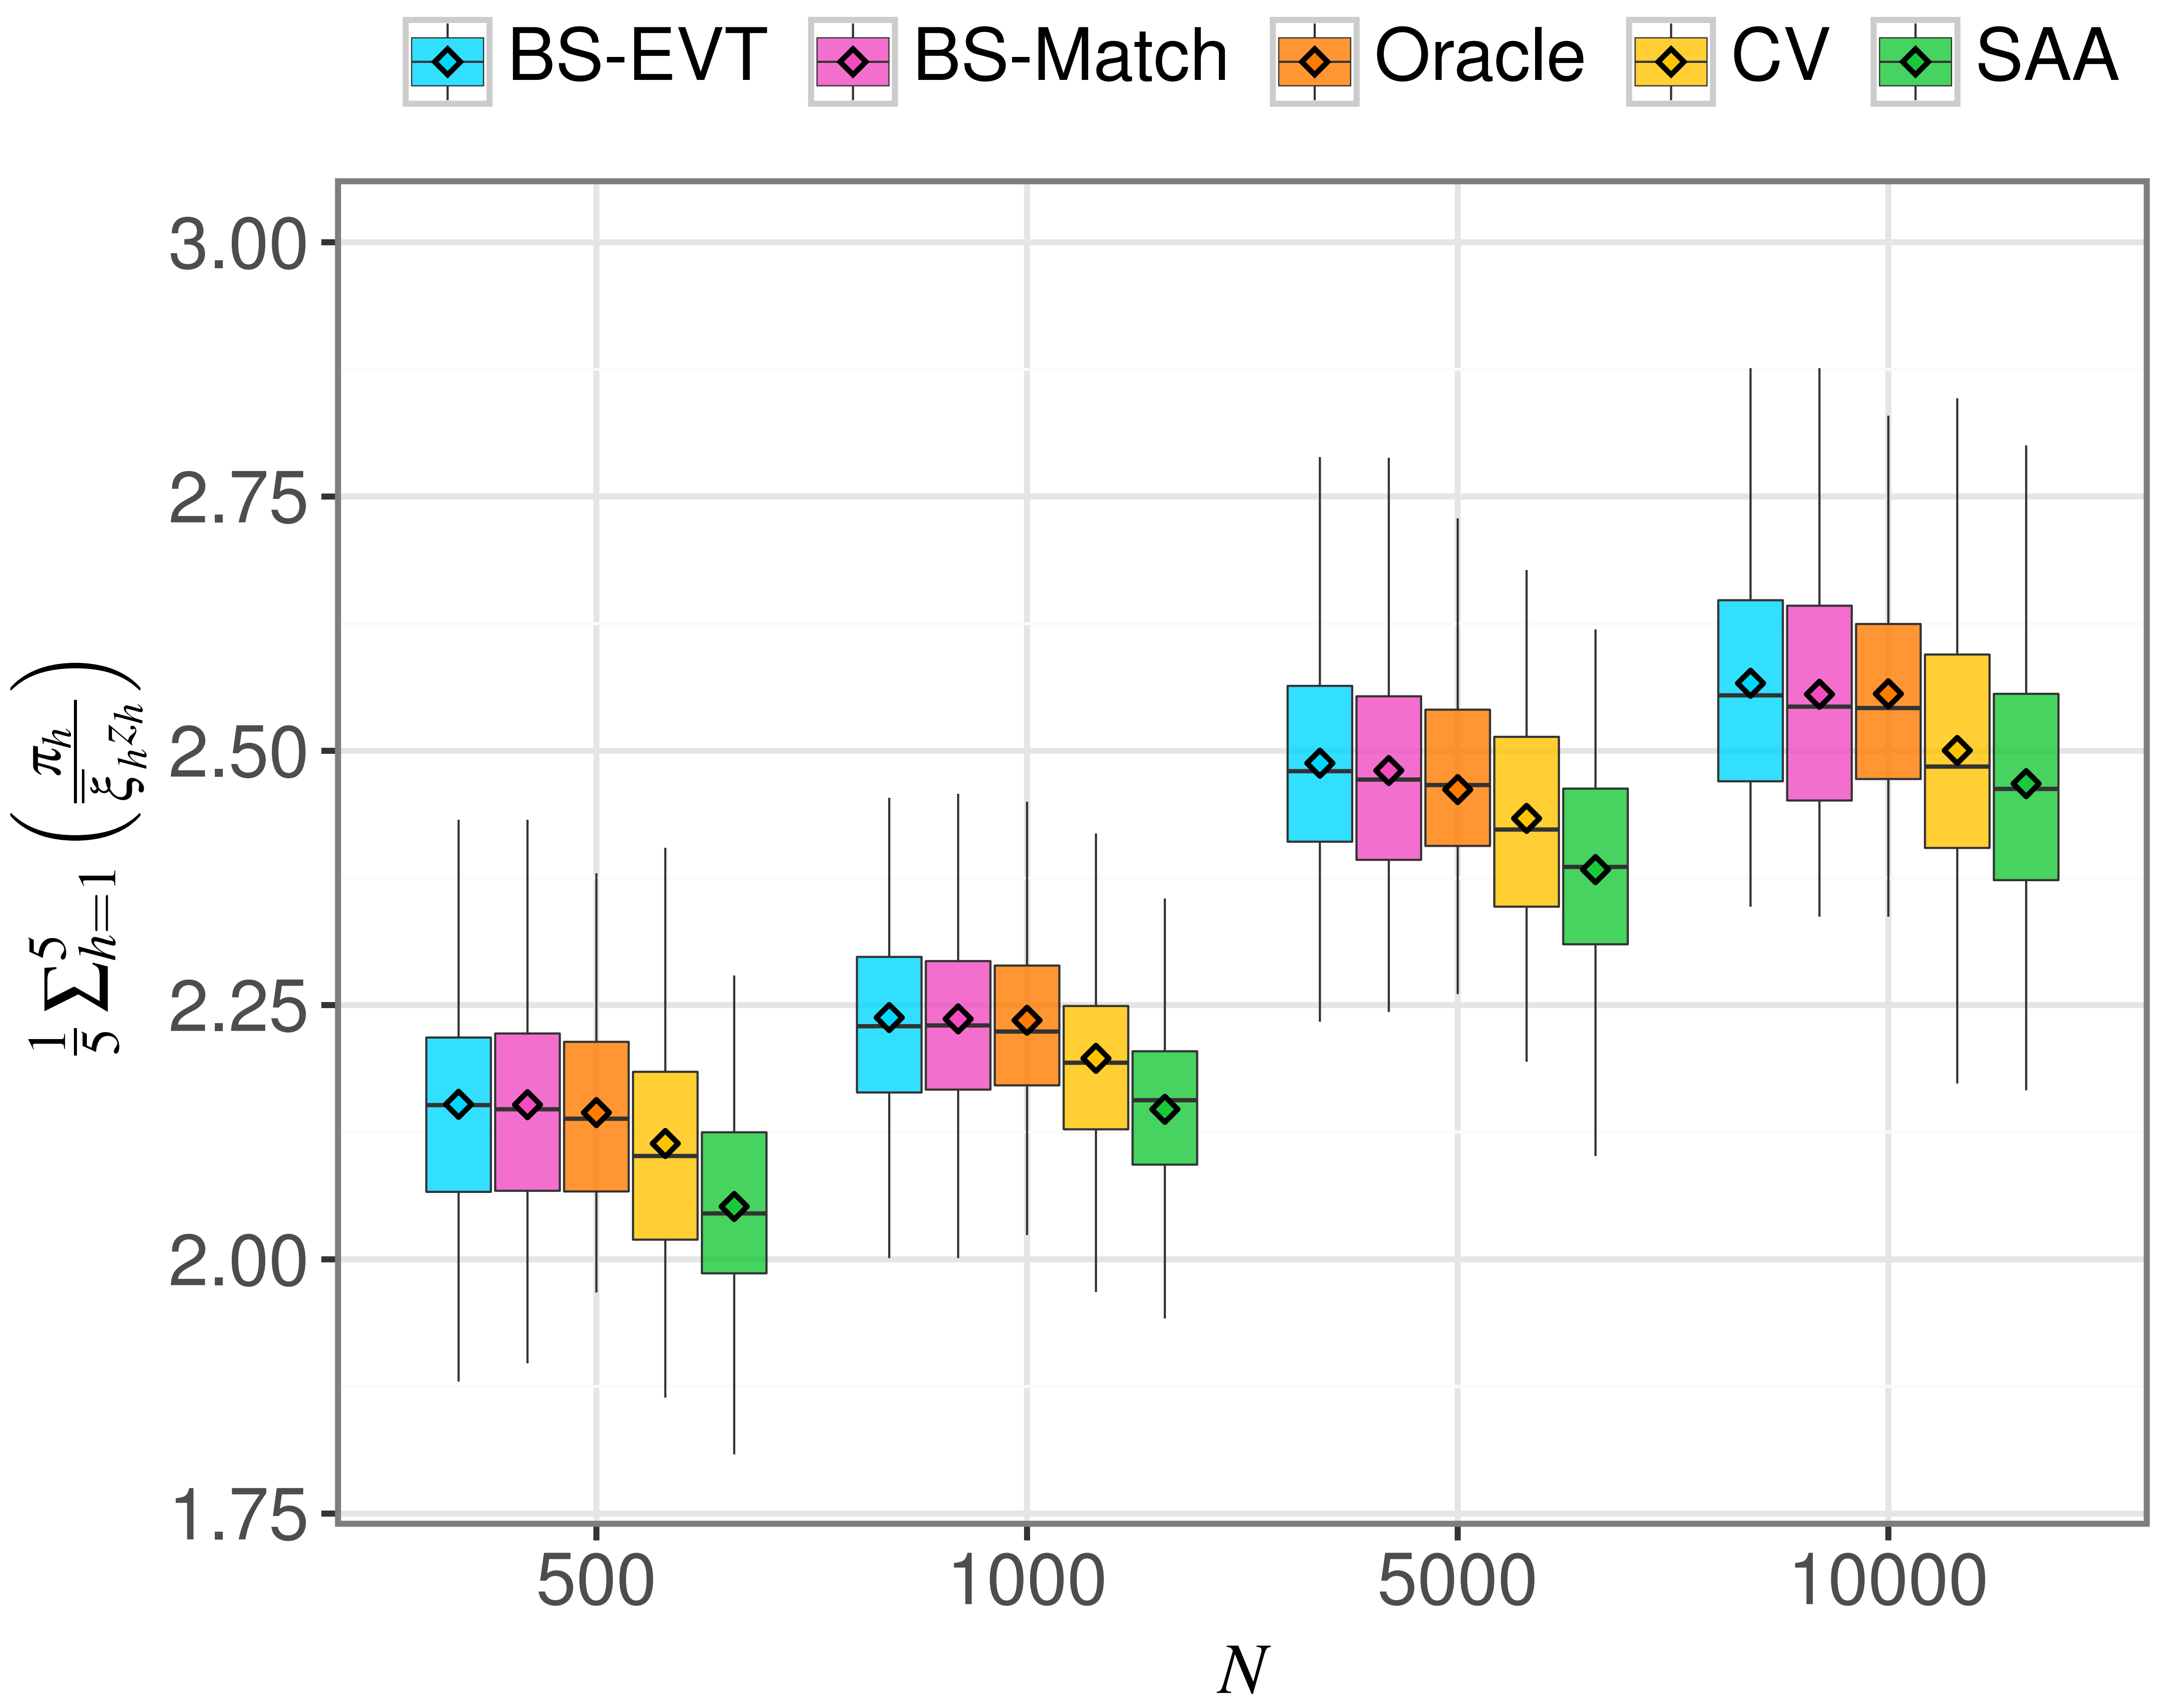

Households pay higher premiums as their estimates of risk improve with N

Out-of-sample risk and radius - vary correlation

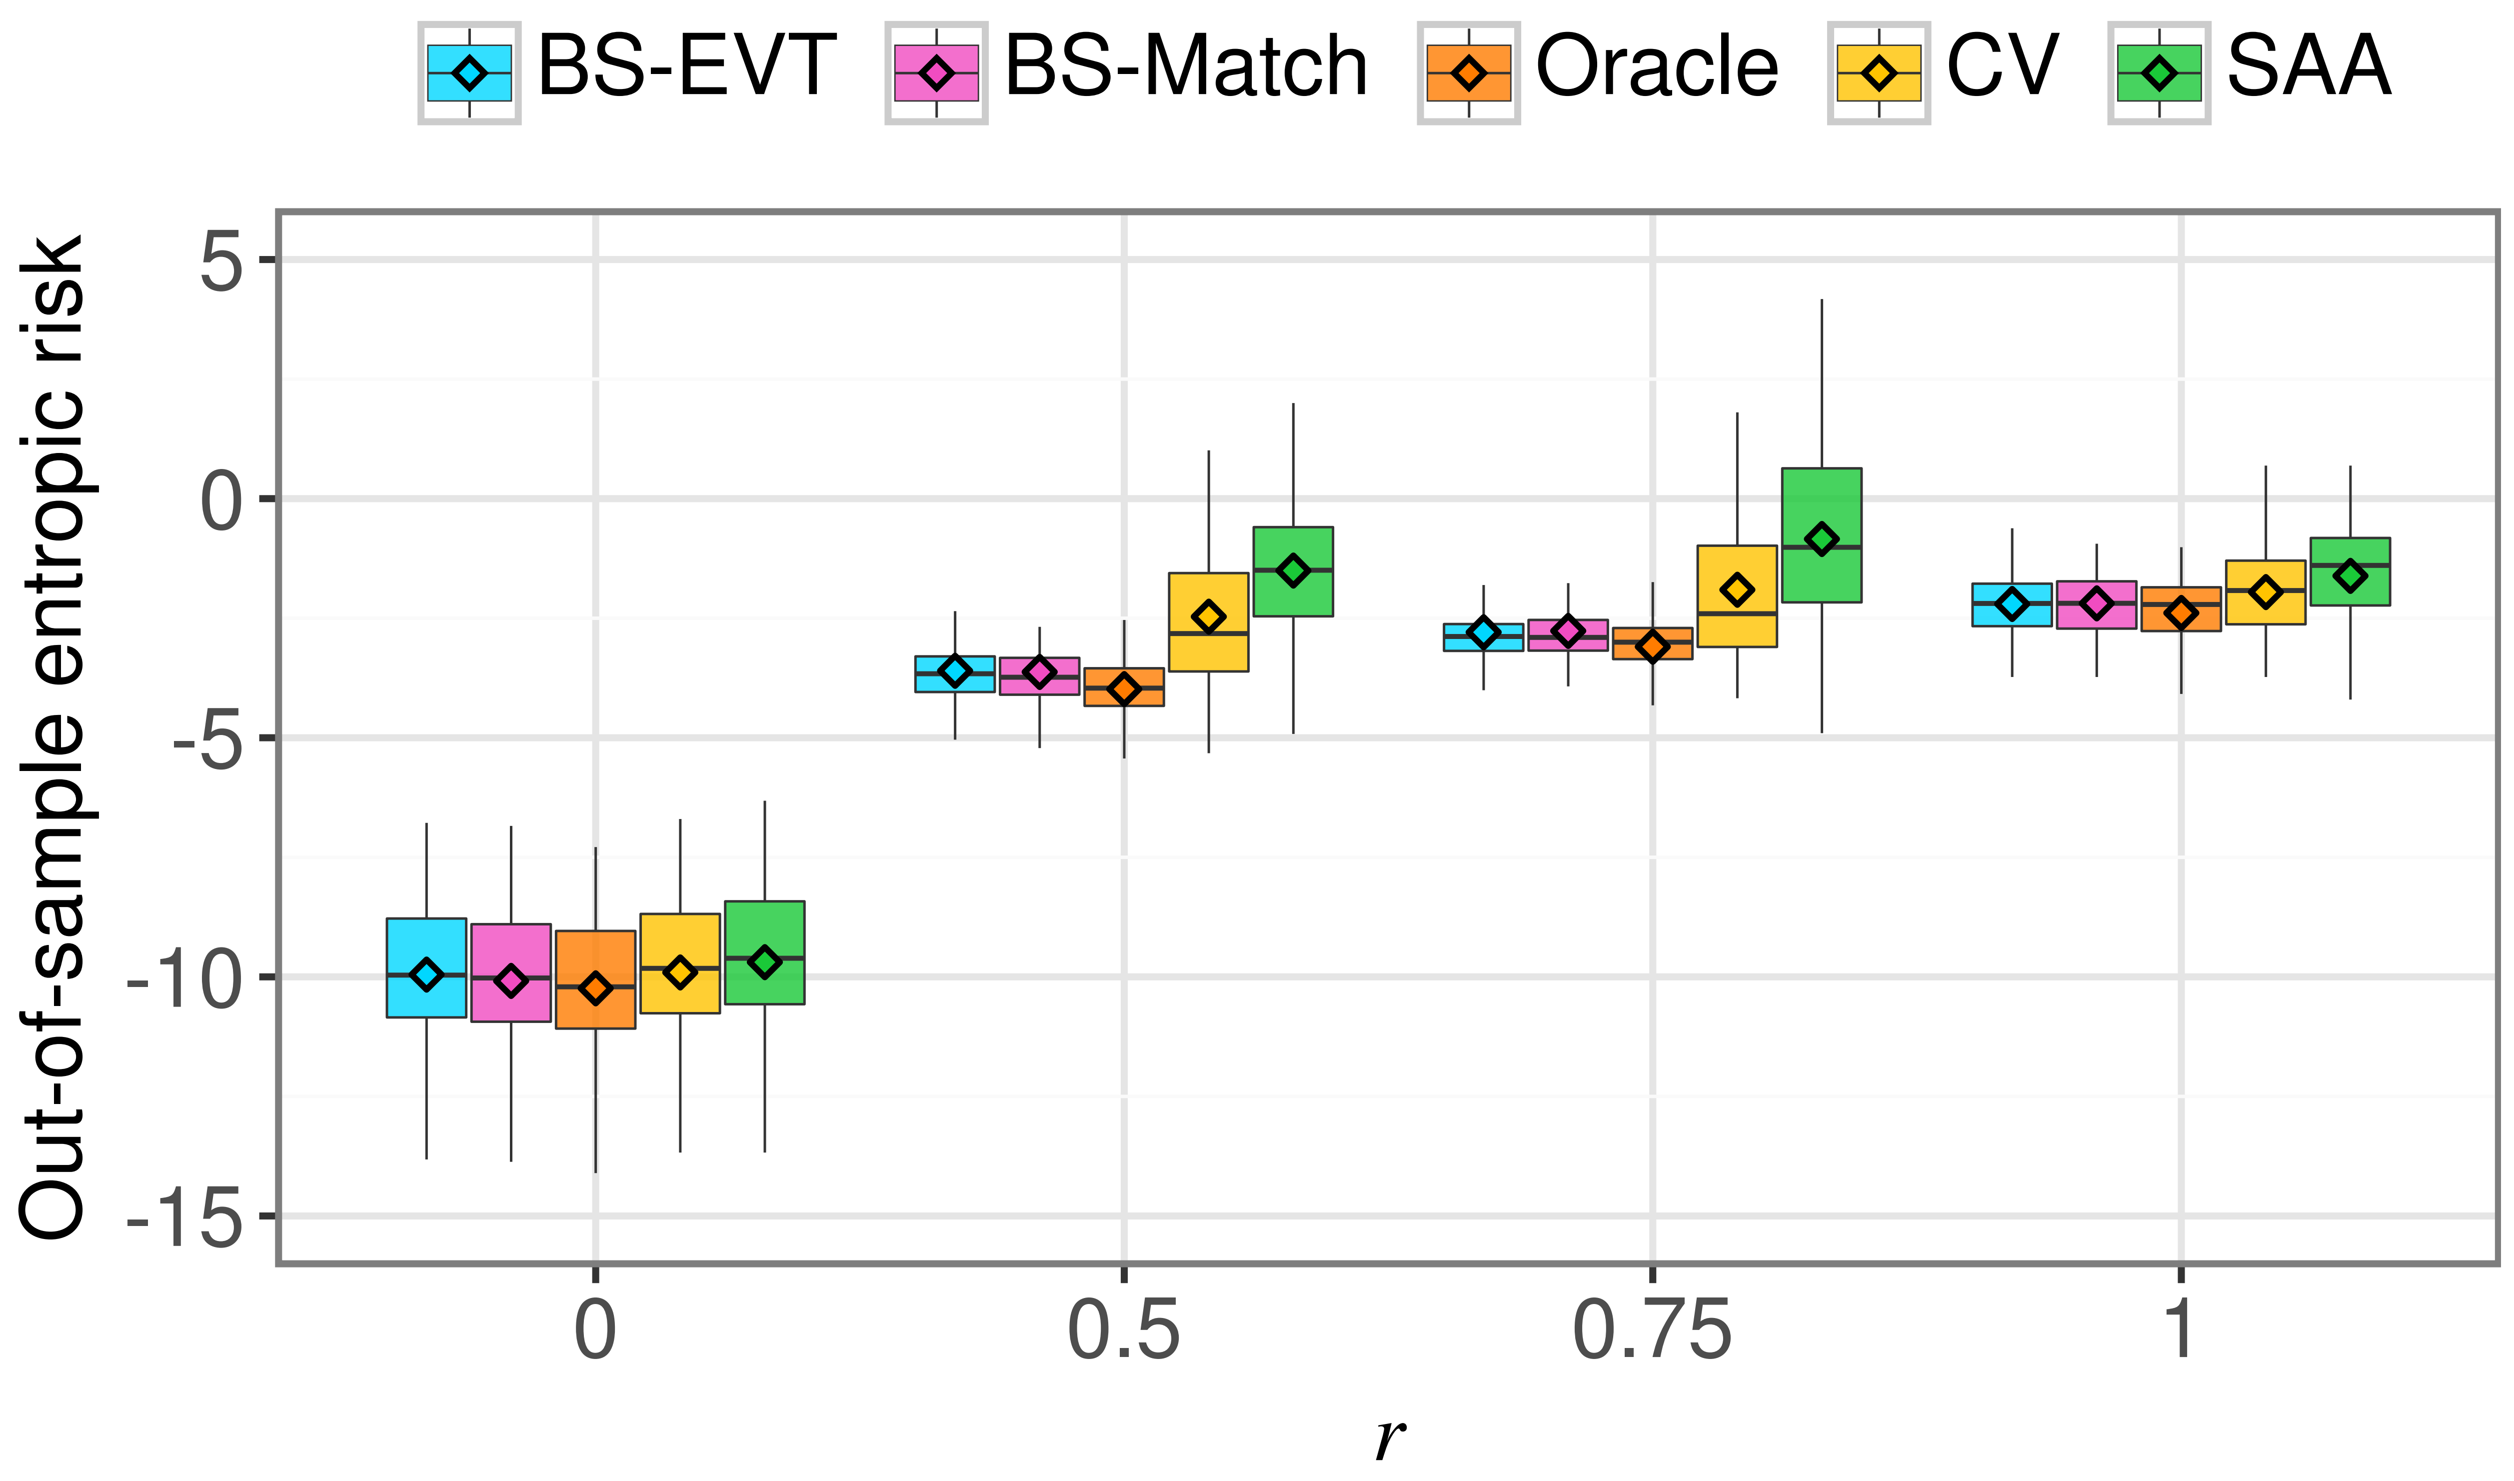

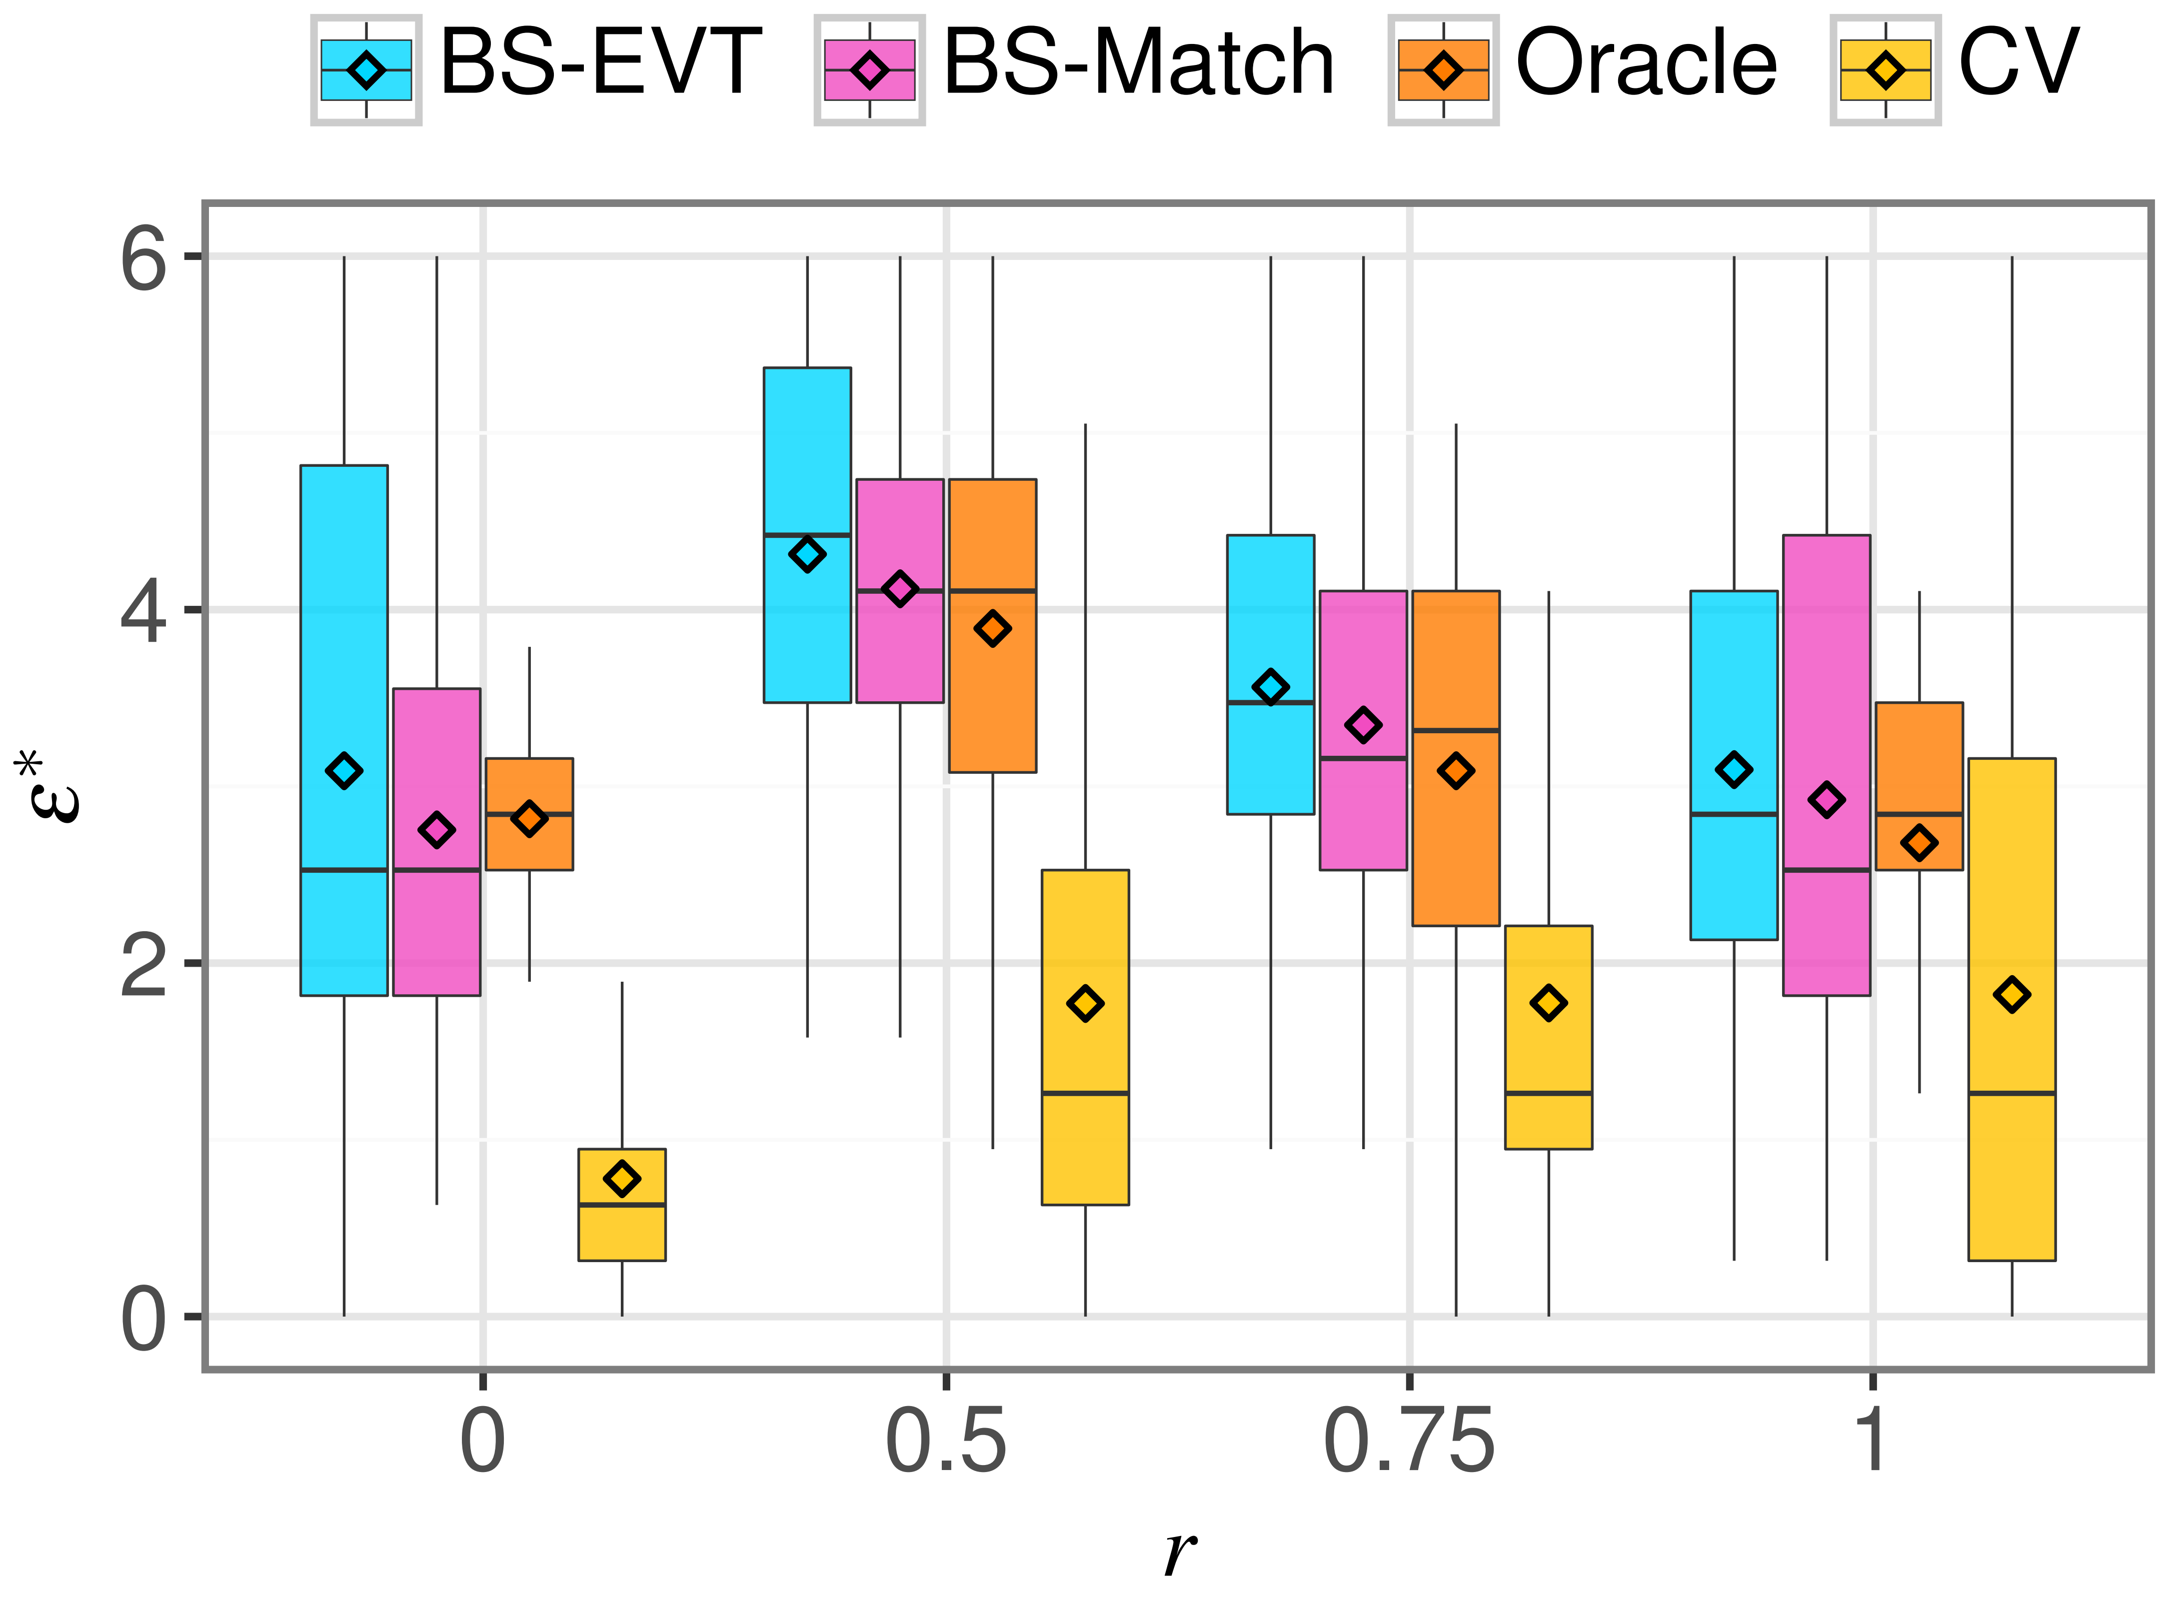

High correlation: extreme loss events more likely to occur simultaneously, increasing insurer’s risk exposure

Premium per unit coverage - vary correlation

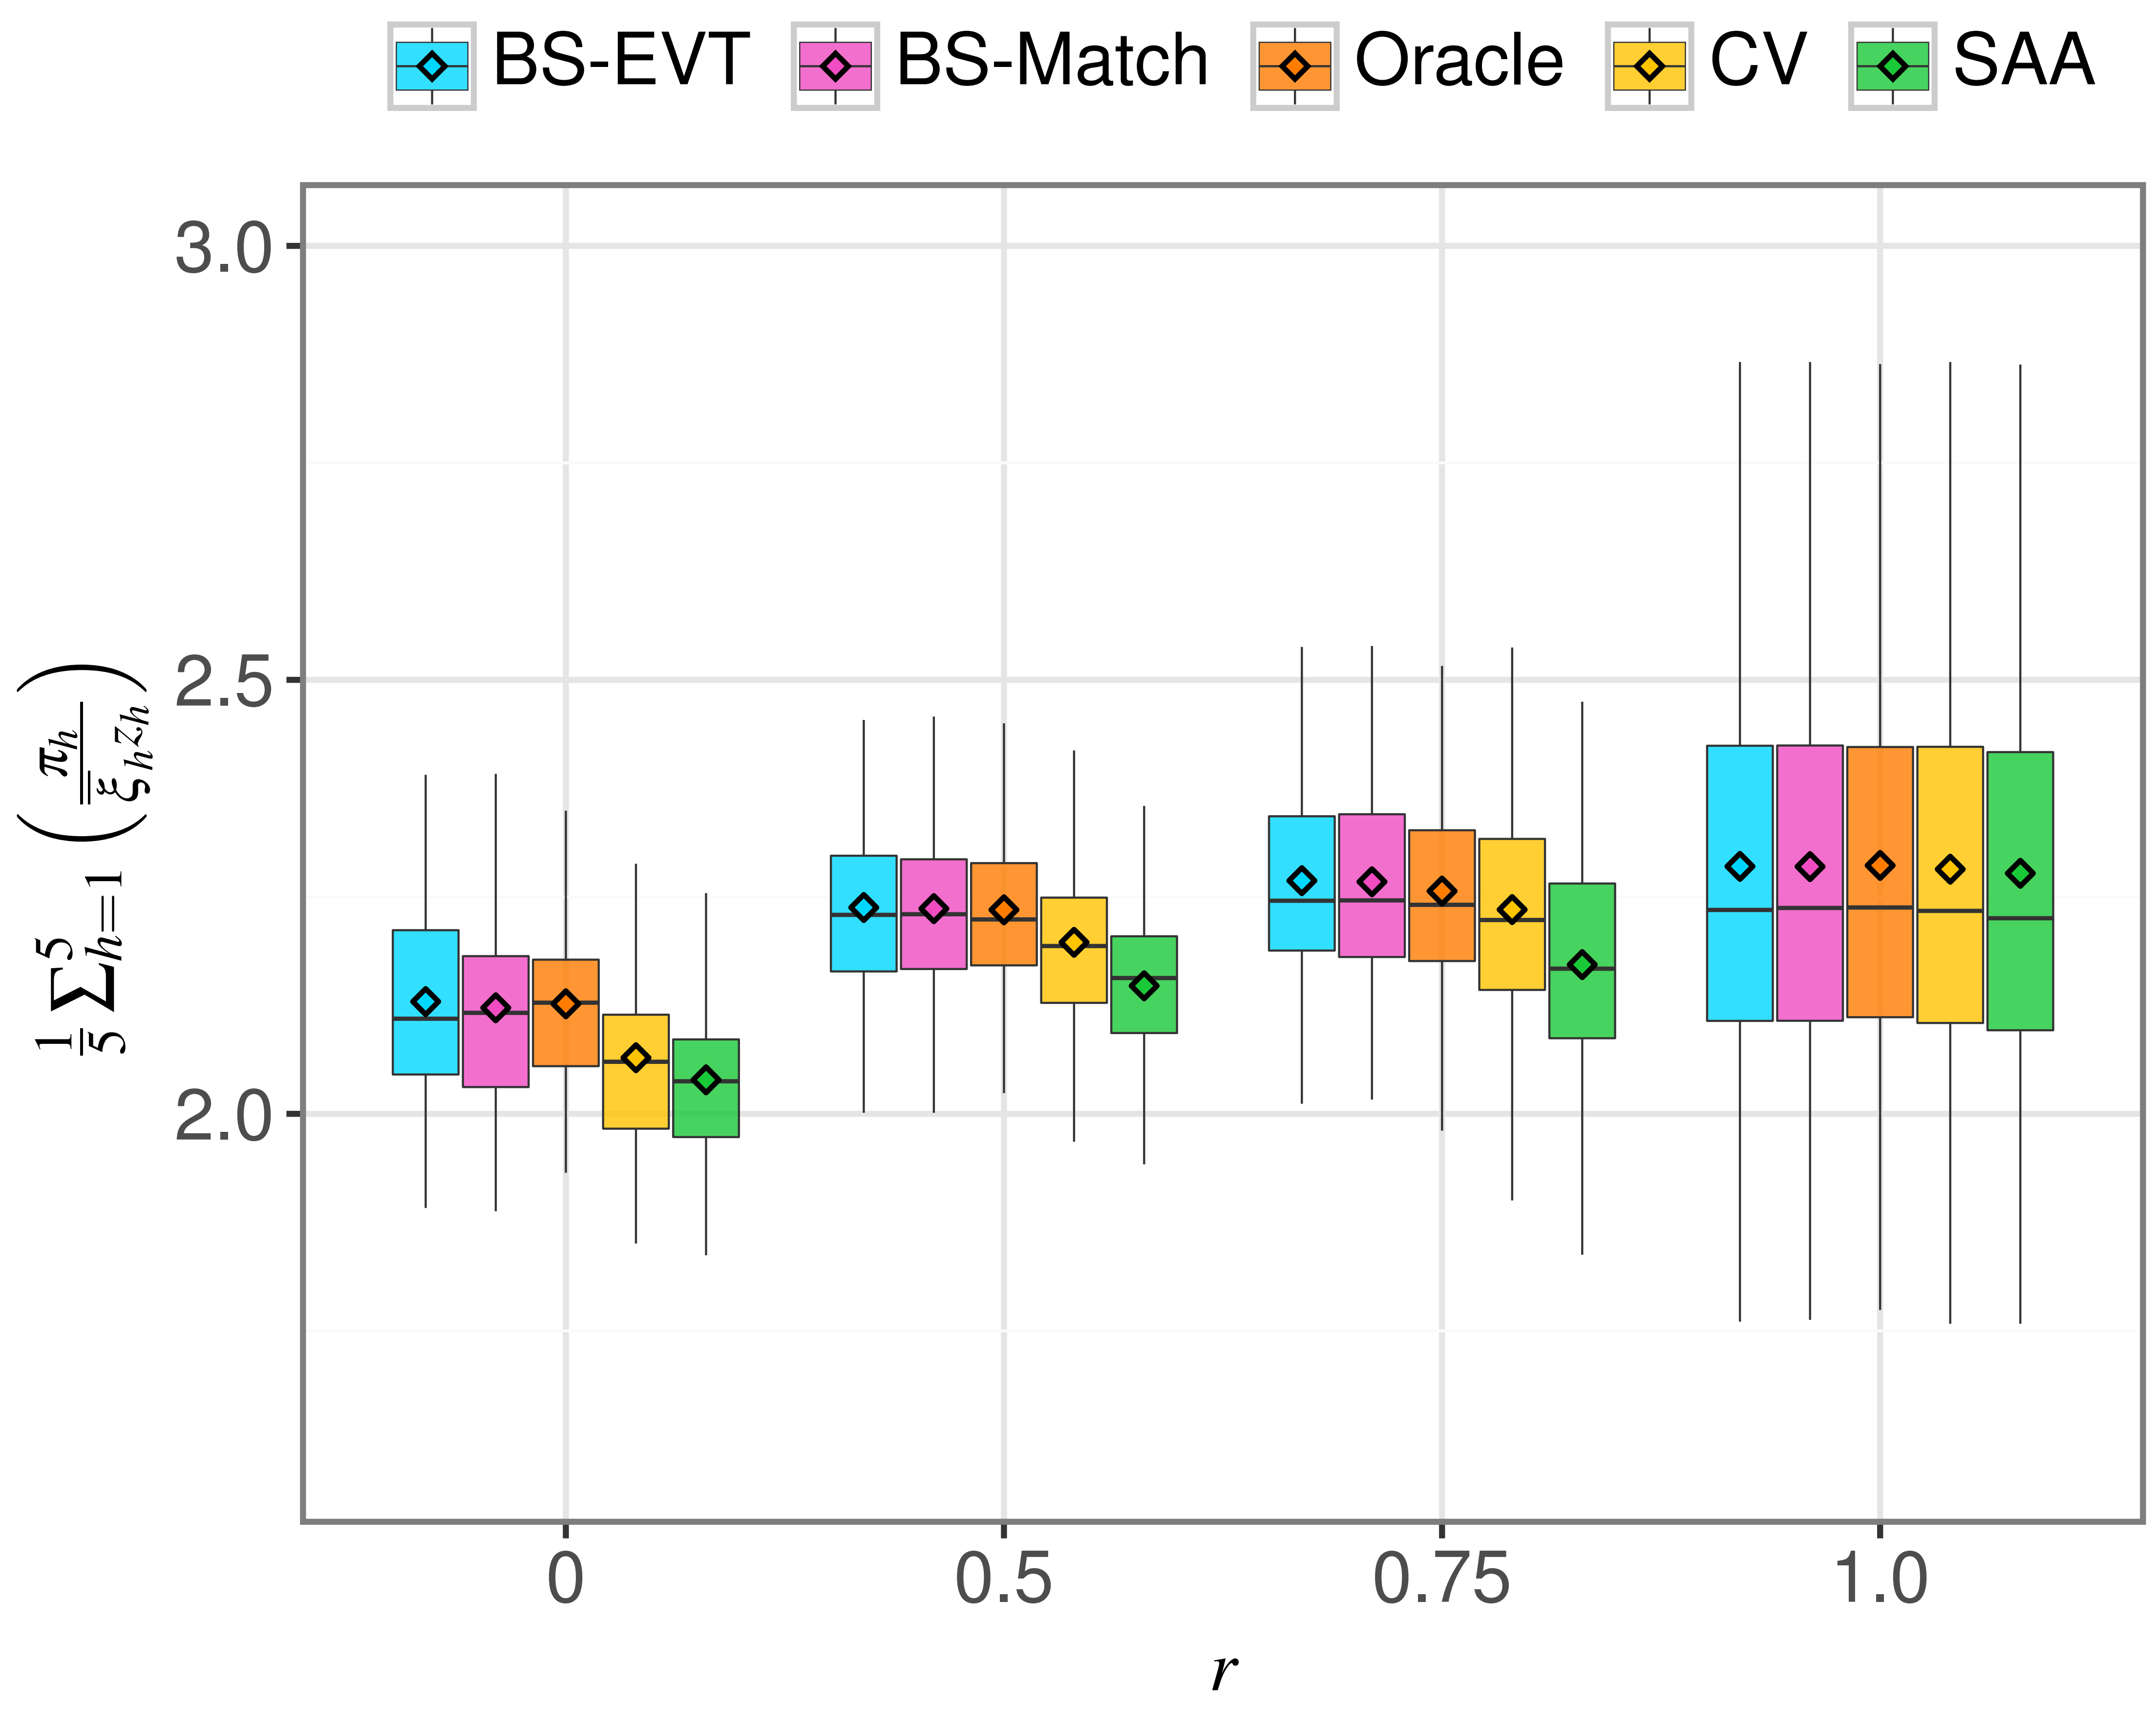

High correlation: benefits of risk pooling diminish, reduce coverage significantly to reduce risk exposure

Why our models identify better radius?

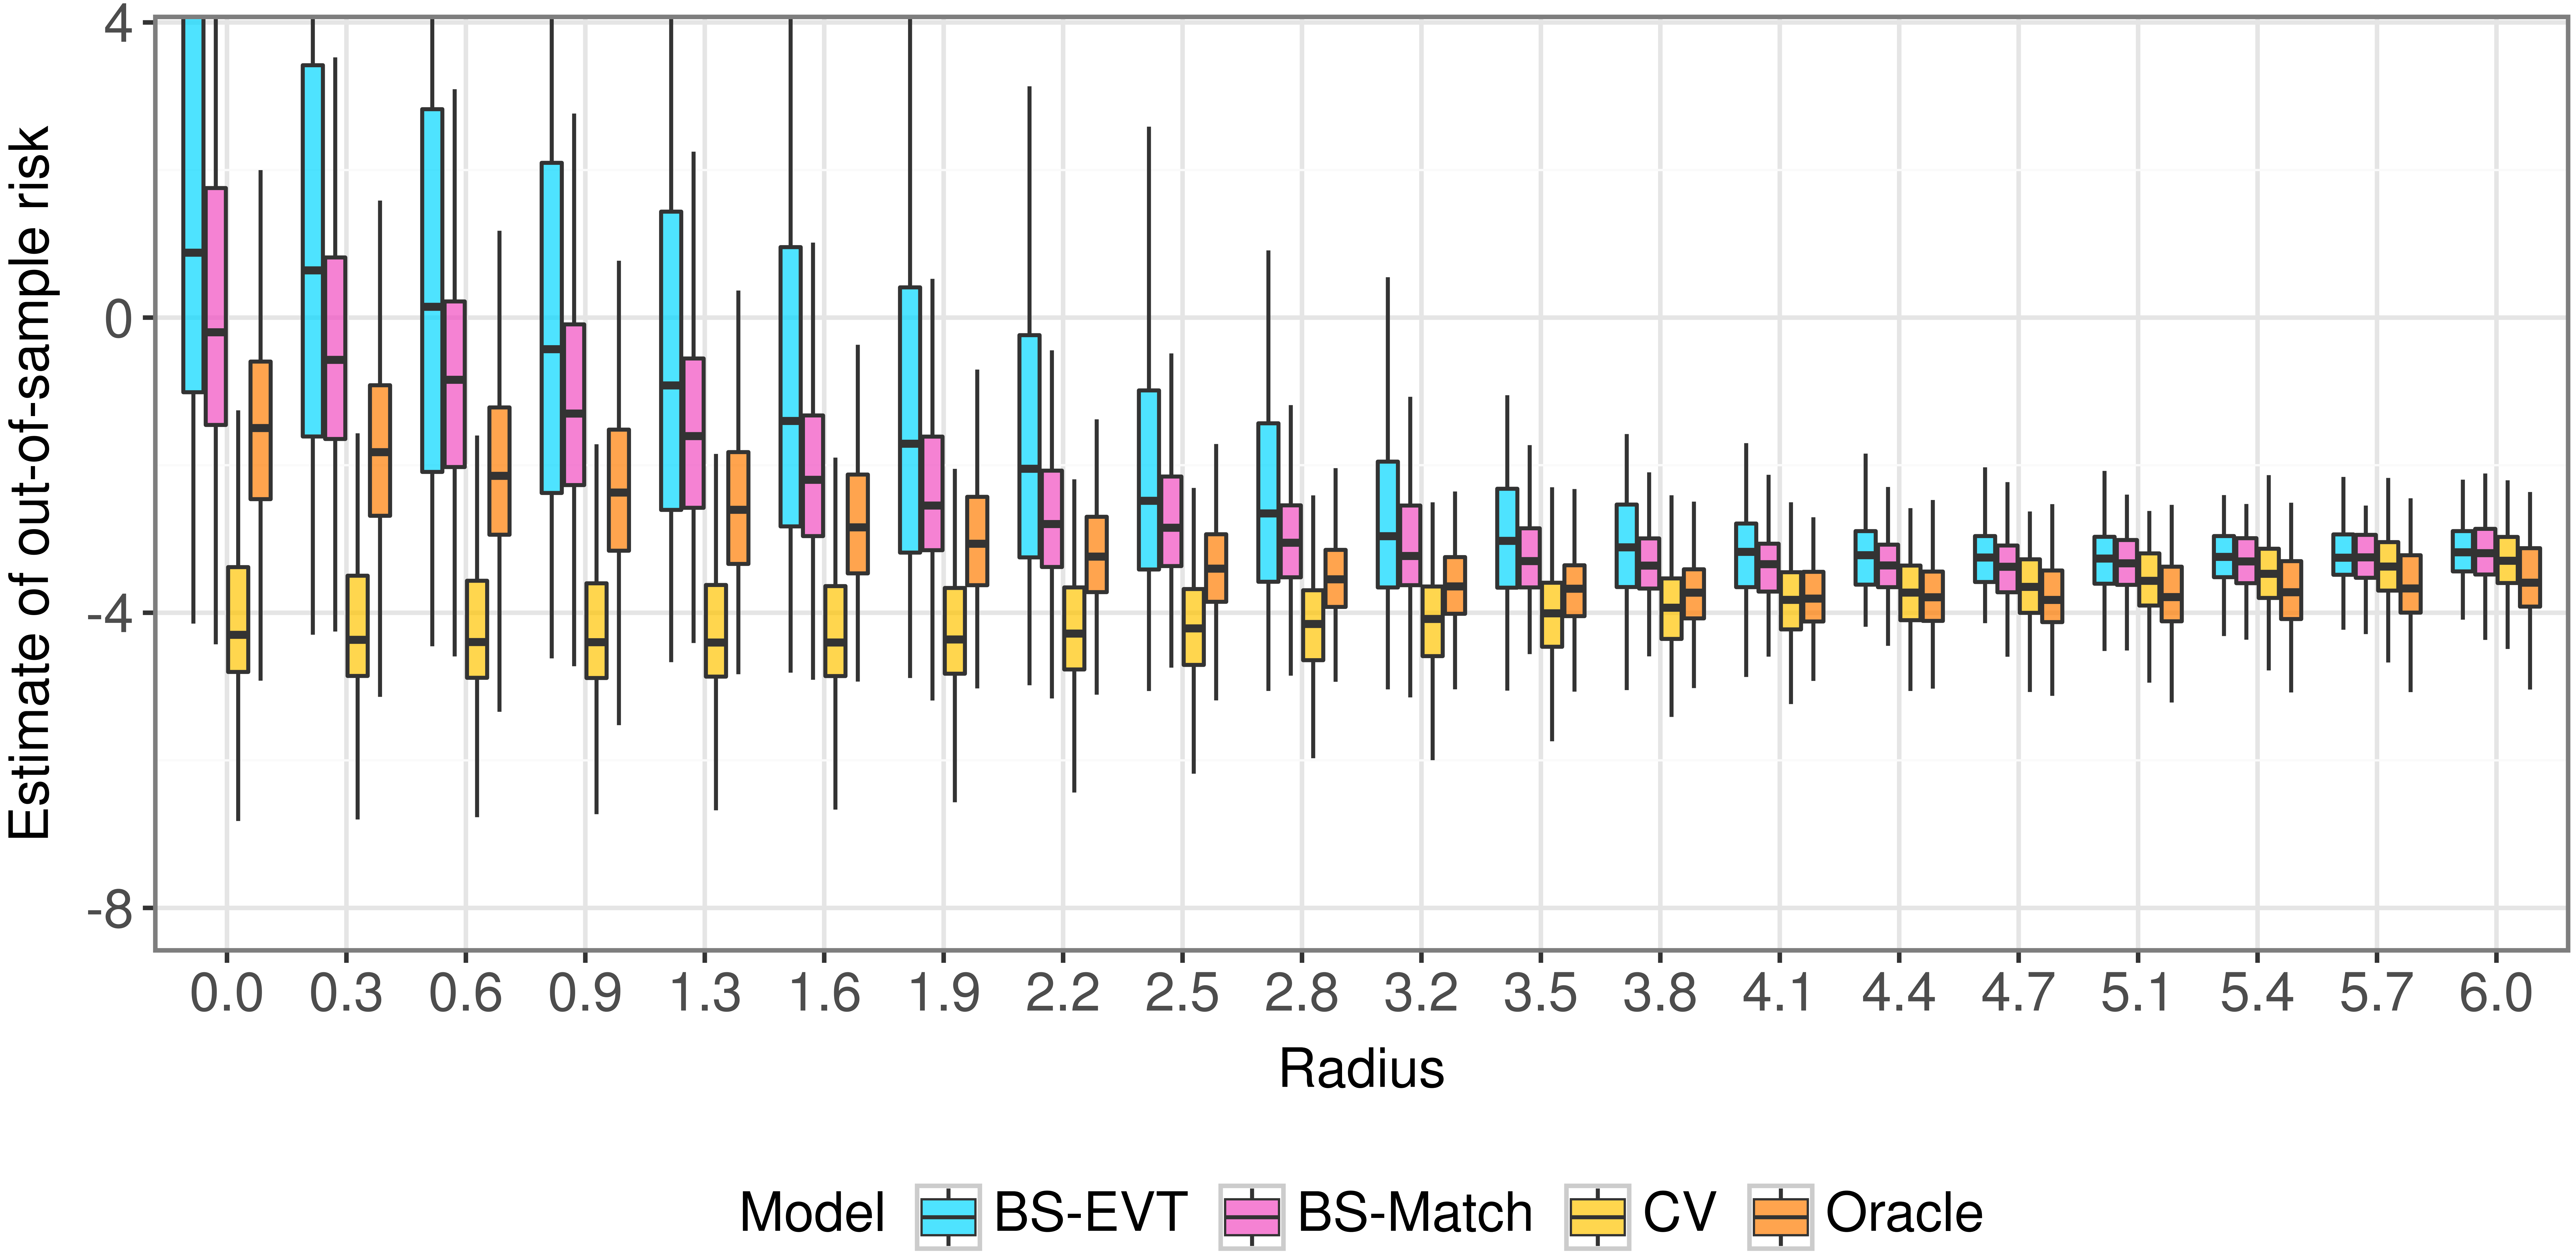

Take-away message

- Entropic risk estimation and optimization
  - Two practical approaches to reduce optimistic bias
- Future research:
  - Extend to CVaR
  - Solve exponential cones faster

Link to paper

Take-away message

- Entropic risk estimation and optimization
  - Two practical approaches to reduce optimistic bias
- Future research:
  - Extend to CVaR
  - Solve exponential cones faster

Link to paper